

MAT267: Advanced Ordinary Differential Equations

Rishabh Prakash

January 2022

Contents

0	An Unnecessary Remark From Me	3
1	Introduction	3
1.1	Solving Differential Equations	3
1.2	Standard Trick	4
2	Simple Examples	5
3	Useful Pictures	6
4	The Logistic Equation	8
4.1	Parameterising the general solution	9
5	Linear System of ODEs	10
5.1	Example	11
5.2	Superposition/Linearity Principle	11
5.3	Types of Systems	13
5.3.1	Saddle Point	13
5.3.2	Unstable Node	14
5.3.3	Stable Node	14
5.3.4	Center	14
5.3.5	Spiral	17
5.4	Repeated eigenvalues (Part I)	18
5.5	Repeated Eigenvalues (Part II)	19
6	Trace-Determinant Plane	20
7	Canonical Forms and Genericity	22
7.1	Canonical Forms	23
7.1.1	Example	23
7.2	Solving Jordan Blocks	24
8	Matrix Exponentials	24
9	Inhomogeneous Linear Systems	26
9.1	Duhamel's Principle	27
10	Linearisation	27
10.1	Example	28
11	Dynamical Systems	29

12 Existence and Uniqueness of solutions	30
12.1 Preliminary Theory	31
12.2 Existence of solutions	34
12.3 Uniqueness of solutions	40
13 Continuous Dependence	45
14 Flow of Dynamical Systems	50
14.1 Classification of Flow	53
14.1.1 Example 1	54
14.1.2 Example 2	54
14.1.3 Example 3	56
14.1.4 Example 4	57
14.1.5 Example 5	58
14.2 Stability	58
15 Techniques for non-linear systems	63
15.1 Nullclines	64
15.1.1 Example 1	64
15.2 Lyapunov functions	64
15.3 Gradient Systems	65
15.4 Hamiltonian Systems	66
15.5 Wronskian	67
16 Limit Sets	68
16.1 Examples	69
16.2 Poincaré-Bendixson Theorem	70
17 Techniques for solving ODES	74
17.1 Separation of variables	74
17.1.1 Example	75
17.2 Homogeneous Functions	75
17.2.1 Example	75
17.3 Exact Differential Equations	75
17.3.1 Example	77
17.3.2 Integrating Factors	77
17.3.3 Bernoulli Equations	78

0 An Unnecessary Remark From Me

I started making these notes mostly for myself. Forcing myself to write out the details was a way of making sure that I understood what was going on. There is this frustratingly elusive difference between understanding something and thinking you understand it. It's the difference between being able to follow a proof and being able to see the logical steps underlying it, seeing how you might have thought of it yourself.

Or maybe the difference is not so elusive after all. There are statements in these set of notes which I can feel to my very bones, while other statements remain on the surface, brushed away easily by time or poor memory. Not many statements fall in the first category; in fact it's a concerningly small amount. So despite my claim that this was made in an effort to understand everything, I don't think I have succeeded in that very much. Nevertheless, I am confident that I have a better understanding of the material as a result of creating these notes and pondering on the proofs and theorems. So I will recommend, as lightly as I can (because hey who the hell am I to tell you what to do), that you do something similar my dear reader. Spend enough time to make everything as simple as can be. Differential equations can get complex, but they should never get complicated. That is what I attempt in these notes. I don't think I achieve that as well as I could but hopefully it is enough for you to get started on. Good luck friend :)

1 Introduction

Differential equations are equations where the unknown is a function and the given equation gives us a relationship between the function and its derivative(s). Such equations are common in physics, economics, and even other areas of math where we know how a quantity changes (i.e. its derivatives) but not necessarily how to determine the quantity at any given instant. Perhaps the most important differential equation is Newton's (second) law:

$$ma = F$$

This equation tells us how force and acceleration, the second derivative of position, are interrelated. A similar example is that of the spring, where Hooke's law (in combination with Newton's law) tells us:

$$m \cdot x''(t) = -k \cdot x(t) \quad (1.1)$$

Without too much effort, we see that

$$x(t) = \cos\left(\sqrt{\frac{k}{m}} t\right)$$

forms a solution to (1.1), by which we mean that it satisfies the differential equation. Playing around a bit more we find that

$$x(t) = A \cos\left(\sqrt{\frac{k}{m}} t\right) + B \sin\left(\sqrt{\frac{k}{m}} t\right) \quad (1.2)$$

where A and B are any real numbers, all form solutions to (1.1). As it turns out, this is what *all* solutions to (1.1) look like, although this is something we will prove later.

The parameters A and B are (uniquely) determined by the initial conditions that the differential equation needs to satisfy (in general 2 unknowns will require 2 initial conditions). Hence A and B in some sense parametrise the solution space. In general, we would like to have the parameters cover the entire solution space and for each set of parameters to correspond to a different solution.

1.1 Solving Differential Equations

There are two types of differential equations: ordinary differential equations, ODEs, (where the unknown function is in one variable) and partial differential equations, PDEs, (where the unknown function is in several variables). Our goal always is to understand the quantity that the unknown function measures. Of

course the best way to understand this quantity is by finding the unknown function. For example, the solutions above (see [Equation 1.2](#)) instantly tell us that x has periodic behaviour, not something immediate from the differential equation itself.

This is why there is so much interest in solving differential equations. ODEs can sometimes be solved analytically (see the above example), PDEs, almost never. However, we can often analyse the differential equations themselves in order to make qualitative statements about the function/quantity being measured (how it changes, its limiting behaviour, points of equilibrium, etc.) and still learn meaningful information about the quantity being measured.

But perhaps before we go any further, we should define what ODEs are

Definition 1.1 (Ordinary Differential Equation). An ODE is an equation of the form

$$F(t, x(t), \dots, x^{(k)}(t)) = 0$$

where x is a vector valued function on an open interval $I \subset \mathbb{R}$ which is k -times continuous differentiable¹. This is known as the implicit form of the ODE.

If the codomain of F is \mathbb{R}^m with $m > 1$, we get a system of equations. Sometimes we can express the k -th derivative as a function of the lower order derivatives which gets the standard or explicit form of the differential equation.

We should also probably define what it means to solve an ODE.

Definition 1.2. A (classical) solution of an ODE $F(t, x(t), \dots, x^{(k)}(t)) = 0$ is a function $\phi : I \rightarrow \mathbb{R}^n$ of class C^k , where $I \subset \mathbb{R}$ is an open interval, such that $F(t, \phi(t), \dots, \phi^{(k)}(t)) = 0$ for all $t \in I$.

Note that not all ODEs can be solved. Consider, for example,

$$\left| \frac{dy}{dx} \right| + |y| + 1 = 0$$

Now, consider the following non-example. Suppose we are given the ODE

$$x + y \cdot y' = 0$$

We can define $y(x) = \sqrt{-(1+x^2)}$ and we see that

$$y'(x) = \frac{-x}{\sqrt{-(1+x^2)}}$$

which means that $x + y \cdot y'$ is certainly equal to 0. However, y is not defined in the reals! Therefore the given y is not a solution to the ODE.

Definition 1.3. The general solution of an ODE is a formula for all possible solutions.

The solution for x given in (1.2) is an example of a general solution. It is normally no easy task to find the general solution to a differential equation as one needs to prove that they have indeed found *all* possible solutions.

1.2 Standard Trick

There is a standard trick to turn a higher order ODE into a system of first-order ODEs which is a bit simple-minded but surprisingly useful.

Example 1.1. Suppose we are given the equation

$$mx'' = -kx - cx'$$

¹In principle we don't need continuity of the k -th derivative, only its existence. However this does make things a lot nicer

(one may think of this as introducing drag into our spring example). We can introduce new variables

$$x_1 = x, x_2 = x'$$

allowing us to write

$$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix}' = \begin{pmatrix} x_2 \\ -\frac{k}{m}x_1 - \frac{c}{m}x_2 \end{pmatrix}$$

■

In the general case of $F(t, x, x', \dots, x^{(k)}) = 0$, we define

$$\begin{aligned} x_1 &= x \\ x_2 &= x'_1 \\ x_3 &= x'_2 \\ &\vdots \\ x_k &= x'_{k-1} \end{aligned}$$

thus allowing us to write $F(t, x_1, x_2, \dots, x_k) = 0$.

2 Simple Examples

We list a few ‘classic’ examples of ODEs here.

The first example is perhaps the simplest one, one could think of

$$x' = 0 \tag{2.1}$$

It is easy to see that this is (only) solved by $x(t) = c$ where c is any real constant. Indeed the Fundamental Theorem of Calculus (and Mean Value Theorem) tell us that this is the general solution, thus we get a solution space of 1-dimension. This means that a single parameter dictates every possible solution, in this case that parameter is c .

The second example we look at is a slightly generalised version of this

$$x' = f(t) \tag{2.2}$$

Once again using the Fundamental Theorem of Calculus and the Mean Value Theorem, we find that the general solution is

$$x(t) = x_0 + \int_0^t f(s) ds$$

where once again our space of solution is one-dimensional, governed in this case by the constant x_0 .

The third example is one where things get interesting (and also an incredibly important example).

$$x' = ax \tag{2.3}$$

In fact this is a whole family of differential equations for every $a \in \mathbb{R}$. Some minor knowledge of calculus tells us that

$$x(t) = ce^{at}$$

is a solution to this differential equation (once again c can be any real constant). However it is not immediately obvious that this is the general solution to the differential equation. Let us show that this is the case.

Let $\tilde{x}(t)$ be any solution to the equation. Our claim is that \tilde{x} is a constant multiple of e^{at} . We show this by proving that the ratio of the two functions is always constant

$$\begin{aligned}\frac{d}{dt}(\tilde{x}(t)e^{-at}) &= \tilde{x}(t)(-ae^{-at}) + \tilde{x}'(t)(e^{-at}) \\ &= -a\tilde{x}(t)e^{-at} + a\tilde{x}(t)e^{-at} \\ &= 0\end{aligned}$$

This example illustrates a general principle for solving ODEs: guess and check. This is to say that it is often easier to guess an answer to an ODE and then verify that this solution works than find one analytically.

3 Useful Pictures

We've said many times that solving ODEs is a difficult, often impossible, task. We reiterate that here: solving ODEs is a difficult, often impossible, task. However, the ODE can give us a lot of information about its solution(s) which we often summarise in various pictorial formats.

A differential equation gives us a way of computing the slope of the tangent to the function at any given point (this is, after all, what the derivative measures).




Figure 1: Drawing the slope at one point

Thus one thing we can do to try and visualise the behaviour of the function is to find its derivative at every point ². This called the slope field or direction field. ³

Seeing the slope field, we can turn the question of solving a differential equation to a visual one: we need to find a function which is tangent to every slope line. Of course there are a variety of possible solutions depending on where one starts from (one would hope that after deciding where to start from, there is only *one* solution. we will see that in sufficiently nice conditions, this is the case). These solutions are also called integral curves. We illustrate a few examples below in [Figure 3](#).

We see that the slopes are all the same along any horizontal line. This is because the differential equation is independent of t . Such an ODE is called autonomous. For more general (i.e. non-autonomous) ODEs, we can still ask the question of what is the set of points (t, x) such that $f(t, x) = k$ for some fixed

²In practice, one often limits themselves to a finite set of points

³Note the multiple names which will be a common theme here. It seems that the only thing harder than solving a differential equation is naming things related to them consistently.




Figure 2: Drawing a slope field on a grid for $x' = 2x$

constant k . We call this set of points *isoclines* (in particular then horizontal lines are isoclines for autonomous ODEs).

Consider the equation

$$x' = 2x$$

It is clear that the case with $x = 0$ (horizontal line in Figure 3) is a special case as this is the only constant solution. In this case we call 0 an equilibrium or steady state or stationary point. Consider what happens we are slightly above $x = 0$. In this case as t goes to infinity (one often thinks of this time evolving), $x(t)$ gets further and further away from 0 (if this is modelling the position of a particle, then this shows that the particle grows further and further from its initial position (and at an increasing pace)). In particular if $x(t)$ is positive, it becomes more and more positive and if $x(t)$ is negative, it becomes more and more negative. We often represent this in what is called a phase line or phase portrait (see Figure 4).

The fact that solutions near 0 move away from 0 as t increases, means that although 0 is an equilibrium point, it is an unstable one, also called a source (the name inspired by a common interpretation of differentiation equations as modelling the flow of a fluid). The opposite is of course a stable equilibrium (or sink) where solutions within some neighbourhood of the equilibrium point tend towards the equilibrium (these are defined more precisely in Subsection 14.2: Stability). Such a case occurs if we consider Equation 2.3 for $a < 0$. In this case 0 is again an equilibrium point, but now it is a stable one.

Notice the difference in behaviour between $a > 0$ and $a < 0$. Behaviour for all $a > 0$ is qualitatively the same and the same holds true for $a < 0$. In such cases we say that the equation $x' = ax$ for $a > 0$ (or $a < 0$) is stable (this is different from 0 being a stable point; in this case we are talking about the family of equations being a stable one.).

It seems suggestive that we left out the case for $a = 0$ above. Indeed that is because the behaviour of the function is completely different for $a = 0$, since $x(t)$ is a constant in this case (this is exactly the very first example of an ODE we considered in Section 2: Simple Examples). It is at this point that 0 goes from being a source to a sink where the slightest change in a causes it to go one way or the other. We say then that $a = 0$ is a *bifurcation* in the 1-parameter family of equations $x' = ax$.




Figure 3: Slope field with some solutions (solutions illustrated in black)




Figure 4: Phase line/portrait for $x' = 2x$

4 The Logistic Equation

Imagine you are trying to model the growth of a population. We know that if a population is small and is in ideal conditions (easily accessible food, few predators, lots of space, etc.), a population will grow exponentially. However, we also know that this cannot continue forever. As a population grows larger and larger, it will start pushing towards the limits of the available resources. In fact if the population grows too large for its environment (for example if there's not enough food or too many predators), then one would expect the population to decrease. A simple equation, known as the logistic equation, that models this behaviour is

$$x' = ax \left(1 - \frac{x}{N} \right) \quad (4.1)$$

where $a \in \mathbb{R}$ is the growth rate of the population and N is the carrying capacity or the ideal population size. Notice that if x is small then $x' \approx ax$ and if $x > N$ (i.e. the population is greater than the carrying capacity) then $x' < 0$ so the population decreases. We immediately see that this equation is still autonomous (we still have no t in the equation) but it is no longer linear. This might already suggest that this is a somewhat more difficult problem than before (it is). Nevertheless, we can make our life a bit easier with one small assumption: without loss of generality we can take N to be 1 (we just choose appropriate units of x to make this work. One can think of x as modelling the proportion of the ideal population rather than the

actual size of the population). We thus define

$$f_a(x) = ax(1-x) \quad (4.2)$$

We see that $f_a(x)$ is 0 for $x = 0$ and $x = 1$ (these are hence our equilibrium points), positive if $x \in (0, 1)$ and negative otherwise. This already tells us that $x \equiv 0$ is an unstable solution while $x \equiv 1$ is a stable one. We can solve for other solutions to this differential equation (for other initial conditions) using a technique known as separation of variables.

We solve the ODE using separation of variables (see [Subsection 17.1: Separation of variables](#)).

$$\begin{aligned} x' &= ax(1-x) \\ \frac{x'}{ax(1-x)} &= 1 \\ \frac{1}{a} \int \frac{1}{x(1-x)} dx &= \int 1 dt \\ \frac{1}{a} \int \frac{1}{x} + \frac{1}{1-x} dx &= t + c \\ \frac{1}{a} \ln \left| \frac{x}{1-x} \right| &= t + c \\ \left| \frac{x}{1-x} \right| &= e^{at+ac} \\ \frac{x}{1-x} &= c_2 e^{at} \\ x &= \frac{c_2 e^{at}}{1 + c_2 e^{at}} \end{aligned}$$

Thus our solution is

$$x(t) = \frac{c_2 e^{at}}{1 + c_2 e^{at}} \quad (4.3)$$

4.1 Parameterising the general solution

Recall that that we also had special solutions to the equation, namely $x \equiv 1$ and $x \equiv 0$. We then see that $c_2 = 0$ recovers one of the solutions! Unfortunately the same is not true for the other solution (we would require $c_2 = \infty$). We can divide the numerator and denominator by c_2 to recover $x \equiv 1$ but we now lose $x \equiv 0$. This suggests that there might be a better parameterisation of the solutions. We see that

$$x_0 := x(0) = \frac{c_2}{1 + c_2}$$

We can solve for c_2 and substitute that into (4.3) to get

$$x(t) = \frac{x_0 e^{at}}{1 - x_0 + x_0 e^{at}} = \frac{x_0}{(1 - x_0)e^{-at} + x_0} \quad (4.4)$$

With this parametrisation, $x_0 = 0$ and $x_0 = 1$ get us the two special solutions.

Let's analyse this equation to see what we can learn of it. We will consider the case $a > 0$ since it is clear from the differential equation that cases for $a < 0$ are simply going to be reflections of the positive case (besides it is remarkably rare for populations to have a negative growth rate). Suppose the initial point x_0 is between 0 and 1. In this case, $x'(0)$ is positive so x will increase as t increases (as we already predicted) and x' will tend toward 0. However, looking at (4.4) we can see that no value of t will make $x(t) = 1$. Hence $x = 1$ forms an asymptote. Suppose x_0 is greater than 1. Then $x'(0)$ is negative. In this case, as before, x will continually decrease, approaching $x = 1$ but never intersecting it. The most interesting case is when $x_0 < 0$. In this case, we know that $x(0) = x_0 < 0$. Note that at $t = 0$ the denominator starts of as a positive number, 1. However as t grows larger and larger, $(1 - x_0)e^{at}$ will tend towards 0. Since x_0 is a fixed negative number, this means the denominator will be 0 at some point in time. In fact, we can work quite easily that this occurs at $t = \frac{1}{a} \ln(\frac{-x_0}{1-x_0})$. This means that in this case the function will blow up to $-\infty$ in finite time.⁴

⁴We let the reader decide how they feel about negative population sizes going to $-\infty$.




Figure 5: Solutions to the logistic equation (solutions in black), [source](#)

5 Linear System of ODEs

Now instead of looking at ODEs with real-valued solutions, we look at ODEs with vector-valued solutions. We often think of this as a *system* of first-order ODEs. As before, the simplest (and arguably most important) case will be the linear one.

A differential equation of the form

$$X' = A(t)X + f(t)$$

where $A(t)$ is an $n \times n$ matrix that may vary with t and f is a map from an interval $I \subset \mathbb{R}$ to \mathbb{R}^n is known as a linear system of ODEs. $A(t)$ is called the matrix of coefficients (its entries being the coefficients) and the function f is called the inhomogeneity. If $f = 0$ then the equation is called homogeneous and if $A(t)$ is a constant matrix then we have what is called a constant coefficient ODE. As usual, if $X' = F(X)$ for some F , then all X_0 that satisfy $F(X_0) = 0$ are called equilibrium points of a system (in the case of a linear system, $F(X) = A(t)X + f(t)$, but the definition applies more generally).

We will start by considering the simplest case of a linear system: a homogeneous, constant coefficient ODE (as we will see solving a homogeneous ODE allows us to solve the general system, see [Subsection 5.2: Superposition/Linearity Principle](#)). That is we are considering equations of the form

$$X' = AX \tag{5.1}$$

where A is an $n \times n$ matrix. In such cases $X = 0$ is always a equilibrium point. If $\det(A) = 0$ then we have a linear subspace of equilibria. Reducing further to case of A being 2×2 , if $\det(A) = 0$ we have a straight line of equilibrium points in \mathbb{R}^2 (we of course ignore the uninteresting case of $A = 0$).

As mentioned, we will start by simply considering the case in \mathbb{R}^2 . In other words our differential equation is of the form

$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \underbrace{\begin{pmatrix} a & b \\ c & d \end{pmatrix}}_A \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} ax + by \\ cx + dy \end{pmatrix}$$

The system would be easy to solve if $b = 0$ and $c = 0$ as we would left with equations of the form $x' = ax$ and $y' = dy$ which we already know how to solve. This occurs if A is a diagonal matrix. If A is diagonalizable, then we can apply a change of coordinates to make A diagonal and solve the system. In either case, we more or less get $x = x_0 e^{at}$ and $y = y_0 e^{dt}$ as our general solutions. Moreover, in the one-dimensional case the homogeneous linear system would simply be $x' = ax$, which again has the solution $x(t) = x_0 e^{at}$. This might inspire us to define the answer for the 2×2 case (and the general case) to be

$$X(t) = e^{At} X_0$$

Despite the nonsense that this looks like, there is a way of interpreting what it means to “exponentiate a matrix”; we use the Taylor expansion of e^x . In other words, we define

$$e^{At} := \sum_{k=0}^{\infty} \frac{t^k}{k!} A^k = I + tA + \frac{t^2}{2} A^2 + \frac{t^3}{3!} A^3 + \dots$$

Of course this is an infinite series so we need to decide whether it converges or not (and what convergence even means in this case), but at least each of the terms in the series makes sense.

As we will see later, this is in fact always convergent and does indeed solve our ODE and forms a general solution to our ODE. Moreover, we have that

$$\frac{d}{dt} e^{At} = A e^{At}$$

(Both of the above statements are proven in [Section 8: Matrix Exponentials](#)).

Another educated guess one may make, once again looking at the one-dimensional case, is that solutions will be of the form

$$X(t) = e^{\lambda t} v$$

where λ is a real number and v is a (constant) vector in \mathbb{R}^n , which are parameters to be determined. Assuming this to be case, we can substitute this in (5.1) to conclude that

$$\lambda e^{\lambda t} v = A(e^{\lambda t} v)$$

for all t . This implies that $Av = \lambda v$ or in other words that λ is an eigenvalue with eigenvector v .

5.1 Example

Let us try our hand with an example (with the second method for now). Suppose we are given

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \underbrace{\begin{pmatrix} 2 & 3 \\ 1 & 0 \end{pmatrix}}_A \begin{pmatrix} x \\ y \end{pmatrix} \quad (5.2)$$

Using our favourite method of finding eigenvalue/eigenvector pairs, we determine that the eigenvalues of A are 3 and -1 with eigenvectors $\begin{pmatrix} 3 \\ 1 \end{pmatrix}$ and $\begin{pmatrix} 1 \\ -1 \end{pmatrix}$ respectively. We then have two solutions

$$\begin{aligned} X_1(t) &= e^{3t} \begin{pmatrix} 3 \\ 1 \end{pmatrix} \\ X_2(t) &= e^{-t} \begin{pmatrix} 1 \\ -1 \end{pmatrix} \end{aligned}$$

However observe that any linear combination of X_1 and X_2 is also a solution. This leads us to superposition principle (also known as the linearity principle).

5.2 Superposition/Linearity Principle

Suppose $X_1(t)$ is such that it solves

$$X' = A(t)X + f_1(t)$$

and $X_2(t)$ solves

$$X' = A(t)X + f_2(t)$$

Then for real a_1, a_2 , $X(t) = a_1 X_1(t) + a_2 X_2(t)$ solves

$$X' = A(t)X + a_1 f_1(t) + a_2 f_2(t)$$




Figure 6: Phase Portrait of Equation 5.2. Solutions along eigenvectors in black. [Source](#)

This is easily verified by substituting the solution into the differential equation (and using the fact that multiplication with a matrix is linear). A consequence of this is that solutions to a homogeneous, linear system of ODEs forms a vector space.

Another consequence is the fact that the general solution $X'(t) = A(t)X + f(t)$ is given by

$$X(t) = X_{genhom}(t) + y(t)$$

where $X_{genhom}(t)$ is the general solution to the homogeneous system of equations $X' = AX$ and $y(t)$ is one *particular* solution to $X'(t) = A(t)X + f(t)$ (recall the similarities to solving a linear system of equations given by $AX = b$, where the solution is given by $b + \tilde{X}$ where \tilde{X} is the space of solutions that solves $AX = 0$). This means that if X maps to \mathbb{R}^n and $A \in \mathbb{R}^{n \times n}$ then the space of solutions to $X' = AX$ forms an n -dimensional vector space. In fact if $X_1(t), \dots, X_n(t)$ are (linearly independent) solutions to $X' = AX$ then

$$X(t) = a_1 X_1(t) + \dots + a_n X_n(t)$$

for $a_1, \dots, a_n \in \mathbb{R}$ is the general solution to $X' = AX$. This in fact holds in general but we will only prove this for the case when A can be diagonalised (for now) and later we will show it holds for all A .

Lemma 5.1 *The general solution to $X' = AX$ where X maps to \mathbb{R}^n and A is a diagonalisable $n \times n$ matrix is given by*

$$X(t) = a_1 X_1(t) + \dots + a_n X_n(t)$$

where the X_i themselves are linearly independent solutions (meaning $X_1(t), \dots, X_n(t)$ are linearly independent for all t) to the system of equations.

Proof. Let v_1, \dots, v_n be eigenvectors of A that form a basis for \mathbb{R}^n . Suppose their respective eigenvalues are $\lambda_1, \dots, \lambda_n$.

Suppose the initial value problem is given by

$$\begin{cases} X' = AX \\ X(0) = a_1 v_1 + \dots + a_n v_n \end{cases} \quad (5.3)$$

Then

$$Y(t) = a_1 e^{\lambda_1 t} v_1 + \dots + a_n e^{\lambda_n t} v_n$$

solves this initial value problem.

Suppose $Z(t)$ is another solution to this IVP. Since v_1, \dots, v_n is a basis for \mathbb{R}^n , we can write $Z(t) = b_1(t)v_1 + \dots + b_n(t)v_n$ where the b_i are real-valued functions. We know that $b_i(0) = a_i$ since $Z(0) = Y(0) = X(0)$. Also note that

$$\begin{aligned} Z'(t) &= b_1'(t)v_1 + \dots + b_n'(t)v_n \\ AZ(t) &= A(b_1(t)v_1) + \dots + A(b_n(t)v_n) \\ &= b_1\lambda_1v_1 + \dots + b_n\lambda_nv_n \end{aligned}$$

Equating coefficients of the v_i (which are uniquely determined since the v_i form a basis) we get that $b_i'(t) = \lambda_i b_i(t)$ and $b_i(0) = a_i$. We know that for each i , this is uniquely solved by $b_i(t) = e^{\lambda_i t} a_i$ implying that $Z = Y$. This allows us to conclude that the solutions to $X' = AX$ are uniquely determined by the initial value. \square

Remark 5.2. We know that the b_i above are differentiable as they are the composition of two differentiable functions: Z and the linear projection onto v_i which is a linear map with constant coefficients/entries (hence in particular is differentiable with respect to t).

5.3 Types of Systems

We can categorise the different systems of equation based on the eigenvalues.

5.3.1 Saddle Point

Suppose we are given

$$X' = \underbrace{\begin{pmatrix} 2 & 3 \\ 1 & 0 \end{pmatrix}}_A X \quad (5.4)$$

Then we know that its eigenvalues are $\lambda_1 = 3$ and $\lambda_2 = -1$ with eigenvectors $v_1 = (3, 1)$ and $v_2 = (1, -1)$ respectively (seen in [Subsection 5.1: Example](#)). Note that this is a case where the eigenvalues are of opposite sign. In this case the phase portrait would look like [Figure 7](#).




Figure 7: Saddle point, [source](#)

As time evolves, points on the line spanned by v_2 move towards the origin (exponentially) while points on the line spanned by v_1 move away (also exponentially). In this case we call v_2 the stable line and v_1 the

unstable line. Now consider a point that is not on either of these lines but is some linear combination of v_1 and v_2 . In this case as time evolves, the component for v_2 will shrink to 0 while the component for v_1 will shoot off to infinity leading to this above shape. This is called having a saddle point at 0. Any time we have eigenvalues of opposite sign we get something like above.

5.3.2 Unstable Node

The next question, of course, is of course is what happens if we have two eigenvalues of the same sign. We first consider the case of both eigenvalues being positive. As an example, we can consider

$$X' = BX \quad (5.5)$$

where $B = A + 2I$. Then the eigenvalues of B are $\lambda_1 = 5$ and $\lambda_2 = 1$ with v_1 and v_2 as before. In this case everything moves away from the origin giving us what is called an “unstable node at 0”. However since λ_1 is greater, the v_1 component of any point will increase much faster as $t \rightarrow \infty$. So eventually the paths will seem parallel to v_1 . Conversely as $t \rightarrow 0$, the v_1 component will also decrease to 0 faster than v_2 so the integral curves (aka solutions) become tangent to v_2 as t approaches 0. These features of the phase portrait can be verified in [Figure 8](#).




Figure 8: Unstable Node, **source**

5.3.3 Stable Node

We next look at the case when both eigenvalues are negative. Perhaps unsurprisingly, the picture will be quite similar to the previous with some slight modifications. Continuing our tradition of having a concrete example, we consider

$$X' = CX \quad (5.6)$$

where $C = A - 5I$. Our eigenvectors remain v_1 and v_2 as usual but their corresponding eigenvalues are now $\lambda_1 = -2$ and $\lambda_2 = -6$. Now as we evolve time, everything will approach the origin. Thus we call this situation having a “stable node at 0”. However now v_2 approaches the origin faster than v_1 so the integral curves will reflect this by becoming tangent to v_1 as $t \rightarrow \infty$. See [Figure 9](#).

5.3.4 Center

Of course not every real matrix will have eigenvalues. To be precise, a given matrix may not have *real* eigenvalues but it will certainly have *complex* eigenvalues. In fact complex eigenvalues for matrices with




Figure 9: Stable Node, [source](#)

real entries come in conjugate pairs. However we can close our eyes, pretend everything is real, and in the end split things into their real and complex components to get genuinely real solutions. Let us demonstrate what this means. Suppose we are given

$$X' = \underbrace{\begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}}_J X \quad (5.7)$$

In this case, the characteristic equation is $p(\lambda) = \lambda^2 + 1$. The roots of this polynomial are i and $-i$ which are thus our complex eigenvalues. As mentioned, we don't worry about these being complex and proceed as normal. We see that

$$J - iI = \begin{pmatrix} -i & -1 \\ 1 & -i \end{pmatrix}$$

and that the vector $v_1 = \begin{pmatrix} i \\ 1 \end{pmatrix}$ lies in its kernel hence is the corresponding eigenvector for $\lambda_1 = i$. Similarly

we conclude $v_2 = \begin{pmatrix} -i \\ 1 \end{pmatrix}$ is the eigenvector with eigenvalue $\lambda_2 = -i$ (in fact this can be concluded without any further calculations. If A is a matrix with real entries and has an eigenvector v with eigenvalue λ , then \bar{v} is an eigenvector with eigenvalue $\bar{\lambda}$ where \bar{v} is defined in the obvious way: taking the complex conjugate of each entry. See [Lemma 5.3](#) for details and proof). Thus our two solutions are

$$z_1(t) = e^{it} v_1, z_2(t) = e^{-it} v_2$$

Expanding z_1 using Euler's formula, we get

$$\begin{aligned} e^{it} v_1 &= (\cos t + i \sin t) \begin{pmatrix} i \\ 1 \end{pmatrix} \\ &= \begin{pmatrix} i \cos t - \sin t \\ \cos t + i \sin t \end{pmatrix} \\ &= \begin{pmatrix} -\sin t \\ \cos t \end{pmatrix} + i \begin{pmatrix} \cos t \\ \sin t \end{pmatrix} \end{aligned}$$

We claim that

$$X_1(t) = \begin{pmatrix} -\sin t \\ \cos t \end{pmatrix}, X_2(t) = \begin{pmatrix} \cos t \\ \sin t \end{pmatrix}$$

are (linearly independent) solutions to Equation 5.7 (note: we would have gotten the same solutions if we had chosen to expand z_2 instead). The fact that these are solutions can be verified easily by substituting them into (5.7). Linear independence follows from the fact that $X_1(t)$ and $X_2(t)$ are orthogonal for all t (verify by taking the dot product).




Figure 10: Center at 0, [source](#)

It is easy to see that the solutions will travel in a circular path: counterclockwise as time moves forward and clockwise as time moves backward. This is known as having a center at 0. This occurs anytime the eigenvalues are purely imaginary numbers. See Figure 10.

A few claims were made up in the previous example. Let us prove them formally.

Lemma 5.3 Suppose $A \in \mathbb{R}^{n \times n}$. Suppose v is an eigenvector with an eigenvalue λ . Then

- $\bar{\lambda}$ is an eigenvalue of A with eigenvector \bar{v}
- If λ is not real, then v is not in \mathbb{R}^n (it has complex entries). Moreover, $\text{Re}(v)$ and $\text{Im}(v)$ are linearly independent.

Proof. The first statement is easily verified by noting that $\bar{A} = A$ since A has real entries. Therefore

$$Av = \lambda v \Leftrightarrow \overline{Av} = \overline{\lambda v} \Leftrightarrow A\bar{v} = \bar{\lambda}\bar{v}$$

In order to verify the second statement suppose $v = u + iw$ where $u, w \in \mathbb{R}^n$. We will prove that u and w are linearly independent via contradiction. So suppose there exists real number s and t and some $v_0 \in \mathbb{R}^n$ such that $u = sv_0$ and $w = tv_0$. Then $v = u + iw = (s + it)v_0$. Since v_0 is a multiple of v , it must also be an eigenvector of A with eigenvalue λ . This means that

$$Av_0 = \lambda v_0$$

The left side of this equation is always in \mathbb{R}^n however if λ is not real then λv_0 will not be. Thus we get a contradiction if λ is non-real. This shows that u and w are linearly independent so neither can be 0, thus v must have complex entries. \square

Lemma 5.4 $Z(t)$ is a complex solution to $X' = AX$ (where A is a real matrix) if and only if $\operatorname{Re}(Z(t))$ and $\operatorname{Im}(Z(t))$ are also solutions.

Proof.

$$Z'_{\operatorname{Re}}(t) + iZ'_{\operatorname{Im}}(t) = Z'(t) = AZ(t) = AZ_{\operatorname{Re}}(t) + iZ_{\operatorname{Im}}(t)$$

□

Next we want to show that we have found the general solution to (5.7).

Lemma 5.5 The general solution (5.7) is given by $x(t) = a(-\sin t, \cos t) + b(\cos t, \sin t)$

Proof. Suppose $y(t) = (u(t), v(t))$ is another solution to the differential equation. Let $f(t) = (u(t) + iv(t))e^{-it}$. Then

$$\begin{aligned} f'(t) &= (u'(t) + iv'(t))e^{-it} - ie^{-it}(u(t) + iv(t)) \\ &= (-v(t) + iu(t))e^{-it} + e^{-it}(-iu(t) + v(t)) \\ &= 0 \end{aligned}$$

Therefore $y = \alpha e^{it}$ where α is some complex number implying that y is a linear combination of $X_1(t)$ and $X_2(t)$ as given above. □

5.3.5 Spiral

Suppose we have

$$X' = \begin{pmatrix} 2 & -1 \\ 1 & 2 \end{pmatrix} X \quad (5.8)$$

We find that the eigenvalues are $\lambda_1 = 2 + i$ and $\lambda_2 = 2 - i$ with corresponding eigenvectors $v_1 = \begin{pmatrix} i \\ 1 \end{pmatrix}$ and $v_2 = \begin{pmatrix} -i \\ 1 \end{pmatrix}$. Thus we find that a solution is given by

$$\begin{aligned} Z(t) &= e^{(2+i)t} \left(i \begin{pmatrix} 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} \right) \\ &= e^{2t} (\cos t + i \sin t) \left(i \begin{pmatrix} 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} \right) \\ &= e^{2t} \begin{pmatrix} -\sin t \\ \cos t \end{pmatrix} + ie^{2t} \begin{pmatrix} \cos t \\ \sin t \end{pmatrix} \end{aligned}$$

Thus the two solutions are given by the real and the imaginary parts:

$$X_1(t) = e^{2t} \begin{pmatrix} -\sin t \\ \cos t \end{pmatrix}, X_2(t) = e^{2t} \begin{pmatrix} \cos t \\ \sin t \end{pmatrix}$$

The general solution is then of course some linear combination of X_1 and X_2 .

Consider $X_1(t)$. As t increases the $(-\sin t, \cos t)$ component makes the point go around in the origin (with period 2π) while the e^{2t} causes the magnitude to increase. Thus we get a spiraling out. In order to determine the direction of the spiral (i.e. clockwise or counterclockwise) we can either investigate $(-\sin t, \cos t)$ (as t increases from 0 the x -coordinate becomes negative while the y -coordinate decreases from 1, therefore counterclockwise) or we can try a point and determine which way the tangent vector points. For example substituting $X = (1, 0)$ into (5.8) we find that $X' = (2, 1)$ implying that the spiral must be going counterclockwise.




Figure 11: Spiral (counterclockwise), [source](#)

5.4 Repeated eigenvalues (Part I)

Suppose we have

$$X' = \begin{pmatrix} 3 & -1 \\ 4 & -1 \end{pmatrix} X \quad (5.9)$$

where we will denote the matrix by A , as usual. The characteristic polynomial of A is $\lambda^2 - 2\lambda + 1$ which has a repeated root for $\lambda = 1$. Hence 1 is an eigenvalue with algebraic multiplicity 2. However $A - I$ is a matrix of rank 1, hence there is only one eigenvector namely $v = \begin{pmatrix} 1 \\ 2 \end{pmatrix}$. This gives one solution with

$$X_1(t) = ce^t \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$

However we expect a two dimensional solution space (a statement which we will justify later) hence we know there should be one other linearly independent solution.

From the study of linear algebra, we recall in this case A must have a basis in generalised eigenvectors. We compute that the generalised eigenvector of A is $w = \begin{pmatrix} 0 \\ -1 \end{pmatrix}$. Hence we might expect the second solution to be of the form $X(t) = \alpha(t)v + \beta(t)w$ with $\beta \neq 0$ (the case with $\beta = 0$ is covered by X_1 above). Assuming that $X(t)$ does solve our differential equation we see that

$$\begin{aligned} X'(t) &= A(\alpha(t)v + \beta(t)w) \\ &= \alpha(t)Av + \beta(t)Aw \\ &= \alpha(t)\lambda v + \beta(t)(\lambda w + v) \\ &= (\lambda\alpha(t) + \beta(t))v + \lambda\beta(t)w \end{aligned}$$

On other hand we also know that

$$X'(t) = \alpha'(t)v + \beta'(t)w$$

equating coefficients (since (v, w) is a basis, the coefficients are unique), we get the following system of

equations

$$\begin{cases} \alpha' = \lambda\alpha + \beta \\ \beta' = \lambda\beta \end{cases}$$

The second equation is easily solved by $\beta(t) = ce^{\lambda t}$ where c can be any constant (in fact, as we know, the equation is *only* solved by this). Recall that we are trying find a *particular* solution to the differential equation (one that is linearly independent of X_1). Hence, for simplicity, we can take $c = 1$ above giving us $\beta(t) = e^{\lambda t}$.

This, however, leaves us to solve for α . One might expect that we can simply use the theory built up so far since this is simply a homogeneous, linear system of equations. However, one can check that this system is exactly the case we are considering right now: the case with repeated eigenvalues with a basis in generalised eigenvectors. Therefore we will have to resort to some combination of guesswork and being clever.

In this case, we recall that the homogeneous equation $\alpha' = \lambda\alpha$ would be solved by $\alpha(t) = ce^{\lambda t}$ where c is a constant. One might guess that if instead we vary the constant in some appropriate manner (i.e. make c a function of t), we might be able to solve for $\alpha'(t) = \lambda\alpha(t) + e^{\lambda t}$. Therefore suppose $\alpha(t) = y(t)e^{\lambda t}$. Assuming this to be a solution, we get

$$\alpha'(t) = y'(t)e^{\lambda t} + \lambda y(t)e^{\lambda t}$$

It is then clear that if $y'(t) = 1$ then we have a solution. In other words we can take $y(t) = ct$ for any constant c . Once again we are only interested in one particular solution. Therefore, we take the simplest one to conclude that $\alpha(t) = te^{\lambda t}$. We then finally have a second, linearly independent solution to our ODE:

$$X_2(t) = te^{\lambda t}v + e^{\lambda t}w$$

In other words the general solution to the ODE is given by

$$X(t) = a_1X_1(t) + a_2X_2(t)$$

where a_1, a_2 are constants determined by the initial conditions, as usual.

Let us substitute the specific values of λ, v and w to get a clearer picture of the solutions (everything thus far will hold for any repeated eigenvalues with a basis in generalised eigenvectors). We have that our solution is given by

$$\begin{aligned} X(t) &= a_1e^t \begin{pmatrix} 1 \\ 2 \end{pmatrix} + a_2 \left(te^t \begin{pmatrix} 1 \\ 2 \end{pmatrix} + e^t \begin{pmatrix} 0 \\ -1 \end{pmatrix} \right) \\ &= a_1e^t \begin{pmatrix} 1 \\ 2 \end{pmatrix} + a_2e^t \left(t \begin{pmatrix} 1 \\ 2 \end{pmatrix} + \begin{pmatrix} 0 \\ -1 \end{pmatrix} \right) \\ &= (a_1e^t + a_2te^t) \begin{pmatrix} 1 \\ 2 \end{pmatrix} + a_2e^t \begin{pmatrix} 0 \\ -1 \end{pmatrix} \end{aligned}$$

We see that as $t \rightarrow \infty$ the component corresponding to $\begin{pmatrix} 1 \\ 2 \end{pmatrix}$ increases much more quickly than its partner, implying that the solution curves become increasing parallel to the line spanned by v . On the other hand as $t \rightarrow -\infty$, this component also approaches 0 more quickly, so in the limit the solution curves become tangent to v at the origin. If the eigenvalue were negative, then things remain essentially identical except the arrows are reversed. See [Figure 12](#) for phase portrait.

5.5 Repeated Eigenvalues (Part II)

It's possible that one has repeated eigenvalues *and* a basis of eigenvectors. An example is

$$X' = \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix} X \tag{5.10}$$




Figure 12: Repeated (positive) eigenvalues, [source](#)

where the standard basis vectors are eigenvectors themselves. In this case every vector is an eigenvector. Although the example might seem like a special case, it in fact highlights the general case since the assumptions on the matrix mean that it is a multiple of the identity. If it's a positive multiple of the identity (like the example above), then everything tends away from the origin in the direction parallel to itself. Note how extremely unstable this is since the mildest of perturbations cases the long term behaviour to be completely different. If we have a negative multiple of the identity, then the situation reverses and everything approaches the origin (and of course this is a very stable situation: regardless of where you start, you tend towards the origin).

6 Trace-Determinant Plane

We are interested in classifying the different dynamical systems, roughly based on what the phase portraits look like. For example, all centers are roughly the same (the only thing that varies is the frequency and the direction of the rotation); all saddle points are essentially the same (up to some rotation and stretching), etc. What we realise is that almost all of this information is determined by the eigenvalues and in particular by the sign of the eigenvalues, whether they are real or complex, etc. If we can then work out the relationships between the eigenvalues (ideally without solving for them), we can classify the dynamical systems more easily.

This is where we remember from our study of linear algebra that the eigenvalues of a matrix can be found by looking at the roots of the characteristic polynomial of the matrix and in the 2×2 case, this polynomial is completely determined by the trace and determinant. Indeed we have that the characteristic polynomial p_A of a 2×2 matrix A is given by

$$p_A(\lambda) = \lambda^2 - \text{Tr}(A)\lambda + \det(A)$$

We know then, for example, that if the discriminant of the above quadratic is positive, then we have 2 distinct real eigenvalues. If in addition the determinant is positive, then we know the eigenvalues share sign. Finally we can use the trace to determine what exactly their sign is. This gives us a near complete description of the qualitative behaviour of A . Repeating this analysis for the other cases, we summarise our findings below.

Eigenvalues	Normal form	How to detect	Shape
$\lambda_1 > \lambda_2 > 0$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$	$\text{Tr}(A)^2 > 4 \det(A)$	Unstable node
$\lambda_1 > 0 > \lambda_2$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$	$\det(A) < 0$	Saddle
$0 > \lambda_1 > \lambda_2$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$	$\det(A) > 0, \text{Tr}(A) < 0$	Stable node
$\lambda = \alpha + i\beta, \alpha > 0$	$\begin{pmatrix} \alpha & -\beta \\ \beta & \alpha \end{pmatrix}$	$\text{Tr}(A) > 0, 4 \det(A) > \text{Tr}(A)^2$	Spiral source
$\lambda = \alpha + i\beta, \alpha < 0$	$\begin{pmatrix} \alpha & -\beta \\ \beta & \alpha \end{pmatrix}$	$\det(A) > 0, \text{Tr}(A) < 0$ $4 \det(A) < \text{Tr}(A)^2$	Spiral sink

Table 1: Generic cases for $X' = AX$

These are the generic cases (the word generic is defined more precisely in [Section 7: Canonical Forms and Genericity](#)), but roughly speaking, this is what ‘most’ 2×2 matrices look like: they either have 2 distinct eigenvalues or 2 complex eigenvalues (which are conjugates). This can be verified by the fact that the above cases have covered almost all of the trace-determinant plane.

There do remain, however, some degenerate or exceptional cases that we still need to verify. In general these occur when we have an equality of some kind (in other words something is equal to 0 and as one can imagine this is rarely a good thing). These are summarised in [Table 2](#).

Eigenvalues	Normal form	How to detect	Shape
$\lambda = i\beta$	$\begin{pmatrix} 0 & -\beta \\ \beta & 0 \end{pmatrix}$	$\text{Tr}(A) = 0, \det(A) > 0$	Center
$\lambda_1 = \lambda_2 > 0$	$\begin{pmatrix} \lambda_1 & 1 \\ 0 & \lambda_1 \end{pmatrix}$	$\text{Tr}(A)^2 = 4 \det(A),$ $\text{rank}(A - \lambda_1 I) = 1, \text{Tr}(A) > 0$	Unstable, source
$\lambda_1 = \lambda_2 < 0$	$\begin{pmatrix} \lambda_1 & 1 \\ 0 & \lambda_1 \end{pmatrix}$	$\text{Tr}(A)^2 = 4 \det(A),$ $\text{rank}(A - \lambda_1 I) = 1, \text{Tr}(A) < 0$	Stable, sink
$\lambda_1 = \lambda_2 > 0$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_1 \end{pmatrix}$	$A = \lambda_1 I, \text{Tr}(A) > 0$	(Very) unstable, source
$\lambda_1 = \lambda_2 < 0$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_1 \end{pmatrix}$	$A = \lambda_1 I, \text{Tr}(A) < 0$	Stable, sink
$\lambda_1 > 0, \lambda_2 = 0$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & 0 \end{pmatrix}$	$\det(A) = 0, \text{Tr}(A) > 0$	Unstable
$\lambda_1 < 0, \lambda_2 = 0$	$\begin{pmatrix} \lambda_1 & 0 \\ 0 & 0 \end{pmatrix}$	$\det(A) = 0, \text{Tr}(A) < 0$	Stable
$\lambda_1 = \lambda_2 = 0$	$\begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}$	$\text{Tr}(A) = \det(A) = 0, \text{rank}(A) = 1$	Unstable
$\lambda_1 = \lambda_2 = 0$	$\begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix}$	$A = 0$	Constant

Table 2: Degenerate cases for $X' = AX$

We say a bifurcation occurs in an ODE when change the parameters a bit causes the qualitative behaviour to change drastically (for example changing from a source to a sink). This is a definition we will make more precise later. This information is often summarised in a bifurcation diagram. Note that the trace-determinant plane is an example of a bifurcation diagram where the axes and parabola are used

to group similarly behaved equations together and bifurcations occur when we go from one region to another.




Figure 13: Trace determinant plane, Source: Differential Equations, Dynamical Systems, and an Introduction to Chaos

7 Canonical Forms and Genericity

Definition 7.1 (Hyperbolicity). The origin is called a hyperbolic equilibrium for the ODE $X' = AX$ if all eigenvalues of A have non-zero real part.

In the case when A is a 2×2 matrix, hyperbolicity holds whenever $\det(A) < 0$ or if $\det(A) > 0$ and $\text{Tr}(A) \neq 0$ (see trace-determinant plane).

Definition 7.2 (Genericity). A property (for example of matrices) is called generic if it is satisfied on a dense, open subset (of $\mathbb{R}^{n \times n}$ for example).

Theorem 7.3 *The property of having n distinct eigenvalues is generic for $n \times n$ matrices. In other words the subset of $n \times n$ matrices which have n distinct eigenvalues is open and dense in $\mathbb{R}^{n \times n}$.*

Corollary 7.3.1 (Cayley-Hamilton Theorem) *Let A be an $n \times n$ matrix and let $p_A(\lambda)$ be its characteristic polynomial. Then $p_A(A) = 0$.*

Proof. The statement is easily verified for diagonal matrices (the diagonal entries are the eigenvalues and also the zeroes to the characteristic polynomial). Also note that $p_S(T^{-1}ST) = T^{-1}p_S(S)T$ for every invertible $T \in \mathbb{R}^{n \times n}$ and every $S \in \mathbb{R}^{n \times n}$. Thus the statement holds true not only for diagonal matrices but for diagonalisable matrices in general. Since the subset of matrices with n distinct eigenvalues is dense in $\mathbb{R}^{n \times n}$, there exists a sequence of diagonalisable matrices $(A_i)_{i=1}^{\infty}$ (with n distinct eigenvalues) that converges to A . In particular this means that the entries of the sequence of the matrices approach the entries of A . The coefficients of the characteristic polynomial of a matrix depend continuously on the entries of the matrix. Thus as $A_i \rightarrow A$ we have that $p_{A_i} \rightarrow p_A$. Since $p_{A_i}(A_i) = 0$ we get $p_{A_i}(A_i) \rightarrow p_A(A) = 0$. \square

7.1 Canonical Forms

Suppose we are studying our good old friend $X' = AX$. Suppose we set $X = TY$ where T is some invertible matrix. Then

$$Y' = (T^{-1}X)' = T^{-1}X' = T^{-1}AX = T^{-1}ATY$$

Defining $C := T^{-1}AT$ we get another differential equation $Y' = CY$. Importantly if we choose T cleverly we can make C a ‘simple’ matrix so that $Y' = CY$ is easily solved. Once we have a solution for Y , we can easily find X , the solution to our original differential equation, since by definition $X = TY$. What does it mean for C to be simple? Ideally we would want it to be diagonal of course. But this of course not always possible. Thus we go for the next best thing: the Jordan Canonical Form (JCF). Technically even this is not always possible over the reals. However, what we can put the matrix into JCF as if we were over the complex numbers and use the complex solutions to get (pairs of) real solutions, as we’ve done before.

The Jordan Canonical form allows to perform a change of basis to write every matrix as a block matrix of the form

$$C = \begin{pmatrix} J_1 & & \\ & \ddots & \\ & & J_k \end{pmatrix}, J_i = \begin{pmatrix} \lambda_i & 1 & & \\ & \lambda_i & \ddots & \\ & & \ddots & 1 \\ & & & \lambda_i \end{pmatrix}$$

where each J_i is an $l_i \times l_i$ matrix. The J_i ’s are called Jordan blocks. A diagonal matrix is a special case of the above form where each Jordan block is 1×1 .

Remark 7.4. The same eigenvalue could appear in different Jordan block. The total number of times that λ appears in C (along the diagonal of course) is exactly the algebraic multiplicity of λ .

Remark 7.5. If $\lambda \in \mathbb{C}$ is an eigenvalue, we know that $\bar{\lambda}$ is as well. This also means that \bar{J} is a Jordan block in A .

7.1.1 Example

Suppose we know that a 5×5 matrix has an eigenvalue λ that has algebraic multiplicity 5 as well. Here are some possible Jordan blocks.

$$\begin{pmatrix} \lambda & 1 & & & \\ & \lambda & 1 & & \\ & & \lambda & 1 & \\ & & & \lambda & 1 \\ & & & & \lambda \end{pmatrix}, \begin{pmatrix} \lambda & & & & \\ & \lambda & 1 & & \\ & & \lambda & 1 & \\ & & & \lambda & 1 \\ & & & & \lambda \end{pmatrix}, \begin{pmatrix} \lambda & 1 & & & \\ & \lambda & & & \\ & & \lambda & 1 & \\ & & & \lambda & 1 \\ & & & & \lambda \end{pmatrix}, \begin{pmatrix} \lambda & & & & \\ & \lambda & 1 & & \\ & & \lambda & 1 & \\ & & & \lambda & 1 \\ & & & & \lambda \end{pmatrix}$$

The question then is how do we know which of the Jordan blocks we have? This takes a bit more work.

Note that any Jordan block of size l can be written as

$$J = \lambda I + N$$

where N is a nilpotent matrix (this means that $N^l = 0$ but $N^{l-1} \neq 0$). As a consequence then we can distinguish which of the canonical forms by looking at the kernel of $(A - \lambda I)$ raised to various power. To be specific, $\dim \ker((A - \lambda I)^m)$ is the number of Jordan blocks of size less than or equal to m . In first case given above, we see that $\dim \ker(A - \lambda I) = 1$. This is enough to characterise this matrix completely since this is the only way to get a 5×5 Jordan block with one eigenvector. However for the middle two examples, we see that $\dim \ker(A - \lambda I) = 2$ for both of them (they both have two eigenvectors). We then try the next power: $\dim \ker(A - \lambda I)^2$ is 3 for the first of the two (second matrix in the list) and 4 for the other matrix (third in the list).

7.2 Solving Jordan Blocks

The question now boils down to how to solve a system of the form $Y' = CY$ where C is in canonical form and possibly complex. One thing to note is that Jordan decomposition allows us to break the space into invariant subspaces (that is after all one of the motivations for the decomposition). What this means is that solutions for the Jordan blocks can be found independently and ‘stacked’ up together (think of the analogy with diagonal matrices where we get n distinct equations that can be solved independently and when combined give us a solution to the entire system). Thus we only need solutions to $Y' = JY$ where the matrix J is of the form

$$J = \begin{pmatrix} \lambda & 1 & & \\ & \lambda & \ddots & \\ & & \ddots & 1 \\ & & & \lambda \end{pmatrix} = \lambda I + \underbrace{\begin{pmatrix} 0 & 1 & & \\ & 0 & \ddots & \\ & & \ddots & 1 \\ & & & 0 \end{pmatrix}}_N$$

Suppose $Y(t) = e^{t\lambda} Z(t)$ is a solution to $Y' = JY$. Then

$$\begin{aligned} Z'(t) &= -\lambda e^{-t\lambda} Y + e^{-t\lambda} Y' \\ &= -\lambda e^{-t\lambda} Y + e^{-t\lambda} (\lambda Y + NY) \\ &= N(e^{-t\lambda} Y) \\ &= NZ \end{aligned}$$

Thus we solve for $Z' = NZ$, where N is a nilpotent matrix as shown above. This is equivalent to writing

$$\begin{pmatrix} z_1 \\ \vdots \\ z_{n-1} \\ z_n \end{pmatrix}' = \begin{pmatrix} z_2 \\ \vdots \\ z_n \\ 0 \end{pmatrix}$$

Since $z_n' = 0$, we know that $z_n(t) = c_n$ where c_n is some arbitrary constant. Then $z_{n-1} = c_n t + c_{n-1}$ where c_{n-1} is again some real constant. Continue this way we get $z_1(t) = p(t)$ where p is some polynomial of degree $n-1$ (since n -th derivative is 0). This is the general solution (which makes sense, the space of polynomials of degree at most $n-1$ is of dimension n). Then $Y(t) = e^{t\lambda} Z(t)$ solves $Y' = JY$, where we may split Y into its real and imaginary components if necessary. We have thus found a solution for all linear systems! A pat on the back is well-deserved but we postpone that for after a discussion of matrix exponentials.

8 Matrix Exponentials

There is a second method of reaching the same answer with a bit less work (or rather with most of the work swept under the rugs of past theorems and lemmas). We claim that a (in fact the) solution to $Y' = (\lambda I + N)Y$ (where N is nilpotent) with $Y(0) = y_0$ is given by

$$e^{(\lambda I + N)t} y_0 = e^{\lambda t} \sum_{k=0}^{\infty} \frac{t^k}{k!} N^k y_0$$

Normally we would have to worry about convergence of infinite series, however since N is nilpotent the above series is finite. Once again we have a product of $e^{\lambda t}$ with some polynomial in t as we did before (the coefficients come from y_0). This might lead one to conjecture that the general solution to $X' = AX$ for *any* A is given by

$$e^{tA}$$

where once again we are more or less using the exponential notation as shorthand for its Taylor expansion. We first need to ensure that this actually makes sense, that is we really do have convergence. We first need a norm. It is a fact that in finite dimensions that all norms (and also all inner products) are equivalent (any norm can be bounded by a multiple of another norm and similarly with inner products). Thus we can quite frankly choose any norm we want. We will choose one that is quite common in this setting, called the operator norm which is defined as

$$\|A\| := \sup_{\|v\| \leq 1} \|Av\| = \sup_{v \in \mathbb{R}^n, v \neq 0} \frac{|Av|}{|v|}$$

We may equivalently define it as the largest singular value of A (in the appropriate setting).

By definition we have the fact that $\|A\| \leq \|A\| \|v\|$. This leads to the lovely fact that $\|AB\| \leq \|A\| \|B\|$ since

$$\|AB\| = \sup_{\|v\| \leq 1} \|ABv\| \leq \|A\| \sup_{\|v\| \leq 1} \|Bv\| = \|A\| \|B\|$$

Remark 8.1. $\mathbb{R}^{n \times n}$ with the usual operations of matrix addition and multiplication and the operator norm forms a Banach algebra.

Proposition 8.2 *The series*

$$e^{tA} := \sum_{k=0}^{\infty} \frac{t^k}{k!} A^k$$

converges absolutely.

Proof. We need to show that

$$\sum_{k=0}^{\infty} \left\| \frac{t^k}{k!} A^k \right\|$$

is finite. With our previous statement this is easy to see, since

$$\begin{aligned} \sum_{k=0}^{\infty} \left\| \frac{t^k}{k!} A^k \right\| &\leq \sum_{k=0}^{\infty} \frac{|t|^k}{k!} \|A^k\| \\ &\leq \sum_{k=0}^{\infty} \frac{|t|^k}{k!} \|A\|^k \\ &= e^{|t| \|A\|} \end{aligned}$$

□

We also have the following lovely statements about matrix exponentials. The first property is often called the semigroup property and the latter property is an if and only if (we only prove one direction).

Lemma 8.3 *Let $s, t \in \mathbb{R}$ and $A, B \in \mathbb{R}^{n \times n}$. Then*

1. $e^{(t+s)A} = e^{tA} e^{sA}$
2. $e^{t(A+B)} = e^{tA} e^{tB}$ if $AB = BA$

Proof. For the first statement (and for the second statement, the proofs are near identical), we see that

$$\begin{aligned} e^{(t+s)A} &= \sum_{k=0}^{\infty} \left(k! \sum_{i+j=k} \frac{s^i}{i!} \frac{t^j}{j!} \right) \frac{A^k}{k!} \\ &= \sum_{k=0}^{\infty} \left(\sum_{i+j=k} \frac{s^i}{i!} A^i \frac{t^j}{j!} A^j \right) \\ &= \left(\sum_{i=0}^{\infty} \frac{s^i}{i!} A^i \right) \left(\sum_{j=0}^{\infty} \frac{t^j}{j!} A^j \right) \\ &= e^{sA} e^{tA} \end{aligned}$$

The same proof works for the second case since we can use the binomial theorem again as the matrices A, B commute. \square

A consequence of the second property is that the exponential of a matrix is always invertible (we take $B = -A$).

Finally we want to make a comment on the differentiability of this map to make sure that it is what we expect.

Lemma 8.4 *Let A be an $n \times n$ matrix. Then we can define a map $t \mapsto e^{tA}$. The derivative of this map is Ae^{tA} . In symbols we write*

$$\frac{d}{dt} e^{tA} = Ae^{tA}$$

Proof.

$$\begin{aligned} \frac{d}{dt} e^{tA} &= \lim_{h \rightarrow 0} \frac{\exp((t+h)A) - \exp(tA)}{h} \\ &= \lim_{h \rightarrow 0} \frac{\exp(tA) \exp(hA) - \exp(tA)}{h} \\ &= \exp(tA) \lim_{h \rightarrow 0} \frac{\exp(hA) - I}{h} \\ &= \exp(tA) \lim_{h \rightarrow 0} \frac{1}{h} \left(hA + \frac{h^2}{2} A^2 + \frac{h^3}{3!} A^3 + \dots \right) \\ &= \exp(tA) A \end{aligned}$$

\square

Much like the exponential function with real numbers, matrix exponentiation gives us a unique solution to a differential equation. In particular, $X(t) = e^{tA} X_0$ is the unique solution to the initial value problem $X' = AX$ with $X(0) = X_0$. The proof is the exact same as with real numbers: suppose $Z(t)$ is another solution. Define $W(t) = Z(t)e^{-tA}X_0$. Then

$$\begin{aligned} W'(t) &= \frac{d}{dt} (Z(t)e^{-tA}X_0) \\ &= (-Z(t)Ae^{-tA})X_0 + (Z'(t)e^{-tA})X_0 \\ &= (-Ae^{-tA}Z(t) + e^{-tA}AZ(t))X_0 \\ &= 0 \end{aligned}$$

(we use the fact that A commutes with its exponential). Therefore $W(t)$ is constant. Since $Z(0) = e^{0A}X_0 = X_0$, Z and X agree everywhere and we are done.

9 Inhomogeneous Linear Systems

Suppose we have an equation of the form

$$x' = Ax + f(t) \tag{9.1}$$

Then we know from the superposition principle (see [Subsection 5.2: Superposition/Linearity Principle](#)) that the general solution to this system is given by

$$x(t) = y(t) + e^{tA}v$$

where $y(t)$ is a particular solution to the ODE and v is some vector that is determined by the initial conditions. Thus our goal is to find just *one* solution, y , to this ODE. This leads us to Duhamel's Principle.

9.1 Duhamel's Principle

We will guess that $y(t) = e^{tA}v(t)$ is a solution where $v(t)$ is a function to be determined (this is the technique of variation of constants). Assuming y is a solution we can plug this in (9.1) to find the left and right hand sides are

$$y' = Ae^{tA}v(t) + e^{tA}v'(t) \quad (\text{LHS})$$

$$Ay + f(t) = Ae^{tA}v(t) + f(t) \quad (\text{RHS})$$

Equating the two we find that

$$v'(t) = e^{-tA}f(t)$$

Hence by the Fundamental Theorem of Calculus we find that

$$v(t) = \int_0^t e^{-sA}f(s)ds$$

(recall we only need a particular solution so we can ignore the constant of integration by setting it to 0). Therefore our particular solution y is given by

$$y(t) = e^{tA}v(t) = e^{tA} \int_0^t e^{-sA}f(s)ds = \int_0^t e^{(t-s)A}f(s)ds$$

Remark 9.1. Note that $t - s$ is always positive. Although not particularly relevant to this example, this is an important note in other contexts.

Remark 9.2. What we mean by integrating a vector-valued function is to integrate each of the component functions and 'stack' them together to get another vector.

10 Linearisation

The idea with linearisation is to try and use the framework we've built up thus far to make statements about non-linear systems. Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is a smooth vector field and we have the differential equation

$$x' = f(x) \quad (10.1)$$

Although in general it's hard to give an explicit solution to this, we can still try and determine it's qualitative behaviour, at least locally. Let $p \in \mathbb{R}^n$ be arbitrary. If $f(p)$ is non-zero, then for q near p we are going to have that $f(q)$ is close to $f(p)$. Therefore the flow is going to look like (almost) parallel lines in this neighbourhood. If $f(p) = 0$, we need to do a bit more work. One thing we can do is consider the Taylor expansion of f . We know that

$$f(x) = f(p) + f'(p)(x - p) + \underbrace{O(|x - p|)}_{\text{error}}$$

Since we assume $f(p) = 0$, the first term disappears and by substituting $y = x - p$, (10.1) becomes

$$y' = Df(p)y + O(|y|)$$

Since $Df(p)$ is a linear map, this looks just like the linear equations we have studied thus far, except there is the added error term. The hope is that this error terms is going to be small so by studying the linear system

$$y' = Df(p)y$$

we can get a pretty good idea of how the true system behaves. For example if there are sources or sinks in the linearised system, they will also appear in the true system. The phase portraits will also be similar (similarity will be defined more precisely later) if the origin is hyperbolic (recall this means that $\Re(\lambda) \neq 0$ for every eigenvalue λ of $Df(p)$).

This determines a procedure that we can use to study non-linear equations

1. Find the steady states/equilibria (the points where $x' = f(x)$ is 0). These are points where the solutions are constant and don't change over time.
2. Linearise near these steady states and determine qualitative behaviour in small neighbourhoods around them
3. Tie everything up into a big picture

The cryptic 'tie everything up into a big picture' can be best illustrated with an example.

10.1 Example

We will consider the case of swinging a pendulum, kind of. More precisely we will look at the case when a mass is attached to a (rigid) rod allowed to swing freely in a vertical circle (we use a rod instead of a string because we don't want to worry about cases where the string may fold onto itself or something). See [Figure 14](#).




Figure 14: Mass pendulum setup

The equation modelling this situation is given by

$$mx'' + rx' = -c \sin x$$

where x is the angle made by the rod with the vertical. Here m is the mass of the mass (names are difficult for physicists), r is the constant of friction (therefore $r \geq 0$) and c is some arbitrary constant (the exact details, such as length of the rod, strength of gravity, etc, are used to set c). To make the analysis a bit simpler we will assume m and c to be 1. Then in particular we have the equation

$$x'' + rx' + \sin x = 0 \tag{10.2}$$

We can use our standard trick to convert this second order equation to a system of first order equation: let $v = x'$. Then we have

$$\begin{cases} x' = v \\ v' = -rv - \sin x \end{cases}$$

We first find the the steady states or in other words where v is 0 (since $v = x'$ this naturally means that when $v = 0$, $v' = 0$). These are the points where $\sin x = 0$ or in other words where $x = k\pi$ for $k \in \mathbb{Z}$. Even values of k correspond with the mass hanging on the bottom and k being odd is when the mass is at the top in a perfectly vertical position. Simple intuition tells us that the the former equilibria should be stable (at least if r is positive) and the latter should be unstable (see [Figure 15](#)). Let us see if the equations agree with this. First we see that

$$\begin{pmatrix} x \\ v \end{pmatrix}' = f(x, v) := \begin{pmatrix} v \\ -rv - \sin x \end{pmatrix}$$

Then

$$Df(k\pi, 0) = \begin{pmatrix} 0 & 1 \\ -\cos x & -r \end{pmatrix} \Big|_{x=k\pi}$$

$$= \begin{cases} \begin{pmatrix} 0 & 1 \\ 1 & -r \end{pmatrix}, & k \text{ odd} \\ \begin{pmatrix} 0 & 1 \\ -1 & -r \end{pmatrix}, & k \text{ even} \end{cases}$$

Let us consider the case with k odd first. In this case the determinant of the matrix is -1 so the eigenvalues are real and of opposite sign. We know this will result in a saddle and hence be unstable. In fact around these points we expect the phase portrait to (roughly) have a saddle as well. This lines up with our intuition above.

Now let us consider the case with k even. Then we know the eigenvalues are

$$\lambda_{1,2} = \frac{-r \pm \sqrt{r^2 - 4}}{2}$$

If $0 < r < 2$ then we will have complex eigenvalues implying that we will have a spiral. Since the real part is negative, solutions are going to spiral in and by looking at the first column we can even infer that spiral is going to be clockwise. This corresponds with the angle tending towards 0 and its speed decreasing, as we would expect the pendulum to behave. If $r > 2$ then we get 2 real eigenvalues, both of which are negative. In this case all equilibria will be stable. This corresponds with the friction becoming so strong that the pendulum can actually become stuck and hence 'stable' even when upright.




Figure 15: Equilibria in mass pendulum setup

We get some rather interesting behaviour at $r = 0$, in the frictionless case. In this case we get a center which again should make sense. If there is no friction, then the pendulum continues swinging on its path ad infinitum.

11 Dynamical Systems

Consider our go-to differential equation:

$$X' = AX$$

where A is some fixed matrix. We know the general solution is given by

$$X(t) = e^{tA}v$$

where v is determined by the initial conditions. There are two ways we may wish to study this equation. We could fix v and consider how the equation behaves as we vary t . This gives us exactly the trajectory or orbit of a solution.

On the other hand, we could fix t and consider what happens as we vary v . Such a map, often denoted Φ_t , is known as the *flow* of the system (this function maps some initial point to where it would be sent by the solution at time t). What we get then is a map from v to $e^{tA}v$, which in this case is simply a linear isomorphism on \mathbb{R}^n . We in fact have a family of isomorphisms $\{e^{tA}\}_{t \in \mathbb{R}}$ that satisfy the following properties: $e^{0A} = \text{Id}$ and $e^{(t+s)A} = e^{tA}e^{sA}$. A family of bijective maps $\{\phi_t\}_{t \in \mathbb{R}}$ on \mathbb{R}^n that satisfies these properties (i.e. $\phi_0 = \text{Id}$ and $\phi_{s+t} = \phi_s \circ \phi_t$) is called a dynamical system. In different areas of math, we make different assumptions on the properties that the ϕ_t must satisfy. For our purposes, we generally assume them all to be smooth. In principle, instead of indexing over all of \mathbb{R} , we could index simply over the positive reals or the natural numbers (which would be like looking at discrete time). If we were feeling particularly adventurous, we could replace \mathbb{R}^n with a manifold instead. But for now, we leave things as they are.

It is perhaps useful to consider why the ‘semigroup property’ of dynamical systems (i.e. that $\phi_{s+t} = \phi(s) \circ \phi(t)$) is important. The claim is that with this property, by only looking at the initial conditions, we can in some way study all solutions. This is because the semigroup property allows us to ‘stitch together’ solutions in a certain sense. Suppose we have a solution that begins at x and passes through y_1 in time t . Then suppose we have another solution that starts at y_1 and reaches y_2 in time s . Then one would hope that the original solution that began at x also reaches y_2 , but at time $s + t$ (and in fact hopefully the path from x to y_2 would be given by combining/stitching together the two solutions in the obvious manner, should this be possible). This is exactly the property characterised by the semigroup property.




Figure 16: The semigroup property of flow. Evolving for time s from $\Phi_t(x)$ corresponds exactly with evolving for time $s + t$ from x .

12 Existence and Uniqueness of solutions

The existence and uniqueness of solutions is of course a key point with differential equations. We have already said that most differential equations can't be solved explicitly. Part of the problem is that there may be no analytic solution. However, the situation may be even more dire than that: there may genuinely be no solution. Suppose we are given that

$$x' = \begin{cases} -1, & x \geq 0 \\ 1, & \text{otherwise} \end{cases}$$

Consider the situation at $x = 0$. At this point, the point ‘wants’ to decrease due to its negative gradient. However as soon as it does so, the derivative becomes positive causing it to move up. Hence why the above differential equation has no solutions. We can argue more precisely using a statement in analysis which tells us that the derivative of a function can never have a jump discontinuity (roughly speaking a jump discontinuity means that the left and right hand limits of a function differ, however the derivative of

a function at a point is a limit so we know if it exists, the left and right hand limits are the same). Hence we know that if a solution to $x' = f(x)$ is to exist then f must at the very least be continuous.

Sometimes solutions may exist but need not be unique. Suppose we have

$$x' = 3x^{2/3}$$

with initial condition $x_0 = 0$. Then $x(t) = 0$ and $x(t) = t^3$ are both solutions that satisfy this initial value problem. The problem here is that f is continuous but not smooth (or even differentiable everywhere). In fact what we will see is that given $x' = f(x)$, we are guaranteed to have unique solutions to the ODE, if f is continuously differentiable. Even then, we cannot be sure that we have a solution that is defined everywhere as is illustrated by equation

$$x' = x^2$$

whose general solution is given by

$$x(t) = \frac{x_0}{1 - x_0 t}$$

We see that this is not defined at $t = x_0^{-1}$. This leads us to the local uniqueness and existence theorem.

Theorem 12.1 (Local Existence and Uniqueness) *Suppose we are considering the ordinary differential equation $X' = F(X)$ where $F : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is continuously differentiable. Then for all $v \in \mathbb{R}^n$ there exists an interval $I = (\alpha, \beta)$ containing 0 such that $X' = F(X)$ has a unique solution $X : I \rightarrow \mathbb{R}^n$ satisfying $X(0) = v$.*

Although we will get on to proving this theorem soon enough, let us consider some of its consequences. For example, if we have two solutions where the corresponding intervals intersect, is it the case that the solutions agree on the intersection? The answer is yes. This follows from uniqueness of the solutions: take a point in the intersection and use this to define our initial condition. We know that both solutions solve the differential equation, therefore must be equal (on the intersection). This suggests that we can ‘patch’ together local solutions (as discussed previously) to get solutions defined on a larger interval. By pasting together all the local solutions, we can get a maximal interval of existence.

12.1 Preliminary Theory

There are several results we need before we can prove the local existence and uniqueness theorem. On account of the author being too small brain, we largely only include the definitions and statements of the theorems we need and omit the proofs.

Because we will refer to it many times, we also include the definition of uniform continuity and uniform convergence. Throughout, let E be a subset of \mathbb{R}^n .

Definition 12.2 (Uniform Continuity). A function $f : E \rightarrow \mathbb{R}^m$ is said to be uniformly continuous if for every $\epsilon > 0$ we can find a $\delta > 0$ such that given $x, y \in E$ satisfying $|x - y| < \delta$, we have $|f(x) - f(y)| < \epsilon$ (in particular the choice of δ is independent of x and y).

Definition 12.3 (Uniform convergence). A sequence of functions $\{f_n : E \rightarrow \mathbb{R}^m\}$ is said to converge uniformly to a function $f : E \rightarrow \mathbb{R}^m$ if for every $\epsilon > 0$ there exists some $N \in \mathbb{N}$ such that

$$\sup_{x \in E} |f_n(x) - f(x)| < \epsilon$$

for every $n \geq N$.

Recall that if a sequence of continuous functions converges uniformly then the function they converge to is also continuous.

Definition 12.4 (Lipschitz continuous). A function $f : E \rightarrow \mathbb{R}^m$ is said to be Lipschitz continuous if there exists a constant L such that

$$|f(x) - f(y)| \leq L|x - y|$$

for all $x, y \in E$. L is sometimes called a Lipschitz constant.

Definition 12.5 (Uniformly Bounded). Let $\{f_\alpha : E \rightarrow \mathbb{R}^m\}$ be a family of functions. Suppose there exists some M such that $|f_\alpha(x)| < M$ for all $x \in E$ and all α . Then we say that the $\{f_\alpha\}$ are uniformly bounded (by M).

Definition 12.6 (Equicontinuous). Let $\{f_\alpha\}$ be a family of functions where each f_α is a map from $E \subset \mathbb{R}^n$ to \mathbb{R}^m . The $\{f_\alpha\}$ are said to be equicontinuous if for every $\epsilon > 0$ we can find some $\delta > 0$ such that for all $x, y \in E$ and for all f_α , if we have $|x - y| < \delta$ then $|f_\alpha(x) - f_\alpha(y)| < \epsilon$ (in particular the choice of δ is independent of x, y and α).

Example 12.1. The family of functions $f_\alpha(x) = \alpha x$ for $\alpha \in \mathbb{R}$ is not equicontinuous. If we instead restrict α to be in $[3, 5]$ (or any bounded interval) then the family of functions is equicontinuous. ■

Example 12.2. The family of functions $f_n(x) = x^n$ defined on the unit interval $[0, 1]$ is not equicontinuous. ■

Example 12.3. A family of Lipschitz functions which share the same Lipschitz constant is equicontinuous. ■

Example 12.4. Any finite family of uniformly continuous functions is equicontinuous. ■

We can define the sup norm on a function space where if $f : E \rightarrow \mathbb{R}^m$ then

$$\|f\|_\infty = \sup_{x \in E} |f(x)|$$

The space of continuous functions on $[0, 1]$ is often denoted $C([0, 1])$. We claim that that $C([0, 1])$ equipped with $\|\cdot\|_\infty$ is a complete metric space. This is exactly the statement that the uniform convergence of a sequence of continuous functions is also continuous.

Theorem 12.7 Suppose $f_n : E \rightarrow \mathbb{R}^n$ is a sequence of continuous maps. If E is compact and the $\{f_n\}$ converge uniformly, then the $\{f_n\}$ are uniformly bounded and equicontinuous.

Proof. Let f denote the limit of the f_n . We first show uniform boundedness. Since each f_n is continuous and E is compact, we know there exists some constant M_n such that $|f_n(x)| \leq M_n$ (indeed we can take M_n to be $\|f_n\|_\infty$). Since f is also continuous, it is also going to be bounded by some M_0 .

The idea is that eventually all the f_n are going to be quite close to f therefore we should be able to bound (bind?) all but finitely many of them with M_0 (or something close to it at least). Then we are only left with finitely many bounds so the maximum of all of these should bound all the f_n . Let us formalise this.

Let N be such that for all $n \geq N$ we have that $\|f - f_n\|_\infty < 1$ (equivalently, $|f(x) - f_n(x)| < 1$ for all $x \in E$). The existence of such an N is guaranteed by uniform convergence. Let $M = \max\{M_0 + 1, M_1, \dots, M_N\}$. We claim that M is a bound for all the f_n . Clearly this holds true for $n \leq N$. Suppose $n > N$. Then

$$\|f_n\|_\infty \leq \|f_n - f\|_\infty + \|f\|_\infty < 1 + M_0$$

Thus the f_n are uniformly bounded.

For equicontinuity, we again use the fact that we can use f to approximate all but finitely many of the f_n . In particular, we see that

$$|f_n(x) - f_n(y)| \leq |f_n(x) - f(x)| + |f(x) - f(y)| + |f_n(y) - f(y)|$$

Suppose $\epsilon > 0$ is given. Then there exists some $N \in \mathbb{N}$ such that $\|f - f_n\|_\infty < \frac{\epsilon}{3}$ for $n \geq N$. Additionally since f is continuous on a compact space, it is in particular uniformly continuous. Thus there exists a δ_0 such that if $|x - y| < \delta_0$ then $|f(x) - f(y)| < \frac{\epsilon}{3}$. Furthermore, each of the f_n are uniformly continuous as well, thus there exist similar δ_n for each of them as well. Then we can take $\delta = \min\{\delta_0, \delta_1, \dots, \delta_N\}$. □

We then get to the Ascoli-Arzelà theorem. Roughly speaking, Ascoli-Arzelà gives a complete description of compact subsets of function spaces (with some conditions) analogous to how Heine-Borel gives a complete description of compact subsets of \mathbb{R}^n .

Theorem 12.8 (Ascoli-Arzelà) *Let $\mathcal{F} = \{f_\alpha : E \rightarrow \mathbb{R}^m\}$ where E is a compact subset of \mathbb{R}^n be an infinite family of continuous functions that is uniformly bounded and equicontinuous. Then there exists a sequence of functions in \mathcal{F} that converges uniformly on E .*

Proof. You can see my notes [here](#) or open any functional analysis textbook (I think). \square

Recall that we say a set A is relatively compact if its closure is compact which (for metric spaces) is the same as saying every sequence in A has a convergent subsequence (although the limit may not be in A itself).⁵ We then have the following theorem

Theorem 12.9 *Let $A \subset C(K, \mathbb{R}^m)$ (recall this is the set of continuous functions from K to \mathbb{R}^m) where $K \subset \mathbb{R}^n$ is compact. Then A is relatively compact if and only if it is uniformly bounded and equicontinuous.*

Proof. Once again, my notes above or any functional analysis textbook should work. \square

We also need consider how we may extend functions defined on a subset to be defined on the whole space.

Theorem 12.10 *Let $f : \overline{B(0, r)} \rightarrow \mathbb{R}^m$ (where $B(0, r) \subset \mathbb{R}^n$) be continuous. Then $\bar{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$, where*

$$\bar{f}(x) = \begin{cases} f(x) & \text{if } |x| \leq r \\ f\left(r \frac{x}{|x|}\right) & \text{if } |x| > r \end{cases}$$

is continuous as well.

Proof. Trust me bro. \square

Finally there are various fixed point theorems we should be aware of, none of which we prove.

Theorem 12.11 (Banach's Contraction Mapping Theorem) *Let E be a complete metric space and let $T : E \rightarrow E$ be such that there exists some $0 \leq q < 1$ where $d(T(x), T(y)) \leq qd(x, y)$. Then T has a unique fixed point. In other words there is exactly one point $z \in E$ such that $T(z) = z$.*

Remark 12.12. The proof of this theorem is constructive. In fact the construction is such that you can start with any point in E and construct a sequence (by iteratively applying T) that converges to the fixed point.

Theorem 12.13 (Brouwer's Fixed Point Theorem) *Let $B \subset \mathbb{R}^n$ denote the closed unit ball. Then if $T : B \rightarrow B$ is continuous, it has a fixed point.*

Remark 12.14. The proof of this theorem is a bit less nice unfortunately. Although we know a fixed point exists, in general, we have no way of working out what it might be.

Theorem 12.15 (Schauder-Tychonoff Theorem) *Let B be the (closed) unit ball in $C([0, 1], \mathbb{R}^n)$, equipped with the usual supremum norm (in other words this is the set of continuous functions on $[0, 1]$ that take values in the unit ball in \mathbb{R}^n). Suppose $T : B \rightarrow B$ is continuous map where $T(B)$ is relatively compact. Then T has a fixed point.*

Remark 12.16. We know that $T(B) \subset B$. Since B is bounded, $T(B)$ is always going to be bounded. Thus, by [Theorem 12.9](#), instead of assuming $T(B)$ is relatively compact, we can equivalently assume $T(B)$ to be an equicontinuous family of functions.

⁵For a similar statement in general topological spaces, change the word sequence to net, see [Wikipedia](#)

12.2 Existence of solutions

Given an initial value problem (IVP), there are 3 things we would like to have, the existence of a solution, the uniqueness of a solution and a continuous dependence on its parameters. We begin by showing the first of these. To be precise we want the following.

Suppose $\xi_0 \in \mathbb{R}^n$ and $f : [t_0, t_1] \times \mathbb{R}^n \rightarrow \mathbb{R}^n$ is continuous. We wish to find a function $x : [t_0, t_1] \rightarrow \mathbb{R}^n \in C^1([t_0, t_1]; \mathbb{R}^n)$ such that

$$x'(t) = f(t, x(t))$$

and $x(t_0) = \xi_0$.

We call this the initial value problem of course. There is an equivalent way of formulating this problem, in which we need to solve an integral equation. The statement is as follows.

Let f and ξ_0 be as above. We wish to find $x : C^1([t_0, t_1]; \mathbb{R}^n)$ such that for all $t \in [t_0, t_1]$, we have

$$x(t) = \xi_0 + \int_{t_0}^t f(s, x(s)) ds$$

Using the fundamental theorem of calculus, we can see that the two problems are equivalent. (In other words, if we find an x that solves the integral equation, then it will be a solution to the initial value problem and vice versa). In general, we will work to solve the integral equation as integrating things makes them nicer (discontinuous functions become continuous, continuous functions become differentiable, etc.).

Theorem 12.17 (Cauchy-Peano Theorem) *Suppose $f : [t_0, t_1] \times \mathbb{R}^n \rightarrow \mathbb{R}^n$ is continuous and bounded by M . Suppose we are given the initial value problem*

$$\begin{cases} x'(t) = f(t, x(t)) \text{ for all } t \in [t_0, t_1] \\ x(t_0) = \xi_0 \end{cases}$$

Then the IVP has at least one solution.

Proof. As mentioned previously, we will work to solve the integral equation

$$x(t) = \xi_0 + \int_{t_0}^t f(s, x(s)) ds$$

We first define an operator $T : C^1([t_0, t_1]; \mathbb{R}^n) \rightarrow C^1([t_0, t_1]; \mathbb{R}^n)$ where

$$Ty(t) = \xi_0 + \int_{t_0}^t f(s, y(s)) ds$$

What we wish to find, then, is some x such that $Tx = x$. As one can imagine, we will use the [Schauder-Tychonoff](#) theorem to do this.

In order to apply the theorem, we need to verify that its conditions hold. So let B be the unit ball in $C^1([t_0, t_1]; \mathbb{R}^n)$ (in other words $B = \{y \in C^1([t_0, t_1]; \mathbb{R}^n) : |y(t)| \leq 1\}$; we use C^1 because we want the solution to be C^1). We see that

$$\begin{aligned} |Tx(t)| &\leq |\xi_0| + \left| \int_{t_0}^t f(s, x(s)) ds \right| \\ &\leq |\xi_0| + \int_{t_0}^t |f(s, x(s))| ds \\ &\leq |\xi_0| + M(t - t_0) \end{aligned}$$

We split the remainder of the proof into two cases. First suppose $|\xi_0| + M(t_1 - t_0) \leq 1$. Then clearly $Tx \in B$ implying that we do indeed have a map $T : B \rightarrow B$. Now suppose $t, t' \in [t_0, t_1]$. Then

$$\begin{aligned} |Tx(t) - Tx(t')| &= \left| \left(\xi_0 + \int_{t_0}^t f(s, x(s)) ds \right) - \left(\xi_0 + \int_{t_0}^{t'} f(s, x(s)) ds \right) \right| \\ &= \left| \int_{t'}^t f(s, x(s)) ds \right| \\ &\leq M|t - t'| \end{aligned}$$

This means that $T(B)$ is a family of functions which all share the Lipschitz constant M . Therefore $T(B)$ is equicontinuous. Since we already know it to be bounded (we are assuming that $T(B) \subset B$), we have that $T(B)$ is relatively compact by [Theorem 12.9](#). All that remains to show is that T is continuous. Since we are working in metric spaces, it suffices to show that T maps convergent sequences to convergent sequences. So suppose $\{x_k\}$ is a sequence of functions in B that converge uniformly to some function x . In other words, we have that $\|x - x_k\|_\infty \rightarrow 0$ and we will show that $\|Tx - Tx_k\| \rightarrow 0$ as well. First we note

$$|Tx(t) - Tx_k(t)| \leq \int_{t_0}^t |f(s, x(s)) - f(s, x_k(s))| ds$$

By uniform convergence, the right hand side goes to 0. This is because f is uniformly continuous (this follows from f being continuous on the compact set $[t_0, t_1] \times [-1, 1]^n$. Since we are only concerned with the set B , we only need consider f restricted to this domain) thus we can get the integrand to be as small as we like given that the inputs are close enough and the inputs can be made as close as we like by uniform convergence.

To be precise, suppose $\epsilon > 0$. By uniform continuity of f , we know there is some $\delta > 0$ such that $|(s_1, x_1) - (s_2, x_2)| < \delta$ implies that $|f(s_1, x_1) - f(s_2, x_2)| < \frac{\epsilon}{t_1 - t_0}$. By uniform convergence of x_k , there is some $N \in \mathbb{N}$ such that for $k \geq N$, we have $\|x - x_k\|_\infty < \delta$. Then $|x(t) - x_k(t)| < \delta$ for all t . Thus for $k \geq N$, we have

$$\int_{t_0}^t |f(s, x(s)) - f(s, x_k(s))| ds < \frac{\epsilon}{t_1 - t_0} (t_1 - t_0) = \epsilon$$

With this we satisfy all the conditions of Schauder-Tychonoff which tells us there exists some $x \in B$ such that $Tx = x$, the precise statement we were aiming for.

The second case to consider is when $|\xi_0| + M(t_1 - t_0) > 1$. Let us call the quantity on the left H . Then we define

$$g(t, x) = \frac{1}{H} f(t, Hx), \eta_0 = \frac{1}{H} \xi_0$$

This gets us a new IVP which falls under case 1 and gets us a solution $y(t)$. Then $x(t) = Hy(t)$ solves the original IVP. \square

There is a second proof of the Cauchy-Peano theorem that is inspired by [Euler's polygonal](#).

Proof. Just to make notation easier, suppose we take $t_0 = 0$ and $t_1 = 1$. Let k be some natural number. Then we define

$$x_k(t) = \begin{cases} \xi_0 & \text{if } 0 \leq t \leq \frac{1}{k} \\ \xi_0 + \int_0^{t-\frac{1}{k}} f(s, x_k(s)) ds & \text{otherwise} \end{cases}$$

Note that $x'_k(t) = f(t - \frac{1}{k}, x_k(t - \frac{1}{k}))$ which should be close $f(t, x_k(t))$ provided that k is sufficiently large. Hence the hope is that we can find a converging (sub)sequence among these x_k that will satisfy the IVP. We show that a converging subsequence exists by using [Ascoli-Arzelà](#). Thus we will show that $(x_k)_{k \in \mathbb{N}}$ is family of uniformly bounded, equicontinuous functions.

Since M is a bound for f , we see that

$$|x_k(t)| \leq |\xi_0| + \int_0^{t-\frac{1}{k}} |f(s, x_k(s))| ds \leq |\xi_0| + M$$

Therefore $|\xi_0| + M$ is a uniform bound for the x_k . Additionally

$$\begin{aligned} |x_k(t) - x_k(t')| &= \left| \int_{t'-\frac{1}{k}}^{t-\frac{1}{k}} f(s, x_k(s)) ds \right| \\ &\leq M \left| \left(t - \frac{1}{k} \right) - \left(t' - \frac{1}{k} \right) \right| \\ &= M |t - t'| \end{aligned}$$

This means that all the x_k are Lipschitz with Lipschitz constant M and hence form an equicontinuous family of functions. By Ascoli-Arzelà, we know there exists a convergence subsequence which we will denote x_k again and we denote their limit as x . All that remains to show is that x satisfies the integral equation (this ends up being easier to show than proving it satisfies the IVP). We see that

$$x_k(t) = \xi_0 + \int_0^t f(s, x_k(s)) ds - \int_{t-\frac{1}{k}}^t f(s, x_k(s)) ds$$

Consider what happens as we take the limit as $k \rightarrow \infty$. The left hand side goes to $x(t)$ by definition and the final term on the right goes to 0 (it bounded by $\frac{M}{k}$). Since the convergence of the x_k is uniform, taking the limit commutes with integration⁶ and the continuity of f allows us to move the limit further in, ultimately allowing us to conclude that

$$x(t) = \xi_0 + \int_0^t f(s, x(s)) ds$$

□

Consider again the initial value problem

$$\begin{aligned} x'(t) &= f(t, x(t)) \\ x(t_0) &= \xi_0 \end{aligned}$$

We will try find an interval $I = [t_0, t_0 + h]$ (for $h > 0$) and a solution $x \in C^1(I; \mathbb{R}^n)$ such that x is a solution to the IVP. In this case, we will say that x is a solution to IVP₊ since it only works in the positive time direction. For now we only focus on moving time forward, the theory for working backwards in time is near identical and will be discussed later.

If x is a solution to IVP₊ on I , then we will say it can be continued to the right if there exists a pair (\bar{x}, \bar{I}) such that $\bar{I} \supset I$ and $\bar{x}|_I = x$. We will say a continuation is strict if $\bar{I} \supsetneq I$. This allows us to define a partial order⁷, namely

$$(x_1, I_1) \geq (x_2, I_2) \Leftrightarrow I_1 \supset I_2 \text{ and } x_1|_{I_2} = x_2$$

By Zorn's lemma, there exists a maximal continuation (x^*, I^*) of (x, I) . Note that I^* is the union of all intervals on which we have continuations of x .

Theorem 12.18 Suppose we are given $(t_0, \xi_0) \in \mathbb{R} \times \mathbb{R}^n$. Let $A = [t_0 - h, t_0 + h] \times \overline{B(\xi_0, a)}$ for some $h, a > 0$. Let $f : A \rightarrow \mathbb{R}^n$ be continuous function bounded by M . Then the IVP

$$\begin{aligned} x'(t) &= f(t, x(t)) \\ x(t_0) &= \xi_0 \end{aligned}$$

has at least one solution x defined on

$$I = \left[t_0 - \min\left\{h, \frac{a}{M}\right\}, t_0 + \min\left\{h, \frac{a}{M}\right\} \right]$$

Moreover, any solution to the IVP defined on $J \subset I$, where J is a neighbourhood of t_0 (i.e. J contains an open set that in turn contains t_0), can be continued to a solution on I .

Remark 12.19. Although the statement is similar to Cauchy-Peano, note that f is now local in space (we use a closed ball rather than all of \mathbb{R}^n). Hence we will require a further bit of argument.

Proof. By Theorem 12.10, we can extend f to \bar{f} defined on $[t_0 - h, t_0 + h] \times \mathbb{R}^n$. By definition of the extension, $\|\bar{f}\|_\infty \leq M$ (which is to say its values are contained in the cube of 'radius' M). By Cauchy-Peano, there exists a solution \bar{x} to the IVP

$$\begin{aligned} x'(t) &= \bar{f}(t, x(t)) \\ x(t_0) &= \xi_0 \end{aligned}$$

⁶See Wikipedia

⁷See Wikipedia

Unfortunately \bar{x} is not necessarily a solution to our original IVP since it may take values outside of $\overline{B(\xi_0, a)}$. However this can be easily fixed by using the fact that \bar{x} is continuous. Since \bar{x} is continuous, there exists $j > 0$ (where $j \leq h$) such that if $t \in [t_0, t_0 + j]$ then $|\bar{x}(t) - \bar{x}(t_0)| = |\bar{x}(t) - \xi_0| \leq a$. In other words, since the solution starts in the ball centered at ξ_0 , it must stay in the ball for some amount of time.

Hence we have a solution to the given IVP on $I = [t_0, t_0 + j]$ for some j . Now we want to show that given a solution on any $J \subset I$, the solution can be extended to a solution on I (where of course J must be a neighbourhood of t_0).

Let (x, J) be a solution to the IVP where $J = [t_0, t_1]$ for some $t_1 \in [t_0, t_0 + h]$ (as mentioned, we will only consider the case of moving forward in time for now). We extend it as described above to get a maximal solution (i.e. defined on the maximal time interval). Then we need show that if (x^*, J^*) is this maximal solution, then $J^* \supset I$.

First we claim that J^* contains its right endpoint (in other words J^* is a closed interval). Let $t', t'' \in J^*$ with $t' < t''$. Then, by considering the integral equation, we get that

$$\begin{aligned} |x^*(t'') - x^*(t')| &\leq \left| \int_{t'}^{t''} f(s, x^*(s)) ds \right| \\ &\leq \int_{t'}^{t''} |f(s, x^*(s))| ds \\ &\leq M(t'' - t') \end{aligned}$$

This means that x^* is uniformly continuous (indeed it is even Lipschitz). In particular then x^* maps a Cauchy sequence to a Cauchy sequence. By constructing a Cauchy sequence that converges to $t_1 := \sup(J^*)$, we can either define x^* on it or verify that it is continuous at that point. In principle we also need to check that $x^*(t_1) \in \overline{B(\xi_0, a)}$. However this is clear since the set is closed (and even compact), so it contains its limit points (and by the construction given, we know that $x^*(t_1)$ is indeed a limit point).

Remark 12.20. This means that any solutions (at least the ones obtained via this method) are defined on closed intervals rather than open/half-open intervals.

Then we show that $t_1 \geq t_0 + \min\{h, \frac{a}{M}\}$. If $t_1 = t_0 + h$, we are done, so suppose $t_1 < t_0 + h$. We will show that $t_1 \geq t_0 + \frac{a}{M}$ in this case. First we note that we must have $x^*(t_1)$ on the boundary of $\overline{B(\xi_0, a)}$ (if $x^*(t_1)$ was in the interior, we could run the proof again with $x^*(t_1)$ as our initial starting point and obtain a solution on $[t_1, t_1 + b]$ for some $b > 0$. We could then combine the solutions on $[t_0, t_1]$ and $[t_1, t_1 + b]$ to get a solution on $[t_0, t_1 + b]$, contradicting J^* being the maximal interval). This means that

$$\begin{aligned} a &= |x^*(t_1) - \xi_0| \\ &= |x^*(t_1) - x^*(t_0)| \\ &\leq (t_1 - t_0) \|x'\| \\ &\leq M(t_1 - t_0) \end{aligned}$$

The third line follows from the mean value theorem in higher dimensions⁸. Therefore we get that

$$t_1 \geq \frac{a}{M} + t_0$$

All that remains to show then is that we can continue to the left of t_0 in a similar manner. Of course we could simply run through the proof again changing the signs where necessary to obtain solutions to the left of t_0 . However, if we are clever (and mathematicians are nothing if not clever), we can use the theory built up thus far to find the solution for negative time.

We define $g : [t_0 - h, t_0 + h] \times \overline{B(\xi_0, a)} \rightarrow \mathbb{R}^n$ where $g(t, \xi) = -f(2t_0 - t, \xi)$ (notice how as t increases from $t_0 - h$ to $t_0 + h$, the first input to f decreases from $t_0 + h$ to $t_0 - h$). We then find a solution on some $J' = [t_0, t_0 + b]$ (where b is the minimum value asserted in the statement of the theorem) to the IVP

$$\begin{aligned} y'(t) &= g(t, y(t)) \\ y(t_0) &= \xi_0 \end{aligned}$$

⁸See [Wikipedia](#)

By setting $x(t) = y(2t_0 - t)$ we obtain a solution on $[t_0 - b, t_0]$ as desired. □

Now that we shown a few existence theorems, we should give uniqueness a shot as well. We will in fact give 2 proofs, related in essence but different in flavour.

Theorem 12.21 *Let D be an open subset of $\mathbb{R} \times \mathbb{R}^n$ with $(x_0, y_0) \in D$. Let $f : D \rightarrow \mathbb{R}^n$ be a function that is continuous in x and Lipschitz in y , with Lipschitz constant K . Then there exists some $a > 0$ such that the initial value problem*

$$\begin{aligned} y'(x) &= f(x, y(x)) \\ y(0) &= x_0 \end{aligned}$$

has a unique solutions on $(x_0 - a, x_0 + a)$.

Contraction Mapping. For the first proof, we will use [Banach's contraction mapping theorem](#).

First we choose a rectangle $R' = [x_0 - A, x_0 + A] \times [y_0 - L, y_0 + L]$ (if $y_0 \in \mathbb{R}^n$ for $n > 1$, we do this for every coordinate) that is contained in D . Since f is continuous and R' is compact, f is bounded by some constant M . Let a be a positive real number that is less than $\min\{\frac{L}{M}, A, \frac{1}{K}\}$. Let $R = [x_0 - a, x_0 + a] \times [y_0 - L, y_0 + L]$.

Now we define a space $X := \{y \in C([x_0 - a, x_0 + a]; \mathbb{R}^n) : \|y - y_0\|_\infty \leq L\}$ (where we think of $y_0 \in C([x_0 - a, x_0 + a]; \mathbb{R}^n)$ as the function that is constantly y_0). Note that if $y \in X$, then the graph of y is contained in R' and since $R' \subset D$, in particular the graph is contained in D . Thus $f(x, y(x))$ makes sense for every $y \in X$.

We would like to apply the Banach contraction mapping theorem to X . We know X is complete as it is the closed subset of a complete metric space. Then we need a map $\Gamma : X \rightarrow X$ that is a contraction from which we will get a fixed point. Remembering the integral equation, what we want is to find a $y \in C([x_0 - a, x_0 + a]; \mathbb{R}^n)$ such that

$$y(x) = y_0 + \int_{x_0}^x f(s, y(s)) ds$$

This tells us how to define Γ . In particular,

$$\Gamma(y)(x) = y_0 + \int_{x_0}^x f(s, y(s)) ds$$

If we can find a fixed point of Γ , we will have found a solution to our initial value problem. In order to do so, all we need do is verify that the hypotheses of the Banach contraction mapping theorem hold.

First we want to know that Γ does indeed map X into itself (a priori, it is not obvious that the image of $y \in X$ under Γ is in X). Luckily for us this indeed the case since

$$\begin{aligned} |\Gamma(y)(x) - y_0| &= \left| \int_{x_0}^x f(s, y(s)) ds \right| \\ &\leq \int_{x_0}^x |f(s, y(s))| ds \\ &\leq M|x - x_0| \\ &\leq Ma \\ &\leq L \end{aligned}$$

As this holds for all x , we get that

$$\|\Gamma(y) - y_0\|_\infty \leq L$$

The only thing that remains to be shown then, is that Γ is a contraction. What we will show is that

$\|\Gamma(y) - \Gamma(z)\|_\infty \leq aK \|y - z\|_\infty$. By our choice of a , we know that $aK < 1$ and so the claim will be satisfied.

$$\begin{aligned} |\Gamma(y)(x) - \Gamma(z)(x)| &\leq \int_{x_0}^x |f(s, y(s)) - f(s, z(s))| ds \\ &\leq \int_{x_0}^x K |y(s) - z(s)| ds \\ &\leq K \int_{x_0}^x \|y - z\|_\infty \\ &= K \|y - z\|_\infty (x - x_0) \\ &\leq aK \|y - z\|_\infty \end{aligned}$$

The second line follows from the fact that f is Lipschitz in the second component.

Thus we can safely apply the contraction mapping theorem and conclude the proof. Since the theorem tells us that the fixed point is unique, we know that our solution y is unique. \square

Remark 12.22. Suppose y_1 is any continuous map. Then functions $\Gamma(y_1)$, $\Gamma^2(y_1) = \Gamma(\Gamma(y_1))$, $\Gamma^3(y_1) = \Gamma(\Gamma(\Gamma(y_1)))$, ... are known as the Picard iterates.

Now we present a second proof of the same theorem using Picard iterates. Arguably this is the same thing, since if we were to follow through with the contraction mapping theorem, we would be applying the contraction, i.e. Γ , repeatedly to our initial point which is exactly what the Picard iterates are. However it's still useful to consider different proofs to see the different techniques they employ. Additionally, our solution will be defined on a slightly larger interval which is also nice.

Proof. Let R' be as before (in other words $R' = [x_0 - A, x_0 + A] \times [y_0 - L, y_0 + L]$ such that $R' \subset D$). This time we choose our a to be less than $\min\{\frac{L}{M}, A\}$ (where, as a reminder, M is a bound for f). As before, we will take $R = [x_0 - a, x_0 + a] \times [y_0 - L, y_0 + L]$.

Let $y_1 \equiv y_0$. For $n > 1$, we define $y_n = \Gamma(y_{n-1}) = \Gamma^{n-1}(y_1)$. We claim that the y_n converge uniformly to some $y \in C([x_0 - a, x_0 + a]; \mathbb{R}^n)$ where y is a solution to the integral equation (corresponding to the given IVP).

Perhaps, we should first show that the above makes sense. Which is to say, we need that $(x, y_n(x))$ is always in D so that we can evaluate f on it. But as shown in the previous proof, $\|\Gamma(y) - y_0\|_\infty \leq L$ if $\|y - y_0\|_\infty \leq L$. We know that $\|y_1 - y_0\|_\infty \leq L$ (in fact it's 0!), so its graph is contained in R (which in turn is contained in D). By induction, this holds for all y_n .

Now we get to the actual meat: showing that the y_n converge. First we can immediately compute that

$$|y_2(x) - y_1(x)| \leq M|x - x_0|$$

This a calculation we have done many times previously. Then we see that for $n \geq 2$

$$|y_{n+1}(x) - y_n(x)| \leq \int_{x_0}^x |f(t, y_n(t)) - f(t, y_{n-1}(t))| dt \leq K \int_{x_0}^x |y_n(t) - y_{n-1}(t)| dt$$

By induction then what we get is that

$$|y_{n+1}(x) - y_n(x)| \leq K^{n-1} M \frac{|x - x_0|^n}{n!} \leq K^{n-1} M \frac{a^n}{n!}$$

Since this holds for all $x \in [x_0 - a, x_0 + a]$ what we get is that

$$\|y_{n+1} - y_n\|_\infty \leq \frac{M}{K} \frac{(Ka)^n}{n!}$$

Since the right hand side denotes the terms of a convergent series, we can conclude that the y_n form a Cauchy sequence and hence converge to some y (to be precise, in order to show Cauchy, we should show that $\|y_n - y_m\|_\infty < \epsilon$ eventually, as we let $n, m \rightarrow \infty$. However $\|y_n - y_m\|_\infty \leq \|y_n - y_{n-1}\|_\infty + \|y_{n-1} - y_{n-2}\|_\infty + \dots + \|y_{m+1} - y_m\|_\infty$ (assuming $n > m$). Each of these terms can be made arbitrarily small by the above estimate and so we are done).

The final thing we wish to show is that y does indeed solve the integral equation. We know that

$$y_n(x) = y_0 + \int_{x_0}^x f(t, y_{n-1}(t)) dt$$

We can consider what happens as we let n tend to infinity on both sides. The left hand side goes to $y(x)$ by definition. As for the right hand side, what we see is that

$$\left| \int_{x_0}^x f(t, y_{n-1}(t)) dt - \int_{x_0}^x f(t, y(t)) dt \right| \leq K \int_{x_0}^x |y_{n-1}(t) - y(t)| dt \leq Ka \|y_{n-1} - y\|_{\infty}$$

Therefore the right hand side approaches

$$y_0 + \int_{x_0}^x f(t, y(t)) dt$$

and we are done. □

12.3 Uniqueness of solutions

The reader has the right to complain because the previous theorem guarantees not only the existence of a solution but also its uniqueness. However, this was left unaddressed in the second proof (and arguably swept under the rug of another theorem for the first one!). We remedy that over here. In fact to make up for it, we will provide 2 proofs of uniqueness.

Proposition 12.23 *The solution found using Picard iterates for [Theorem 12.21](#) is unique.*

Proof 1. Suppose z is another solution to the IVP on $[x_0 - a, x_0 + a]$. Then again we have that the graph of z is contained in $R = [x_0 - a, x_0 + a] \times [y_0 - L, y_0 + L]$ (all the quantities are as defined in the proof). This is because

$$\begin{aligned} |z(x) - y_0| &= |z(x) - z(x_0)| \\ &\leq \|z'\| |x - x_0| \\ &\leq Ma \\ &\leq L \end{aligned}$$

where M , as before, is a bound for f and a is chosen to be smaller than $\frac{L}{M}$ ($\|z'\|$ means the operator norm of z').

We will show that $|y(x) - z(x)| = 0$ for every $x \in [x_0 - a, x_0 + a]$. Immediately, computing this quantity, what we find is that

$$|y(x) - z(x)| = \left| \int_{x_0}^x f(t, y(t)) - f(t, z(t)) dt \right| \leq K \int_{x_0}^x |y(t) - z(t)| dt \leq 2LK |x - x_0|$$

Then

$$\begin{aligned} |y(x) - z(x)| &\leq K \int_{x_0}^x |y(t) - z(t)| dt \\ &\leq K \int_{x_0}^x 2LK(t - x_0) dt \\ &\leq \frac{2LK(x - x_0)^2}{2} \end{aligned}$$

By induction what we find is that

$$|y(x) - z(x)| \leq \frac{2LK(x - x_0)^n}{n!}$$

for all n . Clearly the right hand side goes to 0 as $n \rightarrow \infty$ (the series converges to an exponential for example), thus $y(x) = z(x)$. □

Proof 2. For the second proof, we once again reach the same estimate as before

$$|y(x) - z(x)| \leq K \int_{x_0}^x |y(t) - z(t)| dt$$

Now we define $u(x) = |y(x) - z(x)|$ which we will show is 0. Note that by assumption u satisfies the following property

$$u(x) \leq K \int_{x_0}^x u(t) dt$$

This is known as Grönwall's inequality.

Define

$$U(x) = \int_{x_0}^x u(t) dt$$

Then we know that

$$U'(x) = u(x) \leq KU(x)$$

If we had an equality here, we could say that $U(x) = U(x_0)e^{K(x-x_0)}$. Grönwall tells us that the statement remains true if we replace the equalities with inequalities. Intuitively this is because $U(x_0) = U(x_0)e^{K(x_0-x_0)}$, so they 'start' at the same point but the right hand side grows faster (has larger gradient) so must always be greater than the left hand side. For more rigorous reasons, consider the following argument

$$\left(\frac{U(x)}{e^{K(x-x_0)}} \right)' = \frac{U'(x) - KU(x)}{e^{K(x-x_0)}} \leq 0$$

This means that the function is non-increasing so must for $x \geq x_0$, we see that

$$\frac{U(x)}{e^{K(x-x_0)}} \leq \frac{U(x_0)}{e^{K(x_0-x_0)}} \\ U(x) \leq U(x_0)e^{K(x-x_0)}$$

Since $U(x_0) = 0$, we know that $U(x) \leq 0$ for all x (the case for $x < x_0$ is immediate from the definition of U). However, U is also the integral of a non-negative function, so $U(x)$ must always be non-negative as well. This must mean that U is 0 which in turn implies that u is 0. \square

Grönwall's inequality exists in different levels of generality. A slightly more general case (than the one used above) is as follows.

Theorem 12.24 *If $U'(x) \leq K(x)U(x)$ where $K(x)$ is continuous and $U(x)$ is (of course) differentiable, then*

$$U(x) \leq U(x_0) \exp \left(\int_{x_0}^x K(t) dt \right)$$

Proof. Exercise for the student : (\square)

We have seen that under sufficiently good conditions, we can get unique solutions on intervals of the form $[x_0, x_0 + \delta_1]$ for some $\delta_1 > 0$. We can then run the proofs again so that our new solutions are on $[x_0 + \delta_1, x_0 + \delta_1 + \delta_2]$. Since the solutions are unique, we can combine them to get a solution on $[x_0, x_0 + \delta_1 + \delta_2]$. Thus we obtain a sequence of δ_i . There are then 2 possibilities: either the δ_i form a convergent series or they don't. Since they are all positive, if the δ_i don't converge, their sum must tend to infinity. In this case, we get a solution defined on $[x_0, \infty)$. Otherwise the endpoint is simply the value that the series converges to. One such example is when the function is C^1 , which means that function is at least locally bounded (and indeed even has a locally bounded derivatives so is locally Lipschitz). We have seen an example of the latter phenomena with the differential equation

$$x' = x^2$$

which we noted blew up in finite time.

There also exists a more general uniqueness theorem called Osgood's Uniqueness Theorem.

Theorem 12.25 (Osgood's Uniqueness Theorem) Suppose $D \subset \mathbb{R} \times \mathbb{R}^n$ is open and contains some (x_0, y_0) . Assume that for all $(x, y_1), (x, y_2) \in D$ we have that

$$|f(x, y_1) - f(x, y_2)| \leq \varphi(|y_1 - y_2|)$$

where $\varphi : [0, \infty) \rightarrow [0, \infty)$ is continuous and $\varphi(0) = 0$. It also has the properties that for every $a > 0$, we have $\varphi(a) > 0$ and

$$\int_0^a \frac{1}{\varphi(u)} du = \infty$$

Then the initial value problem

$$\begin{aligned} y' &= f(x, y(x)) \\ y(x_0) &= y_0 \end{aligned}$$

has no more than one solution.

Before proving the theorem, a few remarks need to be made. First we will only prove the theorem for $n = 1$. With some more tools and techniques, the given proof can be modified to work in the general case, but we won't worry about that here.

Another important point is that we are not assuming f to be continuous, so we can't use the Cauchy-Peano theorem to assert the existence of a solution. Indeed, we will not prove that a solution exists; only that if it does, it is unique. However, if we had $\varphi(x) = Kx$ for some $K > 0$, then f would be Lipschitz continuous and we reduce back to previous cases. Thus this theorem is certainly more general than what we seen thus far.

Proof. We will prove this using contradiction. Suppose we have two distinct solutions $y_1(x), y_2(x)$ on some interval (α, β) (containing x_0 of course). We define $z(x) = y_1(x) - y_2(x)$. This function satisfies the IVP

$$\begin{aligned} z'(x) &= f(x, y_1(x)) - f(x, y_2(x)) \\ z(x_0) &= 0 \end{aligned}$$

If $z(x) = 0$ everywhere then we are done. So suppose that there is some x_1 such that $z(x_1) \neq 0$. Then importantly

$$z'(x_1) \leq \varphi(|z(x_1)|) < 2\varphi(|z(x_1)|)$$

We now proceed using a comparison argument which we split into cases. The first (and only) case we consider is when $x_1 > x_0$ and $y_1(x_1) > y_2(x_1)$ (the remaining cases are left as exercises for the student(s) : ().

Let v be the solution to the IVP

$$\begin{aligned} v'(x) &= 2\varphi(v) \\ v(x_1) &= z(x_1) =: z_1 \end{aligned}$$

Note that $z_1 > 0$ by the assumption that $y_1(x_1) > y_2(x_1)$. We see that v and z are functions that agree on x_1 but $v' > z'$. Thus the graphs of v and z cannot intersect anywhere else. We will show this leads to a contradiction. In particular we will show that v is an increasing, positive function defined at least on $(-\infty, x_1]$ but since $z(x_0) = 0$, this means there must be another intersection for some $x < x_1$.

In order to verify the above statements for v , it would be nice to have a slightly more explicit formula for it. Since the differentiation equation for v is separable, we can at least get part of the way there. By separating the variables, we know that

$$\int_{v(x)}^{v(x_1)=z_1} \frac{1}{\varphi(v)} dv = \int_x^{x_1} 2 dx = 2(x_1 - x)$$

Or at least this is what v should satisfy as the solution (consider substituting $v(x)$ on the right hand side). Thus we define the map v so that it makes the above statement true. This only makes sense, however, if for every x there exists exactly one v_x such that

$$\int_{v_x}^{z_1} \frac{1}{\varphi(v)} dv = 2(x_1 - x)$$

However this must be the case since $\frac{1}{\varphi}$ is positive so the integral is going to be monotone.

We claim that $\nu(x)$ is defined for all $x < x_1$. We first define

$$\Phi(y) = \int_y^{z_1} \frac{1}{\varphi(u)} du$$

Then in particular $\nu(x)$ is the implicit solution to the equation

$$\Phi(\nu(x)) = 2(x_1 - x)$$

By assumption, $\Phi(y) \rightarrow \infty$ as $y \rightarrow 0$ (recall this was one of the properties of φ). Additionally $\Phi(z_1) = 0$ so $\Phi : (0, z_1] \rightarrow [0, \infty)$. Since $\Phi' = -\frac{1}{\varphi}$, we can conclude that Φ is decreasing. This means that Φ is in particular invertible, allowing us to write

$$\nu(x) = \Phi^{-1}(2(x_1 - x))$$

The domain of Φ^{-1} is $[0, \infty)$ thus the above formula is defined whenever $x_1 - x > 0$ or in other words when $x < x_1$.

Clearly ν is increasing (there are 2 ways of seeing this: for one it is the composition of two decreasing maps and secondly, we can look at the IVP defining it and use the fact that φ is positive on non-zero values). Additionally, since Φ^{-1} takes values in $(0, z_1]$, ν must also be positive. As discussed previously, this implies that the graphs of ν and z intersect (at least) twice: once at z_1 and once prior, leading to the desired contradiction. \square

Remark 12.26. There is a simpler proof if Φ were concave up which would occur if φ were non-decreasing. There is a standard way of turning any function into a non-decreasing function, define

$$\tilde{\varphi}(u) = \sup_{\tilde{u} \in [0, u]} \varphi(\tilde{u})$$

Unfortunately, it is not true that

$$\int_0^a \frac{1}{\varphi(u)} du = \infty \Rightarrow \int_0^a \frac{1}{\tilde{\varphi}(u)} du = \infty$$

Remark 12.27. If $\varphi : [0, \infty) \rightarrow [0, \infty)$ such that $\varphi(0) = 0$, $\varphi(a) > 0$ for $a > 0$ and $\varphi'(0)$ exists, then

$$\int_0^a \frac{1}{\varphi(u)} du = \infty$$

for all $a > 0$

Although we will take a closer look at continuous dependence of solutions on parameters later, let us consider the relatively simple case where we employ the contraction mapping theorem to obtain solutions.

Theorem 12.28 Let (X, d) be a complete metric space. Suppose $F : X \times I \rightarrow X$ (where I is any interval) is such that $F(\cdot, \lambda)$ is a uniform contraction. In other words, there exists some $0 \leq q < 1$ such that for every $\lambda \in I$ and every $x, y \in X$, we have that

$$d(F(x, \lambda) - F(y, \lambda)) \leq qd(x, y)$$

Then there exists a unique fixed point x_λ^* . This defines a map $\lambda \mapsto x_\lambda^*$ that satisfies

$$d(x_\lambda^*, x_{\lambda'}^*) \leq \frac{1}{1-q} d(F(x_\lambda^*, \lambda), F(x_\lambda^*, \lambda'))$$

If F is continuous in λ then the map x^* is continuous (in λ) as well.

Remark 12.29. The word continuous in the final sentence can be replaced with Lipschitz, differentiable, C^k , etc.

Proof. The first statement for the bound on $d(x_\lambda^*, x_{\lambda'}^*)$ can be seen by

$$\begin{aligned} d(x_\lambda^*, x_{\lambda'}^*) &= d(F(x_\lambda^*, \lambda), F(x_{\lambda'}^*, \lambda')) \\ &\leq d(F(x_\lambda^*, \lambda), F(x_\lambda^*, \lambda')) + d(F(x_\lambda^*, \lambda'), F(x_{\lambda'}^*, \lambda')) \\ &\leq d(F(x_\lambda^*, \lambda), F(x_\lambda^*, \lambda')) + qd(x_\lambda^*, x_{\lambda'}^*) \end{aligned}$$

From this it is clear that continuity of F (in λ) implies continuity of x^* . It is not clear how this would generalise for smooth or C^k F . So we present another proof that will give these to us readily.

Suppose F is C^1 . We seek to solve for x such that

$$x = F(x, v)$$

More precisely we know that for every v there is a solution and we wish to decide how these solutions depend on v . Taking everything to one side and taking derivatives what we find is that

$$I - D_x F = 0$$

where I is the identity. Since $\|D_x F\| \leq q < 1$ what we find is that this is invertible (see following lemma). Therefore we can use the inverse function theorem to conclude that the dependence on solutions is smooth/ C^K /etc. \square

Lemma 12.30 *If $A \in \mathbb{R}^{n \times n}$ such that $\|A\| < 1$ then $I - A$ is invertible.*

Proof. The inverse of $I - A$ is given by

$$\sum_{k=0}^{\infty} A^k$$

(note similarity to geometric series). The series converges absolutely since

$$\sum_{k=0}^{\infty} \|A^k\| \leq \sum_{k=0}^{\infty} \|A\|^k$$

Then we see that

$$\begin{aligned} (I - A) \sum_{k=0}^{\infty} A^k &= \lim_{N \rightarrow \infty} (I - A) \sum_{k=0}^N A^k \\ &= \lim_{N \rightarrow \infty} \sum_{k=0}^N A^k - \sum_{k=0}^N A^{k+1} \\ &= \lim_{N \rightarrow \infty} I - A^{N+1} \end{aligned}$$

Since A^N converges to 0 (check norms), we get the desired statement. \square

Remark 12.31. A remark that is entirely unrelated to ODEs but also so delightful that I cannot help but mention it here. We've seen that plugging matrix into e^x and $\frac{1}{1+x}$ kind of makes sense (admittedly the latter only working for some matrices). In both cases we used the fact that these functions can be approximated via polynomials. Thus as you might imagine, any time we can approximate a function via polynomials, we can consider what happens if we input a matrix! The even more amazing fact is that any continuous function on a compact interval can be approximated via polynomials (see [Weirstrass approximation theorem](#)). So in fact we can apply any continuous function to a matrix. This is what leads to the so-called [functional calculus](#).

Definition 12.32 (Locally Lipschitz). A map $f : A \rightarrow \mathbb{R}^n$ is said to be locally Lipschitz, if for every compact set $K \subset A$ there exists a constant L (which may depend on K) such that $|f(x) - f(y)| \leq L|x - y|$ for all $x, y \in K$.

We will finish with a global existence and uniqueness theorem. Essentially, what we want to say is that if we have a solution on the maximal time interval then the solution blows up (in norm) or the solution approaches the boundary of the domain (or both).

Theorem 12.33 Suppose $D \subset \mathbb{R}^n$ is open and connected and $f : D \rightarrow \mathbb{R}^n$ is a locally Lipschitz vector field. Let v be some vector in D . Then there exists a unique maximal interval of existence $I_{\max} = (\underline{T}, \overline{T}) \ni 0$ such that

1. The IVP

$$\begin{aligned} x' &= f(x) \\ x(0) &= v \end{aligned}$$

has a unique solution on I_{\max} and

2. if $\overline{T} < \infty$ then

$$\lim_{t \rightarrow \overline{T}} |x(t)| + \frac{1}{d(x(t), \partial D)} = \infty$$

Proof. We already know that the solution exists and is unique so we only need to verify the second statement. We do so by showing the contrapositive. In particular we will show that if $T < \overline{T}$ such that $|x(t)| \neq \infty$ as $t \rightarrow T$ and $\lim_{t \rightarrow T} d(x(t), \partial D) \neq 0$ then there is a solution that goes past T .

Let $T > 0$ be arbitrary and suppose the two limit conditions from the previous paragraphs hold (that is $|x(t)|$ doesn't tend to infinity and $x(t)$ always remains some distance from the boundary). Then there exists a sequence $(t_j)_{j \in \mathbb{N}}$ that converges to T such that $m := \sup_{j \in \mathbb{N}} |x(t_j)|$ is finite and $s := \inf_{j \in \mathbb{N}} d(x(t_j), \partial D)$ is positive (or in other words is not 0). Define

$$K = \{z \in D : |z| \leq m, d(z, \partial D) \geq s\}$$

K is closed and bounded, hence it is compact (note that the fact that K is closed might not be as apparent as it seems. In particular $K = D \cap \overline{B(0, m)} \cap \{z \in \mathbb{R}^n : d(z, \partial D) \geq s\}$. While the last two sets are certainly closed, since D is open, it is not clear that the intersection should also be closed. Nevertheless it is true that K is closed because $D \cap \{z \in \mathbb{R}^n : d(z, \partial D) \geq s\} = \{z \in \mathbb{R}^n : d(z, D^c) \geq s\}$ which is obviously closed being the preimage of a closed set under a continuous function).

On K then, we have a Lipschitz constant L and hence by theory we have built up so far (see [Theorem 12.21](#)) we know there is some $\tau > 0$ such that for every $w \in K$ we have a unique solution to the IVP

$$\begin{aligned} y' &= f(y) \\ y(0) &= w \end{aligned}$$

on $[-\tau, \tau]$ (to be fair we started with an open subset of \mathbb{R}^{n+1} in that proof whereas we have a compact set here. However, by looking at the proof(s), we realise that we only needed the fact that the set contained a rectangle. Given the definition of K , we can see that this should be true).

Let y_j be the solutions of the above IVP for $w = x(t_j)$. Then $y_j(t - t_j)$ is defined on $[t_j - \tau, t_j + \tau]$ which agrees with $x(t)$ at t_j (note that since K contains all the t_j by its definition, this τ holds for all t_j). Since x is maximal and unique, we conclude that x is defined at least until $t_j + \tau$. Now suppose we pick some t_j such that $T - t_j < \tau$ (we can definitely do so since t_j converge to T). But this means that $x(t)$ is defined at least till $t_j + \tau > T$. This holds for all T so $x(t)$ is defined on $[0, \infty)$. See [Figure 17](#) to see a quick sketch of the above argument. □

13 Continuous Dependence

We know now that under certain conditions solutions exist and are unique and we even have that solutions vary continuously as the initial conditions (or parameters) vary continuously. This means, for example then, if we have a sequence v_k in \mathbb{R}^n that converge to some v then the solution to

$$\begin{aligned} x' &= f(t, x) \\ x(0) &= v \end{aligned}$$



Figure 17: We can get arbitrarily close to T but always extend by a fixed amount. This allows us to go past T

can be found by taking the appropriate limits with the v_k . In particular, if we use x_k to denote the solutions to

$$\begin{aligned} x' &= f(t, x) \\ x(0) &= v_k \end{aligned}$$

Then the solution to the original IVP can be found by $x(t) = \lim_{k \rightarrow \infty} x_k(t)$ (again, assuming suitable conditions for f). This is easier to express using flow notation as we can equivalently express the above as $\lim_{k \rightarrow \infty} \Phi_t(v_k) = \Phi_t(v)$ for all t . (see [Section 11: Dynamical Systems](#) for definition of the flow map). Although this is lovely, it would be even more lovely if we could have an estimate for how fast the convergence is. To get there, we first need a generalised version of Grönwall's lemma.

Theorem 13.1 (Grönwall's Lemma (Generalised)) *Suppose we have continuous functions $f : [a, b] \rightarrow \mathbb{R}$ and $g : [a, b] \rightarrow \mathbb{R}^+$. Suppose $y : [a, b] \rightarrow \mathbb{R}$ is continuous and satisfies*

$$y(t) \leq f(t) + \int_a^t g(s)y(s)ds$$

for all $t \in [a, b]$. Then we have that

$$y(t) \leq f(t) + \int_a^t f(s)g(s) \exp\left(\int_s^t g(u)du\right)ds$$

In particular if $f(t) \equiv k$ for some constant k , then

$$y(t) \leq k \exp\left(\int_a^t g(u)du\right)$$

The statement for f being constant doesn't seem to be an obvious rearrangement of the general statement (although it is). We can verify it by employing a simple technique for computing such double integrals that we will use in the proof so is worth mentioning. Recall we are trying to evaluate

$$k + k \int_a^t g(s) \exp\left(\int_s^t g(u)du\right)ds$$

We define

$$M(s) = \int_s^t g(u)du = - \int_t^s g(u)du$$

Then

$$\begin{aligned}
k + k \int_a^t g(s) \exp\left(\int_s^t g(u) du\right) ds &= k + k \int_a^t -M'(s) \exp(M(s)) ds \\
&= k - k(\exp(M(s))|_a^t) \\
&= k - k(\exp(M(t)) - \exp(M(a))) \\
&= k \exp(M(a)) \\
&= k \exp\left(\int_a^t g(u) du\right)
\end{aligned}$$

Having verified the theorem in some ways, we can try proving it.

Proof 1. For the first proof we will iterate the the assumed inequality. We will see that doing so gives us the Taylor series of an exponential with an additional term that goes to 0 as we continually iterate the inequality.

First note that without loss of generality we can assume $|y(t)| \leq 1$ (y is continuous on a compact set so $\sup_{t \in [a, b]} |y(t)|$ is finite. Dividing by this quantity, we see that the assumed inequality is maintained). Additionally we will only show the case for when f is a constant (the general case is left as an exercise).

By assumption we know that

$$y(t) \leq k + \int_a^t g(s)y(s)ds$$

Applying this inequality to the integral what we conclude is that

$$y(t) \leq k + \int_a^t g(s)y(s)ds \leq k + \int_a^t g(t_1) \left[k + \int_a^{t_1} g(s)y(s)ds \right] dt_1$$

Expanding the right hand side, what we get is

$$k + \int_a^t g(t_1) \left[k + \int_a^{t_1} g(s)y(s)ds \right] dt_1 = k + k \int_a^t g(t_1)dt_1 + \int_a^t g(t_1) \int_a^{t_1} g(s)y(s)ds dt_1$$

Unfortunately we need to apply the inequality again to get something useful. Doing so and expanding it like above, what we find is that

$$y(t) \leq k + k \int_a^t g(t_1)dt_1 + k \int_a^t g(t_1) \int_a^{t_1} g(t_2)dt_2 dt_1 + k \int_a^t g(t_1) \int_a^{t_1} g(t_2) \int_a^{t_2} g(t_3)y(t_3)dt_3 dt_2 dt_1$$

Although this is quite a mess, we won't worry about all of it (for now). Instead just focus on the double integral term. As a first step we will name it (very often a useful first step)

$$I = \int_a^t g(t_1) \int_a^{t_1} g(t_2)dt_2 dt_1 = \int_a^t M'(t_1)M(t_1)dt_1$$

where

$$M(t_1) = \int_a^{t_1} g(s)ds$$

Therefore

$$I = \frac{1}{2}[M(t)]^2 - \frac{1}{2}[M(a)]^2$$

However, $M(a) = 0$ so

$$I = \frac{1}{2}[M(t)]^2 = \frac{1}{2} \left(\int_a^t g(s)ds \right)^2$$

With a bit of working out or a bit of faith, we conclude that repeatedly applying the inequality to the integrals gets us

$$y(t) \leq k + k \int_a^t g(t_1)dt_1 + \frac{k}{2!} \left(\int_a^t g(t_1)dt_1 \right)^2 + \cdots + \frac{k}{n!} \left(\int_a^t g(t_1)dt_1 \right)^n + \int_a^t g(t_1) \cdots \int_a^{t_n} g(t_{n+1})y(t_{n+1})dt_{n+1} \cdots dt_1$$

For completeness, we can verify this holds for the triple integral term which would appear after we apply the inequality to the integral again. In other words we are trying to evaluate

$$J = \int_a^t g(t_1) \int_a^{t_1} g(t_2) \int_a^{t_2} g(t_3) dt_3 dt_2 dt_1$$

Using our previous notation and result, we see that

$$J = \int_a^t M'(t_1) \cdot \frac{1}{2} [M(t_1)]^2 dt_1$$

Therefore

$$J = \frac{1}{3!} [M(t)]^3$$

as was claimed. Hopefully we can convince ourselves now that the above summation is correct. We can see the desired exponential term appearing (as a Taylor series). Thus if the claim is to hold true then we should find that the final term goes to 0 as we repeatedly apply the inequality. Now we verify that this happens.

Since g is always positive what we find is that

$$\left| \int_a^t g(t_1) \cdots \int_a^{t_n} g(t_{n+1}) y(t_{n+1}) dt_{n+1} \cdots dt_1 \right| \leq \int_a^t g(t_1) \cdots \int_a^{t_n} g(t_{n+1}) |y(t_{n+1})| dt_{n+1} \cdots dt_1$$

We assumed that $|y(t)| \leq 1$ for all $t \in [a, b]$. Therefore the right hand side above is less than or equal to

$$\int_a^t g(t_1) \cdots \int_a^{t_n} g(t_{n+1}) dt_{n+1} \cdots dt_1$$

However we already know that this is equal to

$$\frac{1}{(n+1)!} \left(\int_a^t g(s) ds \right)^{n+1}$$

Clearly this goes to 0 as $n \rightarrow \infty$ (it is a term of a convergent infinite series for example). Therefore we conclude that

$$y(t) \leq k \exp \left(\int_a^t g(s) ds \right)$$

□

As usual there is a second proof of the theorem. This one, a bit more indirect but slicker.

Proof 2. We define

$$z(t) = \int_a^t g(s) y(s) ds$$

By assumption we have that

$$y(t) \leq f(t) + \int_a^t g(s) y(s) ds$$

By multiplying both sides by $g(t)$ (remember g is positive) and rearranging what we find is that

$$g(t)y(t) - g(t) \int_a^t g(s)y(s) ds \leq g(t)f(t)$$

Therefore

$$z'(t) - g(t)z(t) \leq f(t)g(t)$$

Seeing the left hand side one is reminded of linear exact differential equations and integrating factors (see [Subsection 17.3: Exact Differential Equations](#)). In particular, we can write the left hand side as the derivative of a product if we multiply both sides by

$$\exp \left(- \int_a^t g(u) du \right)$$

(note the inequality is maintained since the exp is always positive). The inequality now becomes

$$z'(t) \exp\left(-\int_a^t g(u) du\right) - z(t)g(t) \exp\left(-\int_a^t g(u) du\right) \leq f(t)g(t) \exp\left(-\int_a^t g(u) du\right)$$

Although the left hand side looks like something of a mess, we know (by construction) that it is equal to $w'(t)$ where

$$w(t) = z(t) \exp\left(-\int_a^t g(u) du\right)$$

Hence we get

$$\begin{aligned} \int_a^t w'(s) ds &\leq \int_a^t f(s)g(s) \exp\left(-\int_a^s g(u) du\right) ds \\ z(t) \exp\left(-\int_a^t g(u) du\right) &\leq \int_a^t f(s)g(s) \exp\left(-\int_a^s g(u) du\right) ds \\ z(t) &\leq \int_a^t f(s)g(s) \exp\left(\int_s^t g(u) du\right) ds \\ f(t) + \int_a^t g(s)y(s) ds &\leq f(t) + \int_a^t f(s)g(s) \exp\left(\int_s^t g(u) du\right) ds \end{aligned}$$

where for the second line we use the fact that $w(a) = 0$ (which follows from the fact that $z(a) = 0$). Since the left hand side is greater or equal to $y(t)$ by assumption, we are done. \square

Now that we have Grönwall's Lemma, we can use it to provide estimates for how much solutions can differ, if we only change the initial conditions slightly. In particular, we have the following theorem.

Theorem 13.2 Let $x_1 : [a, b] \rightarrow \mathbb{R}^n$ and $x_2 : [a, b] \rightarrow \mathbb{R}^n$ be differential functions such that $|x_1(a) - x_2(a)| \leq \delta$. Let $f : [a, b] \times \mathbb{R}^n \rightarrow \mathbb{R}^n$ be function that is Lipschitz in the second variable with Lipschitz constant L . Suppose

$$\begin{aligned} |x'_1(t) - f(t, x_1(t))| &\leq \epsilon_1 \\ |x'_2(t) - f(t, x_2(t))| &\leq \epsilon_2 \end{aligned}$$

Then

$$|x_1(t) - x_2(t)| \leq \delta e^{L(t-a)} + (\epsilon_1 + \epsilon_2) \frac{e^{L(t-a)} - 1}{L}$$

Remark 13.3. Suppose $\epsilon_1 = \epsilon_2 = 0$. Then x_1 and x_2 satisfy the same differential equation but with slightly different initial conditions. In this case the conclusion of the theorem is that $|x_1(t) - x_2(t)| \leq \delta e^{L(t-a)}$, which is to say that the solutions can differ at most by some exponential term. If we also have $\delta = 0$, we get yet another proof of uniqueness.

Proof. Let $\epsilon = \epsilon_1 + \epsilon_2$. We define $g(t) = x_1(t) - x_2(t)$. We want to show that g is bounded by some kind of exponential term and as we said we want to use Grönwall. For this we need a bound on g in terms of some kind of integral. First we will find an estimate for g' (as you can imagine this will be easier since we have some strong conditions on f which approximates x'_1 and x'_2 fairly well) and then use this to find an estimate for g .

$$\begin{aligned} |g'(t)| &= |x'_1(t) - x'_2(t)| \\ &= |(x'_1(t) - x'_2(t)) - f(t, x_1(t)) + f(t, x_1(t)) - f(t, x_2(t)) + f(t, x_2(t))| \\ &\leq |f(t, x_1(t)) - f(t, x_2(t))| + |x'_1(t) - f(t, x_1(t))| + |x'_2(t) - f(t, x_2(t))| \\ &\leq L|g(t)| + \epsilon \end{aligned}$$

Now that we have an estimate for g' we can try to find an estimate for g as well. In particular we get

$$\begin{aligned} |g(t)| &= \left| g(a) + \int_a^t g'(s) ds \right| \\ &\leq |g(a)| + \int_a^t |g'(s)| ds \\ &\leq \delta + \int_a^t L|g(s)| + \epsilon ds \\ &= \delta + \epsilon(t-a) + \int_a^t L|g(s)| ds \end{aligned}$$

Using [Grönwall](#) then what we conclude is that

$$|g(t)| \leq \delta + \epsilon(t-a) + \int_a^t L(\delta + \epsilon(s-a))e^{L(t-s)} ds$$

Computing the right hand side⁹, we find it to be equal to

$$\delta e^{L(t-a)} + \frac{\epsilon}{L}(e^{L(t-a)} - 1)$$

giving us the stated inequality. □

We now make a statement about solving (nonautonomous) linear systems that could have been stated earlier but is here because we will use it shortly.

Corollary 13.3.1 *Let $(t_0, x_0) \in I \times \mathbb{R}^n$ be arbitrary, where I is some interval in \mathbb{R} . Let $A(t)$ be a continuous family of $n \times n$ matrices. Then the IVP*

$$\begin{aligned} x'(t) &= A(t)x \\ x(0) &= x_0 \end{aligned}$$

has a unique solution on all of I . Moreover this solution is given by

$$x(t) = e^{\int_{t_0}^t A(s) ds} x_0$$

Proof. We already know that the solution exists and is unique from everything we've done. We can check that the solution is as given by verifying that it satisfies the ODE. Indeed we see that

$$x'(t) = A(t)e^{\int_{t_0}^t A(s) ds} x_0 = A(t)x$$

and $x(t_0) = x_0$. By uniqueness, this is the solution to the IVP. □

14 Flow of Dynamical Systems

Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is C^1 . Then we know for every $x \in \mathbb{R}^n$ there is a unique solution satisfying the IVP

$$\begin{aligned} y' &= f(y) \\ y(0) &= x \end{aligned}$$

Suppose also that the solution is defined for all $t \in \mathbb{R}$. Recall from [Section 11: Dynamical Systems](#) that we can define a map $\Phi_t(x) = y(t)$ (where y is the (unique) solution to the IVP). In effect Φ_t maps initial points x to where the differential equation sends them at time t . As mentioned previously, the flow map satisfies the semigroup property: $\Phi_{s+t} = \Phi_s \circ \Phi_t$ (note this property does *not* hold if we do not have uniqueness of solutions).

⁹There's some integration by parts nonsense. If you really want to, feel free to give it a go.

Sometimes we alternatively use the notation $\Phi(t, x)$. This is especially useful when we wish to take the derivative with respect to time for example. In fact let us do exactly this to try and calculate the so-called time derivative $\frac{\partial \Phi}{\partial t}$. This means that we hold the initial value x constant and vary t . Clearly this must be the solution $y(t)$. Therefore

$$\frac{\partial \Phi}{\partial t}(t, x) = f(\Phi(t, x))$$

Additionally, it is clear by definition that $\Phi_0(x) = \Phi(0, x) = x$, so Φ_0 is the identity. Using the semigroup property, we can conclude that Φ_t is a bijection for all t with inverse Φ_{-t} . The injectivity of Φ_t is exactly equivalent to a previous statement we have made of how the uniqueness of solutions implies that the solutions can never intersect. The surjectivity is a confirmation of a long held suspicion of ours that the solution space for $y' = f(y)$ for $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is n -dimensional.

Now that we have considered the time derivative, we should also consider its sibling the space derivative $\frac{\partial \Phi}{\partial x}$. First note we know this exists and is continuous by [Theorem 12.28](#) (or more precisely by the remark following it). As you can imagine however, computing it will be harder.

Consider our standard IVP

$$\begin{aligned} x'(t) &= f(x(t)) \\ x(0) &= x_0 \end{aligned}$$

for $t \in J$ where J is a closed interval containing 0 and f is C^1 . For $t \in J$, we define $A(t) = Df(x(t))$ (this is the Jacobian of f at $x(t)$). Since f is C^1 , we know that A is continuous function. We then define the *variational equation along the solution $x(t)$* to be the unique solution $u(t)$ to

$$\begin{aligned} u' &= A(t)u \\ u(0) &= u_0 \end{aligned}$$

(recall we know that unique solutions exist for every u_0 and are defined on all of J by [Corollary 13.3.1](#)). Note that this is a (admittedly non-autonomous) linear system of equations, hence the [linearity principle](#) still holds. What we will show is that if u_0 is small the map $t \mapsto x(t) + u(t)$ (with $x(t)$ as above) is a good approximation to the solution of

$$\begin{aligned} x' &= f(x) \\ x(0) &= x_0 + u_0 \end{aligned}$$

We formalise this in the following proposition.

Proposition 14.1 *Let $D \subset \mathbb{R}^n$ and $f : D \rightarrow \mathbb{R}^n$ be C^1 . Let J be a closed interval containing 0 and $x(t)$ is the solution to $x' = f(x)$ with $x(0) = x_0$. Let $u(t)$ be the solution to $u' = Df(x(t))u$ satisfying $u(0) = \xi$ (in other words u is the variation equation along $x(t)$) and let $y(t)$ be the solution to $x' = f(x)$ satisfying $y(0) = x_0 + \xi$. Then*

$$\lim_{\xi \rightarrow 0} \frac{|y(t) - x(t) - u(t)|}{|\xi|}$$

converges uniformly to 0.

Remark 14.2. Uniform convergence in this context means that for every $\epsilon > 0$ there exists some $\delta > 0$ such that if $|\xi| \leq \delta$ then $|y(t) - x(t) - u(t)| \leq \epsilon |\xi|$ for every $t \in J$.

Proof. As usual we go to integral equations. We have

$$\begin{aligned} x(t) &= x_0 + \int_0^t f(x(s)) ds \\ y(t) &= x_0 + \xi + \int_0^t f(y(s)) ds \\ u(t) &= \xi + \int_0^t Df(x(s)) u(s) ds \end{aligned}$$

Then

$$|y(t) - x(t) - u(t)| \leq \int_0^t |f(y(s)) - f(x(s)) - Df(x(s))u(s)| ds \quad (14.1)$$

Let us denote the left hand side as $g(t)$. Finding the Taylor expansion of f centered at x what we find is that

$$f(y) = f(x) + Df(x)(y - x) + R(y - x)$$

where $R(y - x)$ is the remainder term such that

$$\lim_{y \rightarrow x} \frac{R(y - x)}{|y - x|} = 0$$

Substituting this expansion of f into the above inequality, we get

$$\begin{aligned} g(t) &\leq \int_0^t |f(x(s)) + Df(x(s))(y(s) - x(s)) + R(y(s) - x(s)) - f(x(s)) - Df(x(s))u(s)| ds \\ &\leq \int_0^t |Df(x(s))(y(s) - x(s) - u(s))| ds + \int_0^t |R(y(s) - x(s))| ds \end{aligned}$$

Since J is a closed interval it is in particular compact hence $|Df(x(s))|_{s \in J}$ achieves its maximum. Let N be this maximum. Then N serves as a Lipschitz constant for f . Let $\epsilon > 0$ be given. By uniform convergence of $\frac{R(y-x)}{|y-x|}$ to 0, we know there is some $\delta_1 > 0$ so that $|y - x| < \delta_1$ then $|R(y - x)| \leq \epsilon |y - x|$. But recall that y and x are solutions to the same differential equation with the differing initial condition. We already have an estimate for how far these can differ, by [Theorem 13.2](#). We know f is C^1 so it is locally Lipschitz. Thus for every compact set C , we can find a constant K such that f restricted to C is Lipschitz with Lipschitz constant K . There is a compact set containing the images of x and y (we can in particular take the union of their images). Let K be the Lipschitz constant for this compact set. Then we know by the aforementioned theorem that

$$|y(s) - x(s)| \leq |\xi| e^{Ks}$$

for every $s \in J$. Since J is compact the right hand attains some maximum. Thus by choosing an appropriate $\delta > 0$ we can ensure that if $|\xi| < \delta$ then $|\xi| e^{Ks} < \delta_1$ for all $s \in J$. Thus if $|\xi| < \delta$, then we know that $|R(y(s) - x(s))| \leq \epsilon |y(s) - x(s)| \leq \epsilon |\xi| e^{Ks}$ for every $s \in J$. Under this condition, what we find then is that

$$\begin{aligned} g(t) &\leq \int_0^t |Df(x(s))(y(s) - x(s) - u(s))| ds + \int_0^t |R(y(s) - x(s))| ds \\ &\leq \int_0^t |Df(x(s))| g(s) ds + \int_0^t \epsilon |\xi| e^{Ks} ds \\ &\leq \int_0^t N g(s) ds + \epsilon |\xi| \int_0^t e^{Ks} ds \end{aligned}$$

Note that the integral in the second term is simply a constant which we will denote c . Thus using [Grönwall](#) (specifically the case when f is a constant function), we conclude that

$$g(t) \leq c\epsilon |\xi| e^{Nt}$$

Thus

$$\lim_{\xi \rightarrow 0} \frac{g(t)}{|\xi|}$$

goes to 0 uniformly since ϵ was chosen arbitrarily. \square

Why is this proposition important? Well, as we've said it is often impossible to solve differential equations explicitly. However, perhaps we can find $x(t)$ for some special x_0 (for example, we might be able to find equilibria, which would correspond to the zeroes of f). This means that we can approximate solutions near these special x_0 to get at least get some idea of the behaviour.

In fact we can compute $\frac{\partial \Phi}{\partial x}$ a bit more explicitly now. First note $y(t)$ and $x(t)$ in the above proposition are simply $\Phi(t, x_0 + \xi)$ and $\Phi(t, x_0)$. We think of $u(t)$ as a function of 2 variables: namely we define $u(t, \xi)$ to

be the solution to the variational equation along $x(t)$ (or $\Phi(t, x_0)$) satisfying $u(0, \xi) = \xi$. The linearity principle (see [Subsection 5.2: Superposition/Linearity Principle](#)) ensures that u is linear in ξ . The definition of the differential of a map g is the (unique) linear map A that satisfies

$$\lim_{h \rightarrow 0} \frac{|g(x+h) - g(x) - Ah|}{|h|} = 0$$

Therefore the map $\xi \mapsto u(t, \xi)$ is the partial derivative of $\Phi(t, x)$ with respect to x , which we may also denote $D\Phi_t(x)$. In fact since we have an explicit solution for u , we can write

$$D\Phi_t(x_0) = e^{\int_0^t Df(x(s))ds} x_0$$

Using this, we can then verify that $D\Phi_t(x)$ satisfies the following IVP

$$\begin{aligned} \frac{\partial}{\partial t}(D\Phi_t(x)) &= Df(\Phi_t(x))D\Phi_t(x) \\ D\Phi_0(x) &= I \end{aligned}$$

where I is the identity map. This matches the IVP for u as well (we use the fact that $A(t) = Df(x(t)) = Df(\Phi_t(x))$ and $u(0, x_0) = x_0$).

To recap, we have managed to find some formulae and descriptions of the variation equation. The reason we wanted to solve the variation equation is because it lets us approximate solutions around known solutions. The question then is, have we actually done this? We can find $u(t)$ by solving an integral or by solving an initial value problem, neither of which is particularly easy in general. And indeed, what we've done is converted one difficult problem into another (hurrah.). However there are some important cases where we can solve for things and are hence worth exploring.

Suppose the known solution $x(t)$ is an equilibrium point. In other words $x(t) \equiv a$ for some a . Then $A(t) = Df(a)$ for all t . Hence what we find is that

$$u(t) = D\Phi_t(a) = e^{tDf(a)}$$

Therefore in a neighbourhood around equilibria, we can approximate the flow using a linear system.

14.1 Classification of Flow

Much like we did with linear systems in the plane, we want to classify all (possibly non-linear) systems in any number of dimensions. This as daunting a task as it sounds. At the very least, we first need to start with how this classification will work, namely when can we say that two flows are 'the same'? This leads us to the notion of similarity.

We will say that two flows Φ and Ψ are similar (or more precisely conjugate) when there is a bijection h such that

$$h^{-1} \circ \Phi \circ h = \Psi$$

Remark 14.3. Note how similar this is to similarity of matrices: we say two matrices A, B are similar if there is an invertible T such that $T^{-1}AT = B$.

Note how we have not made any assumption about h being continuous/differentiable/smooth/etc. Such assumptions can vary based on the field of study. Taking h to be a homeomorphism is quite common (and as we will show, somewhat easy to achieve). If h is indeed just a homeomorphism, we say we have *topological conjugacy*. Getting h to be a diffeomorphism is much nicer, but also a lot harder.

Now that we have the general overview, let's try and be a bit more precise. Suppose x is a solution to the non-linear system

$$\begin{aligned} x' &= f(x) \\ x(0) &= v \end{aligned}$$

and z is its linearization. In other words, z is the solution to

$$\begin{aligned} z'(t) &= Df(x(t))z \\ z(0) &= w \end{aligned}$$

(where w is taken to be small). As we've discussed, $x + z$ is an approximation solution to $x' = f(x)$ with $x(0) = v + w$ (this is exactly [Proposition 14.1](#)). Thus we would like to say that x is conjugate to $x + z$. Unfortunately in general this is not true, even in the simplest case where we approximate near a steady state. Consider the following example.

14.1.1 Example 1

Suppose we are given the differential equation

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \begin{pmatrix} x^2 \\ y^2 \end{pmatrix}$$

Clearly, the only equilibrium point is the origin and

$$Df(0,0) = \begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix}$$

Thus the linearised system has constant solutions. And indeed, if we were to look in a small neighbourhood near the origin the solutions (on a small time interval) do look roughly constant (see [here](#)). However, and this is the key, the constant solution is *not* conjugate to the non-linear system. Intuition is enough to guide us here: if the flow of a system is constant, then there is no motion. We really should not consider this equivalent to a system that does have motion. It might be a useful exercise to prove to yourself that constant systems cannot be conjugate to non-constant ones, before I do that right now.

Quite simply, the key is that if the solutions to a system are constant then the flow map Φ_t is simply the identity for all t . Thus $h \circ \Phi_t \circ h^{-1}$ must be the identity as well.

Nevertheless, there will be cases where we *do* get conjugacy, as in the following example.

14.1.2 Example 2

Consider the system

$$\begin{aligned} x' &= x + y^2 \\ y' &= -y \end{aligned}$$

Then $f(x, y) = (x + y^2, -y)$. Clearly we have only one equilibrium point, which is at the origin again. Thus the linearisation around it is given by

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \underbrace{\begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}}_A \begin{pmatrix} x \\ y \end{pmatrix}$$

Without any further work, we can conclude that the linearised system is unstable and has a saddle. Note that if y is small, then y^2 would be even smaller. Thus we would expect the non-linear system to look quite similar to its linearisation around the origin.

In this case we can actually solve the nonlinear system, allowing us to test our hypotheses. We can easily solve for y to get

$$y(t) = y_0 e^{-t}$$

Then we need to solve

$$x' = x + y_0^2 e^{-2t}$$




Figure 18: Phase portrait of linearised system, [source](#)

We know that the general solution is going to be of the form

$$x(t) = ce^t + x_p(t)$$

where $x_p(t)$ is a particular solution. We will guess x_p to be of the same form as the inhomogeneity (this is often not a bad guess), so in particular we will guess

$$x_p(t) = be^{-2t}$$

where b is a constant to be determined. Substituting this into the ODE (for x), we find

$$b = -\frac{y_0^2}{3}$$

thus

$$x(t) = \left(x_0 + \frac{y_0^2}{3}\right)e^t - \frac{y_0^2}{3}e^{-2t}$$

(the coefficient of e^t is of that form so that $x(0) = x_0$). Plotting this we get the phase portrait in [Figure 19](#).

The parabola going through the origin is given by $x = -\frac{1}{3}y^2$. If our initial conditions lie on this parabola, we get

$$x(t) = -\frac{y_0^2}{3}e^{-2t}, y(t) = y_0e^t$$

which draws out the above parabola.

Comparing phase portraits of the two systems, we can maybe convince ourselves of a certain similarity. The hope is that there is a way of ‘translating’ one to the other (this is after all exactly what the homeomorphism would do). In fact in this case we can check that the change of variables given by

$$T(x, y) = \left(x + \frac{1}{3}y^2, y\right)$$




Figure 19: Phase portrait of non-linear system, [source](#)

maps the flow of the linear system onto the non-linear one.

$$\begin{aligned}
 T^{-1}e^{tA}T(x_0, y_0) &= T^{-1}e^{tA}\left(x_0 + \frac{1}{3}y_0^2, y_0\right) \\
 &= T^{-1}\left(e^t\left(x_0 + \frac{1}{3}y_0^2\right), y_0e^{-t}\right) \\
 &= \left(e^t\left(x_0 + \frac{1}{3}y_0^2\right) - \frac{y_0^2}{3}e^{-2t}, y_0e^{-t}\right)
 \end{aligned}$$

which we know is exactly the flow of the non-linear system. In this case, the map T linearised the system globally and was a diffeomorphism. Most of the time, this is not the case; usually, we only get local homeomorphisms as we see in the following example.

14.1.3 Example 3

Suppose

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \begin{pmatrix} \frac{1}{2} & -1 \\ 1 & \frac{1}{2} \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} - \frac{1}{2}(x^2 + y^2) \begin{pmatrix} x \\ y \end{pmatrix}$$

Once again we see that the (only) steady state is at the origin and we can easily find that the differential at the origin is

$$A := Df(0,0) = \begin{pmatrix} \frac{1}{2} & -1 \\ 1 & \frac{1}{2} \end{pmatrix}$$

thus the linearisation is given by

$$e^{tA} = e^{\frac{t}{2}} \begin{pmatrix} \cos t & -\sin t \\ \sin t & \cos t \end{pmatrix}$$

Hence we conclude that the linearised system is a (counterclockwise) unstable spiral.

With a slight change of variables, we will be able to see the qualitative behaviour of the non-linear system more easily. In particular, we will change to polar coordinates, $x(t) = r \cos \theta$, $y(t) = r \sin \theta$. Using some combinations of the product and chain rules, we find

$$\begin{aligned}x' &= r' \cos \theta - r(\sin \theta)\theta' \\y' &= r' \sin \theta + r(\cos \theta)\theta'\end{aligned}$$

Multiplying and rearranging things appropriately, we get

$$\begin{aligned}xx' + yy' &= rr' \\-x'y + y'x &= r^2\theta'\end{aligned}$$

Therefore

$$r' = \frac{1}{r}(xx' + yy')$$

where upon substituting things from the given ODE, we find

$$r' = \frac{1}{2}r - \frac{1}{2}r^3$$

Similarly we find

$$\theta' = \frac{1}{r^2}(-x'y + y'x) = 1$$

We see that r is stable when $r = 0, 1$ or -1 . Obviously r cannot be negative and we are ignoring the case of $r = 0$ (polar coordinates don't work here anyway and we've divided by r too many times to give that a go), so the only interesting case is when $r = 1$. In this case, the solution is given by the unit circle since θ is a simple linear function.

Note that when $0 < r < 1$, $r' > 0$ thus r is increasing and when $r > 1$, we have $r' < 0$ so r is decreasing. Thus solutions that begin inside the unit circle grow towards it and solutions that begin outside decay towards it. The phase portrait is visualised [here](#).

Certainly then we cannot get a global homeomorphism to a linear system because no linear system has dynamics where it grows near the origin but shrinks far away from it. However we can get a local homeomorphism around the origin. Let Φ be the flow for the non-linear system and Ψ the flow for the linearised system. Fix some $r_0 < 1$. Then for every $x \in B_{r_0}(0) \setminus \{0\}$ there is some unique time τ such that $\Phi_\tau(x) \in B_{r_0}(0)$ (this is because the radius r is increasing on such solutions). Thus we can think of τ as a function of the initial point x . Then we define the homeomorphism h by $h(x) = \Psi_{-\tau}\Phi_\tau(x)$. We will show that this is a homeomorphism later (see [Proposition 14.12](#)).

14.1.4 Example 4

Consider the system

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \begin{pmatrix} -y + \epsilon x(x^2 + y^2) \\ x + \epsilon y(x^2 + y^2) \end{pmatrix}$$

As usual, the only equilibrium point is at the origin where the differential is

$$Df(0,0) = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}$$

The flow of the linearised system is therefore a center. In order to compare this with the original, non-linear system we write the original system in polar coordinates to get

$$r' = \epsilon r^3, \theta' = 1$$

Therefore when $\epsilon > 0$, we have a source at the origin and when $\epsilon < 0$, we have a sink. Neither of these can be conjugate to a center (if they were, then we would conclude that a sink and source are conjugate and certainly no reasonable system of classification should consider such systems to be equivalent).

14.1.5 Example 5

Consider the system

$$\begin{pmatrix} x \\ y \end{pmatrix}' = \begin{pmatrix} x^2 \\ -y \end{pmatrix}$$

Then the linearisation around the equilibrium at the origin is

$$Df(0,0) = \begin{pmatrix} 0 & 0 \\ 0 & -1 \end{pmatrix}$$

In this case the linearisation produces an entire line of stability. However clearly this cannot be the case with the non-linear system, as at least one of y' and x' is non-zero everywhere (besides the origin). Indeed the origin is an unstable equilibrium since solutions on the x -axis for example always move to the right.

In the following section, we will see when the linearisation is actually similar. As you might guess, the linearisation stops being a good description of the system when a certain quantity is 0.

14.2 Stability

We've mentioned stability and instability quite a few times so far, without defining precisely what they mean (although we have an intuitive idea already). Let us remedy this now.

Definition 14.4 (Stable/Unstable Equilibria). Let $(\Phi_t)_{t \geq 0}$ be a dynamical system on a domain $D \subset \mathbb{R}^n$. Let a be an equilibrium point. Then a is said to be *stable* if for every $\epsilon > 0$ there exists some $\delta > 0$ such that $|x - a| < \delta$ implies that $\sup_{t \geq 0} |\Phi_t(x) - a| < \epsilon$.

We can equivalently define stable equilibria in the language of open sets: an equilibrium point a is said to be stable if for every neighbourhood V of a there is a neighbourhood U of a such that if $x \in U$ then $\Phi_t(x) \in V$ for all $t \geq 0$.

An equilibrium u is said to be *unstable* if it is not stable. This is equivalent to saying that there exists some $\epsilon > 0$ and a sequence $(x_j)_{j \in \mathbb{N}}$ in D that converges to u such that

$$\sup_{t \geq 0} |\Phi_t(x_j) - u| \geq \epsilon$$

for all j .

Remark 14.5. The fact that Φ_t is a dynamical system means that $\Phi_0 = \text{Id}$ and $\Phi_{s+t} = \Phi_s \circ \Phi_t$. This is effectively the flow of a system. See [Section 11: Dynamical Systems](#) for more details.

Remark 14.6. In order to see why the above definition of instability is equivalent to the negation consider the following argument: if an equilibrium point is not stable then there is some $\epsilon > 0$ such that for all $\delta > 0$ we have some x such that $|x - u| < \delta$ and $\sup_{t \geq 0} |\Phi_t(x_j) - u| \geq \epsilon$. Taking the x_j to be the corresponding x for $\delta = \frac{1}{j}$ we get the desired sequence. The other direction is immediate (even arbitrarily close to u , we can find points that eventually leave a certain threshold thus we cannot have stability).

Definition 14.7. An equilibrium point a is said to be *asymptotically stable* (or *attractive*) if a is stable and there is some $\delta > 0$ such that

$$\lim_{t \rightarrow \infty} |\Phi_t(x) - a| = 0$$

for all $|x - a| < \delta$.

Remark 14.8. The condition that a be stable is necessary. You can construct counterexamples where a is not stable but the second condition still holds.

We have topological conjugacy between a system and its linearisation if hyperbolicity holds (which is to say, if the eigenvalues of the linearised system have non-zero real part). Although we will not prove this statement in general, we will state it and prove some special cases of it.

Theorem 14.9 (Hartman-Grobman) Consider $x' = f(x)$ for $f \in C^1$ and suppose $f(a) = 0$ (i.e. $x(t) \equiv a$ is an equilibrium). Let $A := Df(a)$ be hyperbolic. Let Φ_t be the dynamical system for the ODE and let $\Psi_t := e^{tA}$ be the dynamical system for the linearisation. Then, there is a neighbourhood U of a and neighbourhood V of the origin and a homeomorphism $h : U \rightarrow V$ such that

$$\Phi_t = h^{-1} \circ \Psi_t \circ h$$

provided t is sufficiently small.

What can be learned by applying this theorem? Well for one thing, we can determine the stability or even asymptotic stability of equilibria. This is quite useful when analysing systems in the real world (gross). We also get the so-called *stable/unstable curve theorem*, which roughly says that if the linearisation (of a planar system) has a saddle then so will the nonlinear system. In particular there will be two curves: the solutions starting on one curve will approach the origin as $t \rightarrow \infty$ (known as the stable curve) and solutions starting on the other curve will approach the origin as $t \rightarrow -\infty$ (known as the unstable curve). In [Figure 19](#), the stable curve is the ‘parabola’ seen in the left half of the plane and the unstable curve is the x -axis. This generalises to the stable/unstable manifold theorems for higher dimensional systems.

The theorem also implies that the filter of topological conjugacy is quite broad. For example, all sinks are topologically conjugate to one another (as are all sources). Thus for example, you wouldn’t be able to tell the difference between a spiral sink or a stable node. In some cases that is fine, in some cases it is not. We could perhaps say a bit more if we instead had that h was a diffeomorphism. When can we get diffeomorphic conjugacy then? We need a certain non-resonance condition to hold.

Nevertheless, the Hartman-Grobman theorem is indicative of a more general principle, if the linearised system is ‘structurally stable’ (in this instance, structurally stable means hyperbolic, although this may change in other contexts), then the non-linear system will look like the linear system, at least locally.

We will only prove Hartman-Grobman in the special case when we have n distinct, negative eigenvalues. But first we need the following lemma.

Lemma 14.10 Consider $x' = f(x)$ where $f : \mathbb{R}^n \supset U \rightarrow \mathbb{R}^n$ is C^1 with an equilibrium at a . If $A := Df(a)$ has n distinct, negative (hence real) eigenvalues. Then a is asymptotically stable for the non-linear system (or equivalently, it is a sink).

Proof. Suppose $A = Df(a)$ has n distinct, negative real eigenvalues. Then we claim that we can assume A to be a diagonal matrix without loss of generality. This can be achieved by doing a change of basis if necessary.

Suppose D is the diagonalisation of A . In other words, there exists an invertible T such that $A = TDT^{-1}$. Suppose y is a solution to $y' = Ay$. Then substituting $y = T\tilde{y}$, we find that $(T\tilde{y})' = T\tilde{y}' = AT\tilde{y}$ or in other words

$$\tilde{y}' = D\tilde{y}$$

We can similarly assume $a = 0$, by taking $\tilde{f}(x) = f(x + a)$ and considering $D\tilde{f}$ instead (then of course $D\tilde{f}(0) = Df(a)$).

Using the Taylor expansion, we have

$$f(x) = f(0) + f'(0)x + \bar{o}(|x|)$$

where $\bar{o}(|x|)$ contains the higher order terms for the error (see [Wikipedia](#) on little-o notation. If you want to see what exactly the remaining terms looks like, consider the [following](#) article). Thus the ODE reduces to

$$x' = Ax + \bar{o}(|x|)$$

Now we define a function $L : \mathbb{R}^n \rightarrow \mathbb{R}$ where

$$L(x) = \frac{1}{2} |x|^2$$

Suppose y is a solution to the linearisation. Intuitively, it is quite clear that the origin is asymptotically stable for the linearisation: we have a basis of eigenvectors and along each of these, the solutions decay. Since

the general solution is a linear combination of these, it must decay as well. We will prove the statement formally using the technique of *Lyapunov functions* (to be discussed shortly), a very useful technique that can be used in very general contexts.

We find that

$$\frac{d}{dt}L(y(t)) = \frac{d}{dt}\left(\frac{1}{2}|y(t)|^2\right) = y^T \cdot y' = y^T Ay = \lambda_1 y_1^2 + \dots + \lambda_n y_n^2 \leq \max\{\lambda_1, \dots, \lambda_n\} \cdot |y|^2$$

where $y^T \cdot y'$ refers to the dot product of the two vectors. In particular we find that if we let c denote the above maximum, then

$$\frac{d}{dt}L(y) \leq 2cL(y) \leq cL(y)$$

where the second inequality holds because c is negative. Then a consequence of (an easy case of) [Grönwall](#) is that

$$L(y(t)) \leq L(y(0))e^{ct}$$

Since c is negative, we know all such solutions decay.

Now suppose x is a solution to the non-linear system. Then

$$\begin{aligned}\frac{d}{dt}L(x(t)) &= x^T \cdot x' = x^T (Ax + \bar{o}(|x|)) \\ &\leq c|x|^2 + \bar{o}(|x|^2)\end{aligned}$$

We choose a neighbourhood V of 0 so that for $x \in V$ we have

$$|x|^2 \leq \frac{|c|}{2} = \frac{-c}{2}$$

Then by definition of little-o, we find that for $x(t) \in V$

$$\frac{d}{dt}L(x(t)) \leq c|x(t)|^2 + \bar{o}(|x(t)|^2) \leq c|x(t)|^2 - \frac{c}{2}|x(t)|^2 = \frac{1}{2}c|x(t)|^2 = cL(x)$$

By Grönwall again, what we find is that

$$L(x(t)) \leq L(x(0))e^{ct}$$

Since c is negative, this goes to 0 as $t \rightarrow \infty$. Since $L(x) = \frac{1}{2}|x|^2$, the same holds true for $|x|$ as well. Thus all solutions that start in V , approach the origin as $t \rightarrow \infty$, proving the claim that the origin is asymptotically stable for the non-linear system as well. \square

Of course, the fact that the two flows are sinks does not (immediately) show that they are conjugate. We will do so shortly. However before that, a comment on the techniques used above. As mentioned, one of the key ingredients was the Lyapunov function.

Definition 14.11 (Lyapunov functions). Consider the differential equation $x' = f(x)$ with an equilibrium at a . Then a Lyapunov function is a real-valued, differentiable function on a neighbourhood of a such that it is 0 on a and positive everywhere else and also satisfies

$$\frac{d}{dt}L(x(t)) \leq 0$$

whenever $x(t)$ is a solution to the ODE.

If we have

$$\frac{d}{dt}L(x(t)) < 0$$

for all non-constant solutions, then L is said to be a strict Lyapunov function.

In general it is difficult to find Lyapunov functions, however in certain cases this can be done by considering things like conservation of energy/momentum/etc (certainly any solution to an ODE in physics is going to satisfy the fact that energy is non-increasing with time).

Now we show that in the case of sinks, the non-linear system is indeed conjugate to its linearisation. This is a special case of Hartman-Grobman

Proposition 14.12 *Consider the differential equation $x' = f(x)$ where f is C^1 and a is an equilibrium point. Suppose $A := Df(a)$ has n distinct, negative eigenvalues. Let Φ_t denote the flow of the differential equation. Then there exist open neighbourhoods V and U of the origin and a respectively and a map $h : U \rightarrow V$ such that*

$$\Phi_t = h^{-1} \circ e^{tA} \circ h$$

Proof. As before, without loss of generality, we can assume that A is diagonal and $a = 0$. By the previous lemma, [Lemma 14.10](#), we know there is a $\rho > 0$ such that for all $x \in B_\rho(0)$, we have $\Phi_t(x) \rightarrow 0$ as $t \rightarrow \infty$. We define $B_\rho(0)$ to be our U . Since all solutions in U , tend toward 0, we conclude that the vector field is always pointing inward, (roughly) toward the origin (to be precise what we want to say is that the dot product of the vector field at a point with the point itself is always negative. This is clear from the fact that solutions tend toward the origin but can also be verified rigorously in a very similar manner to the previous lemma. The dot product of a point (x_1, \dots, x_n) with its corresponding vector is exactly $\lambda_1 x_1^2 + \dots + \lambda_n x_n^2 + o(|x|^2)$. By making the error small, we can ensure this quantity is negative.). In particular this means that $-f$ is always pointing outward (away from the origin). Thus for every $x \in U \setminus \{0\}$, there is some (unique) time $t < 0$ such that $|\Phi_t(x)| = \rho$. This means that $\{t < 0 : \Phi_t(x) \notin U\}$ is non-empty for every non-zero x (there is some time in ‘the past’ where the solution was on the boundary hence not in U) and is of course bounded above (by 0 for example). This means that the supremum always exists, allowing us to define the map

$$\begin{aligned} \tau : U \setminus \{0\} &\rightarrow \mathbb{R} \\ x &\mapsto \sup\{t < 0 : \Phi_t(x) \notin U\} \end{aligned}$$

In particular, $\tau(x)$ corresponds to the first time that the solution starting at x entered U . Then we define

$$h(x) = e^{-\tau(x)A} \Phi_{\tau(x)}(x)$$

with $h(0) = 0$. See [Figure 20](#) for an illustration of the map.




Figure 20: Visualisation of h

We want to show that

$$h \circ \Phi_t = e^{tA} \circ h$$

for all $t \geq 0$. We claim that $\tau(\Phi_t(x)) = \tau(x) - t$. Intuitively, this is clear. We know that $\tau(x)$ is the first time that the solution starting at x entered U (again, remembered that this happened in the past so to speak, so $\tau(x)$ is in fact negative). So if we evolve the system (forward) by time t , in order to find the first time we entered U from this new point $\Phi_t(x)$ we must go back time t and then go further back $\tau(x)$. Thus $\tau(\Phi_t(x)) = -t + \tau(x) = \tau(x) - t$. Luckily, this is one of the (painfully few) times in math when intuition can be converted almost exactly into formulaic manipulations.

$$\tau(\Phi_t(x)) = \sup\{s < 0 : \Phi_s(\Phi_t(x)) \notin U\} = \sup\{s < 0 : \Phi_{s+t}(x) \notin U\} = \sup\{s < 0 : \Phi_s(x) \notin U\} - t = \tau(x) - t$$

We then compute that

$$\begin{aligned} h(\Phi_t(x)) &= e^{-\tau(\Phi_t(x))A} \circ \Phi_{\tau(\Phi_t(x))}(\Phi_t(x)) \\ &= e^{-(\tau(x)-t)A} \Phi_{\tau(x)-t}(\Phi_t(x)) \\ &= e^{tA} \circ e^{-\tau(x)A} \Phi_{\tau(x)}(x) \\ &= e^{tA} \circ h(x) \end{aligned}$$

The fact that h is invertible is clear; we could swap around things appropriately in the construction above. All that remains to check is that it is continuous.

We define the function

$$f(t, x) = |\Phi_t(x)| - r_0$$

Note this function is 0 exactly when the solution starting at x passed through the ball of radius r_0 (some time in the past). We know this function is C^1 if we ignore the origin. Moreover,

$$\frac{\partial f}{\partial t} = \frac{\partial}{\partial t}(\sqrt{x_1(t)^2 + \dots + x_n(t)^2}) = \frac{1}{|x(t)|} \begin{pmatrix} x_1(t) \\ \vdots \\ x_n(t) \end{pmatrix} \cdot \begin{pmatrix} x'_1(t) \\ \vdots \\ x'_n(t) \end{pmatrix}$$

The dot product is always non-zero (the vector field always points inward toward the origin/solutions intersect the boundary of the ball transversely). Since $\frac{\partial f}{\partial t} \neq 0$, we can use the Implicit Function Theorem to find the x that solve $f(t, x) = 0$ in terms of t and conclude that this map is C^1 . This is exactly the map τ .

This shows continuity (and even differentiability) of h everywhere besides the origin. So let's check continuity at the origin as well. Recall we have $h(x) = e^{-\tau(x)A} \Phi_{\tau(x)}(x)$. By definition of $\tau(x)$, we have that $|\Phi_{\tau(x)}(x)| = \rho$. Convince yourself that $\lim_{x \rightarrow 0} \tau(x) = -\infty$. This means that $h(x)$ always corresponds with taking a point on the boundary of U and evolving it forward in time. The closer we start to the origin, the more time we evolve for, so the closer we end up to the origin. Thus we have continuity at the origin as well. In symbols, we have

$$\lim_{x \rightarrow 0} |h(x)| \leq \lim_{x \rightarrow 0} |e^{-\tau(x)A}| |\Phi_{\tau(x)}(x)| = \rho \lim_{x \rightarrow 0} |e^{-\tau(x)A}| = \rho \lim_{t \rightarrow \infty} |e^{tA}| = 0$$

□

Although we assumed that the eigenvalues were real in the above proof, them having negative real part is sufficient. (The only thing we need to check is that the system is still a sink and this can be done by considering the same Lyapunov function).

We also assumed that A was diagonalisable, but even this is not necessary. Suppose A is a Jordan block like so

$$A = \begin{pmatrix} \lambda & 1 \\ 0 & \lambda \end{pmatrix}$$

We claim that A is conjugate to the matrix

$$\tilde{A} := \begin{pmatrix} \lambda & \epsilon \\ 0 & \lambda \end{pmatrix}$$

for all $\epsilon > 0$ (this is sometimes called the ϵ -Jordan form). This is easily verified by taking

$$T = \begin{pmatrix} 1 & 0 \\ 0 & \epsilon \end{pmatrix}$$

and computing $T^{-1}AT$. Suppose y is a solution to $y' = \tilde{A}y$. Then

$$\frac{d}{dt} \left(\frac{1}{2} |y|^2 \right) = \lambda(y_1^2 + y_2^2) + \epsilon y_1 y_2$$

Note that we have

$$0 \geq -\frac{\epsilon}{2}(y_1 + y_2)^2 = -\frac{\epsilon}{2}(y_1^2 + y_2^2) - \epsilon y_1 y_2$$

Now if we have $\lambda < -\frac{\epsilon}{2}$ we would have

$$\epsilon ab \geq -\frac{\epsilon}{2}(y_1^2 + y_2^2) > \lambda(y_1^2 + y_2^2)$$

Thus by taking $0 < \epsilon < -2\lambda$, we can ensure that $\frac{d}{dt}(\frac{1}{2}|y|^2) < 0$ on all non-equilibria solutions, allowing us to use this as a Lyapunov function (as before) and then proceeding with the proof as before. In this manner, we can show that [Hartman-Grobman](#) holds in the case of Jordan blocks as well (provided that all eigenvalues are real and negative).

Although we won't prove the Hartman-Grobman theorem in general, we briefly outline the strategy for doing so. The generalisation of the above to $n \times n$ Jordan blocks is clear (we replace all the off diagonal 1's with ϵ). We can construct very similar arguments for sources by going back in time. For saddles we can decompose our space into a direct sum of sources and sinks (that is we write our total space as the sum of subspaces where solutions approach the equilibria as $t \rightarrow \infty$ and solutions that approach the equilibria as $t \rightarrow -\infty$) and argue on these subspaces separately.

The final idea leads us to the following two definitions.

Definition 14.13 (Stable/Unstable manifold). Given an equilibrium a , the stable manifold $W^s(a)$ as

$$W^s(a) := \left\{ x \in \mathbb{R}^n : \lim_{t \rightarrow \infty} \Phi_t(x) = a \right\}$$

Similarly, the unstable manifold $W^u(a)$ is defined to be

$$W^u(a) := \left\{ x \in \mathbb{R}^n : \lim_{t \rightarrow -\infty} \Phi_t(x) = a \right\}$$

The fact that these sets are manifolds is in fact a consequence of the [Stable manifold theorem](#). A property that readily follows from the definitions is that the two manifolds are positively invariant (that is if x is an element of either the stable or the unstable manifold then $\Phi_t(x)$ is an element of the same manifold for all $t \geq 0$). Moreover, these manifolds are tangent at $x = a$.

15 Techniques for non-linear systems

We have discussed plenty about how to classify flows and such. So we can know (in certain cases) how the flow of a certain system behaves, up to homeomorphism. Unfortunately, when dealing with systems in the real world (gross) one would like to know how a given system behaves without having to construct homeomorphisms and such (indeed the homeomorphism defined in [Proposition 14.12](#) required one to know the flow of the non-linear system which rather defeats the purpose). So we briefly discuss some techniques for determining qualitative/rough behaviour of non-linear systems.

15.1 Nullclines

Nullclines are perhaps the simplest idea one may think of when trying to determine the qualitative behaviour of some given system. Namely, we find when the different components of the solutions will have derivative 0. This allows us to partition the space into subdomains where we know the signs of the components of the solution. This gives a rough idea as to how initial points in these domains behave and with a bit of creativity and guesswork, one can stitch these together to sketch out the phase portrait of the system.

Definition 15.1 (Nullcline). Given a system $x' = f(x)$ on $D \subset \mathbb{R}^n$, the j -th nullcline is the set

$$\{x \in D : f_j(x) = 0\}$$

15.1.1 Example 1

Consider the system

$$\begin{aligned} x' &= x(a - by) \\ y' &= y(-c + dx) \end{aligned}$$

Remark 15.2. This is the famous **Lotka-Volterra system**

We see that the x -nullcline is at $x = 0$ and $y = \frac{a}{b}$. Similarly the y -nullcline is at $y = 0$ and $x = \frac{c}{d}$. This splits the upper right quadrant (which is our domain if we are indeed thinking of this as the Lotka-Volterra system). Then for $x < \frac{c}{d}$ we have $y' < 0$ and for $x > \frac{c}{d}$ we have $y' > 0$. Similarly for $y < \frac{a}{b}$ we see $x' > 0$ and for $y > \frac{a}{b}$ we have $x' < 0$. Plotting this information we conclude, the system evolves at least somewhat circularly. We don't immediately see that the solutions are periodic (as we know them to be), which already tells us that nullclines aren't *too* informative, but something is better than nothing so one can hardly complain.




Figure 21: Nullclines for Lotka-Volterra system

15.2 Lyapunov functions

We have already defined Lyapunov functions in the previous section. Here we show why they are so useful. They allow us to find stable and asymptotically stable equilibria (especially useful if solving for them directly is difficult).

Theorem 15.3 (Lyapunov's Theorem) *Let x^* is an equilibrium point for $x' = f(x)$. Suppose $L : D \rightarrow \mathbb{R}$ (where $D \subset \mathbb{R}^n$, as usual, is open) is a differentiable function such that $L(x^*) = 0$ and $L(x) > 0$ for all $x \in D \setminus \{x^*\}$. If $DL(x)(f(x)) \leq 0$ for all $x \in D \setminus \{x^*\}$ then x^* is stable and if the inequality is strict, then x^* is asymptotically stable.*

Remark 15.4. We don't strictly need the fact that $L(x^*) = 0$ and $L(x) > 0$ for $x \neq x^*$; we only need x^* to be a strict local minimum. However, we can always assume this to be the case by shifting L appropriately.

Remark 15.5. As a fun little fact, $DL(x)(f(x))$ is exactly the push forward of the tangent vector $f(x)$ at x via L . Perhaps someday I will be able to tell you whether this should be painfully obvious or not.

Proof. We first show stability. That is given a neighbourhood U of x^* , we want to find a neighbourhood V of x^* such that $x \in V$ implies that $\Phi_t(x) \in U$ for all $t \geq 0$.

Without loss of generality, we can assume that U is bounded and that x^* is the only critical point in U (if either of these does not hold, we can always consider a smaller subset where this does hold and since the solutions will always lie in this subset, they will in particular lie in the larger one). Let $m := \min\{L(x) : x \in \partial U\}$. Since x^* is the (only) local minimum, we must have that $L(x^*) < m$. Then $x^* \in V := \{x \in D : L(x) < m\}$. Then if $x(0) \in V$, then $L(x(t)) < m$ for all $t > 0$ since L is non-increasing along solutions (note $(L \circ x)' = DL(x) \circ x' = DL(x)(f(x))$ giving us the non-increasing property by hypothesis of the theorem).

Suppose now that the inequality is strict. In other words, $DL(x)(f(x)) < 0$ for all $x \in D \setminus \{x^*\}$. We wish to show that x^* is asymptotically stable. We know that $L(x(t))$ is a decreasing function for every (non-constant) solution. Thus $\lim_{t \rightarrow \infty} L(x(t))$ is bounded below by $L(x^*)$ implying that the limit exists and is greater than or equal to $L(x^*)$.

We choose some $x_0 \in V$ and let $x(t)$ be the solution to the differential equation satisfying $x(0) = x_0$. Since \bar{V} is compact, we can find an increasing, unbounded sequence t_j such that $\lim_{j \rightarrow \infty} x(t_j) = y_0$ for some $y_0 \in \bar{V}$. Clearly we must have that $L(x(t)) > L(y_0)$ for all t . We will show that this leads to a contradiction. The idea is that L will decrease along the solution starting at y_0 . Continuous dependence implies that L of solutions that begin close to y_0 will also be bounded above by $L(y_0)$. But since $x(t_j)$ lie close to y_0 (for sufficiently large j), we must eventually get $L(x(t_j)) < L(y_0)$ leading to the desired contradiction. Now we write this formally.

Suppose $y_0 \neq x^*$ (if they are equal, we are done). Let $y(t)$ be a solution to the differential equation satisfying $y(0) = y_0$. We know then that for all $t > 0$, we get $L(y(t)) < L(y_0)$ (this is where we use $y_0 \neq x^*$). We fix some $\tau > 0$. Then of course $L(y(\tau)) < L(y_0)$ implying that there is a neighbourhood W_τ of $y(\tau)$ such that for all $w \in W_\tau$, we have $L(w) < L(y_0)$ (we can take $L^{-1}((-\infty, L(y_0)))$ for example).

Continuous dependence of solutions implies that solutions with similar initial conditions will evolve similarly. Thus we can find a neighbourhood W_0 of y_0 such that $\Phi_\tau(w) \in W_\tau$ for all $w \in W_0$. We can find W_0 using [Theorem 13.2](#) for instance which gives a bound for how far solutions can differ as a function of how far the initial conditions are. Since W_0 is a neighbourhood of y_0 , we know there is some $x(t_j) \in W_0$ (by definition of convergence). Therefore $\Phi_\tau(x(t_j)) = x(t_j + \tau) \in W_\tau$. But this means that $L(x(t_j + \tau)) < L(y_0)$ by construction of W_τ , leading to the desired contradiction since we had $L(x(t)) > L(y_0)$ for all t . \square

In fact the above statement, has a converse. Namely if x^* is a sink for a system $x' = f(x)$, then there exists a Lyapunov function $L(x)$ on a neighbourhood of x^* that satisfies all the above hypotheses which would allow us to conclude that x^* is a sink (i.e. we have x^* is a strict local minimum, derivative along non-constant solutions is negative, etc.).

Sometimes one finds that the Lyapunov function is actually constant along solutions. In this case, one finds solutions to the differential equation very easily: they are simply the level sets of the Lyapunov function. This occurs, for example, with the Lotka-Volterra system if we take

$$L(x, y) = d \cdot x - c \log x + b \cdot y - a \log y$$

15.3 Gradient Systems

As mentioned in general it is rather difficult to find Lyapunov functions. However one instance, where the Lyapunov function is particularly nice and easy to find is with gradient systems.

Definition 15.6 (Gradient System). A gradient system is a differential equation of the form $x' = -\nabla V(x)$ where $V : \mathbb{R}^n \supset D \rightarrow \mathbb{R}$ is a C^2 function.

In this case we can take $L(x) = V(x)$ since

$$\frac{d}{dt} V(x(t)) = DV(x) \cdot x' = \nabla V \cdot (-\nabla V) = -|\nabla V|^2$$

Of course this is 0 exactly on the equilibrium and negative everywhere else, as we would expect.

With gradient systems, one often imagines the graph of V as forming a surface (say if $D \subset \mathbb{R}^2$) and solutions to $x' = -\nabla V(x)$ correspond to travelling down this surface according to gravity.

15.4 Hamiltonian Systems

Hamiltonian systems are quite analogous to gradient systems, at least in terms of definitions. Although they are a bit more general, we will restrict our attention to Hamiltonian systems on \mathbb{R}^2 .

Definition 15.7 (Hamiltonian systems). A Hamiltonian system is a differential equation of the form

$$x'' = -V'(x)$$

where $V : D \subset \mathbb{R} \rightarrow \mathbb{R}$ is a C^2 function.

Equivalently a system $z' = f(z)$ is Hamiltonian if there exists a map $H : C \subset \mathbb{R}^2 \rightarrow \mathbb{R}$ such that

$$x' = \frac{\partial H}{\partial y}(x, y), y' = -\frac{\partial H}{\partial x}(x, y)$$

Remark 15.8. The equivalence is not obvious. For one thing, we need to write the second-order differential equation in the first definition as a system of first-order differential equations giving us

$$\begin{cases} x' &= y \\ y' &= -V'(x) \end{cases}$$

Given such a V , we can define the Hamiltonian H as

$$H(x, y) = \frac{1}{2}y^2 + V(x)$$

Then we have

$$\partial_x H = V'(x), \partial_y H = y$$

The other direction remains a mystery to me (for now?).

These are another class of systems where the Lyapunov function is easy to find. Namely we define

$$L(x, y) = \frac{1}{2}|y|^2 + V(x)$$

(where V is as in the first definition). Then if $(x(t), y(t))$ is a solution, we get

$$\frac{d}{dt} L(x(t), y(t)) = yy' + V'(x)x' = 0$$

Thus we find that L is actually constant along solutions. Since L is continuous, if solutions get closer and closer together with time then they must have the same constant. To be precise, we have the following statement.

Suppose Γ is a closed orbit that solves the equation and suppose $y \in D$ is such that $d(\Phi_t(y), \Gamma) \rightarrow 0$ as $t \rightarrow \infty$. Suppose $L|_{\Gamma} \equiv c_{\Gamma}$ and $L|_{\beta} \equiv c_{\beta}$ where β is the orbit (aka image) of y . Then continuity of L implies that $c_{\Gamma} = c_{\beta}$.

15.5 Wronskian

The Wronskian is another tool that can be used to compare solutions of a differential equation (or even compare solutions of similar differential equations).

Definition 15.9 (Wronskian). Given two differentiable functions f, g we define their Wronskian $W(f, g)$ to be a function where

$$W(f, g)(x) = f(x)g'(x) - f'(x)g(x) = \det \begin{pmatrix} f(x) & g(x) \\ f'(x) & g'(x) \end{pmatrix}$$

In order to see how useful this tool is, consider the equation

$$u''(x) + p(x)u'(x) + q(x)u(x) = 0 \quad (15.1)$$

where $p(x)$ and $q(x)$ are some continuous functions. Clearly if f, g are linearly dependent functions then $W(f, g) = 0$. However the converse also holds to some extent: if f, g are linearly independent solutions to (15.1) then $W(f, g)$ is never 0.

Lemma 15.10 Suppose f, g are linearly independent solutions to (15.1) then $W(f, g)$ is never 0.

Proof. If we write (15.1) as a linear system of first order equations, we get

$$\begin{pmatrix} u(x) \\ v(x) \end{pmatrix}' = \begin{pmatrix} 0 & 1 \\ -q(x) & -p(x) \end{pmatrix} \begin{pmatrix} u(x) \\ v(x) \end{pmatrix}$$

If f and g are linearly independent solutions to (15.1), then they are linearly independent solutions to the first-order system above as well (to be precise $u(x) = f(x)$, $v(x) = f'(x)$ and $u(x) = g(x)$, $v(x) = g'(x)$ form the pair of linearly independent solutions). This exactly means that

$$\begin{pmatrix} f(x) \\ f'(x) \end{pmatrix}, \begin{pmatrix} g(x) \\ g'(x) \end{pmatrix}$$

are linearly independent for all x which in turn is equivalent to saying that the determinant above is never zero. \square

This fact allows for an easy proof of the following theorems.

Theorem 15.11 (Sturm Separation Theorem) Suppose f, g are linearly independent solutions to (15.1). Then f must vanish between any two successive zeroes of g . In other words the zeroes of f and g occur alternately.

Proof. Let a, b be 2 successive zeroes of g with $a < b$. Suppose $g > 0$ on (a, b) . Then $g'(a) > 0$ and $g'(b) < 0$ (in principal we only have $g'(a) \geq 0$ rather than have the strict inequality. However if $g'(a) = 0$ then $W(f, g)(a) = 0$ which we know cannot happen. The same is true for $g'(b)$). Note that $W(f, g)(a) = f(a)g'(a)$ and $W(f, g)(b) = f(b)g'(b)$. Since $W(f, g)$ never vanishes, it must maintain its sign. Since $g'(a)$ and $g'(b)$ have opposing signs, the same must be true for $f(a)$ and $f(b)$. Then the Intermediate Value Theorem implies that f is 0 somewhere on (a, b) as desired. A symmetric argument holds if $g < 0$ on (a, b) . \square

Remark 15.12. By noting that $\sin(kx)$ and $\cos(kx)$ are linearly independent solutions to $u'' + k^2u = 0$, we find that their zeroes alternate (as we know them to).

Theorem 15.13 (Sturm Comparison Theorem) Let $f(x)$ and $g(x)$ be non-trivial solutions to $u'' + p(x)u = 0$ and $v'' + q(x)v = 0$ where $p(x) \geq q(x)$, respectively. Then $f(x)$ vanishes at least once between successive zeroes of g , unless $p(x) = q(x)$ and f is a constant multiple of g .

Proof. Let $a < b$ be two successive zeroes of g like before and suppose f does not vanish on (a, b) . By replacing f and/or g with their negatives (those also satisfy the same ODE's), we can assume without loss of generality that f, g are positive on (a, b) . As before, this means that $g'(a)$ is non-negative and $g'(b)$ is non-positive (because f and g no longer solve the same differential equation, we cannot use the same argument to conclude that the inequalities are strict). Thus what we find is that $W(f, g)(a) = f(a)g'(a) \geq 0$ and $W(f, g)(b) = f(b)g'(b) \leq 0$.

Then we compute that

$$\begin{aligned}\frac{d}{dx} W(f, g)(x) &= f'g' + fg'' - g'f' - gf'' \\ &= f(-qg) - (-pf)g \\ &= (p - q)fg\end{aligned}$$

If $p \neq q$, then the above is positive on some neighbourhood so $W(f, g)$ is increasing on that neighbourhood (and remain non-decreasing outside it). But this contradicts $W(f, g)(b) \leq 0$ (since $W(f, g)(a) \geq 0$). Since we have covered the case with $p = q$ in the previous theorem, we are done. \square

16 Limit Sets

Limit sets are useful for understanding the ultimate behaviour of a given solution and thereby gaining insight into the limiting behaviour of the flow.

Definition 16.1 (Limit sets). Given a solution $x(t)$ to the differential equation $x' = f(x)$ with $f : D \rightarrow \mathbb{R}^n$ (with $D \subset \mathbb{R}^n$ open), we define the ω -limit set of x as

$$\omega(x) = \{y \in D : \text{exists an increasing sequence } t_j \text{ to } \infty \text{ such that } \lim_{j \rightarrow \infty} x(t_j) = y\}$$

Equivalently we can write

$$\omega(x) = \bigcap_{s>0} \overline{\{x(t) : t \geq s\}}$$

Analogously, we define the α -limit set of x as

$$\alpha(x) = \{y \in D : \text{exists a decreasing sequence } t_j \text{ to } -\infty \text{ such that } \lim_{j \rightarrow \infty} x(t_j) = y\}$$

or

$$\alpha(x) = \bigcap_{s<0} \overline{\{x(t) : t \leq s\}}$$

Proposition 16.2 The two definitions of $\omega(x)$ (and $\alpha(x)$) are equivalent.

Proof. Clearly if there is an increasing sequence t_j such that $x(t_j) \rightarrow y$ then y is going to be in the closure of all $\{x(t) : t \geq s\}$ (one could take this to be the definition of closure (in metric spaces)). Thus we wish to show inclusion in the other direction. This is also fairly easy to do. For every $j \in \mathbb{N}$, we find t_j such that $t_j \geq j$ and $|y - x(t_j)| < \frac{1}{j}$ (both of these are possible by definition of the sets and the definition of closure). Then by construction $x(t_j) \rightarrow y$. \square

Remark 16.3. For the second definition of $\omega(x)$, we could equivalently take $s \in \mathbb{R}$ or $s \in \mathbb{N}$ instead of $s > 0$; all we need is a set that is unbounded above. The same holds true for the second definition of $\alpha(x)$ (i.e. we could have $s \in \mathbb{R}$ or $s \in -\mathbb{N}$ or any set that is unbounded below instead of $s < 0$).

Note that limit sets could be empty. For example if $x' = 1$, then we know solutions are of the form $x(t) = t + c$ and clearly both limit sets will be empty in this case. This is in fact typical of unbounded solutions or solutions that leave D . However, if $x(t)$ is a bounded solution that remains in D for all $t > 0$, then $\omega(x)$ will be non-empty. This is easily seen by considering a compact set that contains $x(\mathbb{R}^+)$ and verifying that the **finite intersection property** holds for the family of sets in the (second) definition of $\omega(x)$. An entirely analogous argument holds for $\alpha(x)$. For the remainder of the section, you may assume that all limit sets are bounded, even when not explicitly stated.

16.1 Examples

Consider Figure 22. In this case we have $\omega(x) = \{p\}$ and $\alpha(x) = \{q\}$. The situation is reversed for $y(t)$, namely we have $\omega(y) = \{q\}$ and $\alpha(y) = \{p\}$. Such solutions that connect different equilibrium points to one another are called heteroclinic solutions. Solutions that connect the same equilibria are called homoclinic.

We see that $z(t)$ has a closed orbit so, in fact, we get $\omega(z) = \alpha(z) = \{z(t) : t \in \mathbb{R}\}$. Finally we have $r(t)$ with $\omega(r) = \{z(t) : t \in \mathbb{R}\}$ and $\alpha(r) = \{a\}$. Finally Figure 23 highlights another case that may occur with




Figure 22: Example of limit sets, assume p, q, a are equilibria

$\omega(w) = \{x(t) : t \in \mathbb{R}\} \cup \{y(t) : t \in \mathbb{R}\} \cup \{p, q\}$ and $\alpha(w) = \{a\}$. Roughly speaking, these examples illustrate what all bounded limit sets in the plane look like.




Figure 23: Example of limit sets, assume p, q, a are equilibria

Limit sets have some nice properties as well.

Proposition 16.4 Suppose $x(t)$ be a bounded solution to $x' = f(x)$ and let $\omega(x)$ denote the ω -limit set of x . Then $\omega(x)$ is always closed and if $y \in \omega(x)$, then $y(\mathbb{R}) \subset \omega(x)$. Moreover $\omega(x)$ is connected.

Proof. The closure of limit sets is immediate by noting that they are defined by the intersection of closed sets. But we can show this result using the other definition as well. Let us do this as well, if for not other reason because we can.

We show $\omega(x)$ is closed by showing it contains its limits points (where by limit points, we mean the topological notion). Let y_n be a sequence in $\omega(x)$ that converges to some y . We want to show that $y \in \omega(x)$. The procedure will be exactly what you expect: since the y_n converge to y , they come arbitrarily close to y and since each y_n is a ω limit point, the solution through x also gets arbitrarily close to them. By combining these two facts (perhaps you can already see the triangle inequality lurking around the corner), we can find a sequence of points in x that converges to y . All that remains is to translate this into precise mathematics.

By definition of limit sets, for every n , there exists a sequence $(t_{k,n})_{k \in \mathbb{N}}$ such that

$$\lim_{k \rightarrow \infty} x(t_{k,n}) = y_n$$

For each $n \in \mathbb{N}$, define l_n to be such that $|y_n - x(t_{k,n})| < \frac{1}{n}$ for all $k \geq l_n$. We then define $s_n = t_{l_n, n}$. In particular this means that

$$|y_n - x(s_n)| < \frac{1}{n}$$

Let $\epsilon > 0$ be given. We find n_0 such that for all $n \geq n_0$ we have $|y - y_n| < \frac{\epsilon}{2}$. Then for $n \geq n_0$, we get

$$|y - x(s_n)| \leq |y - y_n| + |y_n - x(s_n)| < \frac{\epsilon}{2} + \frac{1}{n}$$

Thus by taking n to be larger than n_0 and $\frac{2}{\epsilon}$, we get $|y - x(s_n)| < \epsilon$.

Suppose $y \in \omega(x)$ and let $t \in \mathbb{R}$ be arbitrary. Then there exists an increasing t_j to infinity such that $\lim_{j \rightarrow \infty} x(t_j) = y$. Then

$$\Phi_t(y) = \lim_{j \rightarrow \infty} \Phi_t(x(t_j)) = \lim_{j \rightarrow \infty} x(t_j + t)$$

where the final term is an element of $\omega(x)$ by definition.

We can prove this using the second definition of $\omega(x)$ as well.

$$\begin{aligned} \Phi_t(\omega(x)) &= \Phi_t\left(\bigcap_{s>0} \overline{\{x(r) : r \geq s\}}\right) \\ &\subset \bigcap_{s>0} \Phi_t(\overline{\{x(r) : r \geq s\}}) \\ &\subset \bigcap_{s>0} \overline{\Phi_t(\{x(r) : r \geq s\})} \\ &\subset \bigcap_{s>0} \overline{\{x(t+r) : r \geq s\}} \\ &= \omega(x) \end{aligned}$$

Recall that a closed set $C \subset \mathbb{R}^n$ is connected if for any partition $C = C_1 \cup C_2$ where C_1, C_2 are closed and disjoint, we have that $C_1 = \emptyset$ or $C_2 = \emptyset$.

Now suppose $\omega(x)$ is not connected. Let C_1, C_2 form a partition of $\omega(x)$. In other words C_1, C_2 are both disjoint, non-empty closed sets whose union is equal to $\omega(x)$. Let $y_1 \in C_1$ and $y_2 \in C_2$. Note that C_1 and C_2 are bounded as well and are thus compact. Thus there exists $m > 0$ such that $|z_1 - z_2| \geq m$ for all $z_1, z_2 \in C_1, C_2$ respectively. By definition of $\omega(x)$, we know there exist increasing sequences s_j and t_j such that $\lim_{j \rightarrow \infty} x(s_j) = y_1$ and $\lim_{j \rightarrow \infty} x(t_j) = y_2$. Without loss of generality we can assume that $s_j < t_j$ for all j . One way to see this is by considering the construction of the sequences in [Proposition 16.2](#). We first construct the s_j , then when constructing the second sequence, we just always ensure to choose t_j to be larger than s_j .

Consider the map on $[s_j, t_j]$ (for every j) where we map r in this open interval to $d(x(r), C_1)$ (note that this distance function is well-defined since C_1 is closed and even compact). We see that $d(x(s_j), C_1) = 0$ and $d(x(t_j), C_1) \geq m$. Thus by the Intermediate Value Theorem, there exists some $r_j \in (s_j, t_j)$ such that $d(x(r_j), C_1) \geq \frac{m}{2}$. Similarly, we can argue that the r_j are such that $d(x(r_j), C_2) \geq \frac{m}{2}$ (in fact there needs to be a slight argument for why we can pick the same r_j to satisfy both inequalities. It's not a terribly difficult argument but is besides the point. Quite frankly all we need is that $x(r_j)$ is some positive distance from both C_1 and C_2 and clearly this can be done). We now have another increasing sequence r_j so we must have $p := \lim_{j \rightarrow \infty} x(r_j) \in \omega(x)$ (if the limit does not exist, we can always pass to a convergent subsequence by sequential compactness). But p cannot be in either of C_1 or C_2 since it is a positive distance away from both (follows from the fact that the map defined above $r \mapsto d(x(r), C_1)$ is actually continuous). This contradicts $\omega(x) = C_1 \cup C_2$. \square

Remark 16.5. The above properties are also held by α -limit sets. This can be verified by tweaking the proofs where appropriate (in the obvious manner) or realising that $\alpha(x) = \omega(y)$ where $y = x(-t)$ (i.e. going back in time) is a solution to $y' = -f(y)$.

16.2 Poincaré-Bendixson Theorem

In general, limit sets can be pretty messy, even bounded ones. However on the plane, bounded limit sets are (relatively) well-behaved. In fact, we have the Poincaré-Bendixson (not a typo) theorem, which roughly says that $\omega(x)$ will either contain equilibrium points or will be a periodic solution to the differential equation. Before we prove this statement (or even state it precisely), a definition.

Definition 16.6. A section γ for the flow of $x' = f(x)$ is a curve segment (i.e. a continuous map $\mathbb{R} \rightarrow \mathbb{R}^2$) such that it intersects the vector field transversally. That is $\det(\gamma'(s), f(\gamma(s))) \neq 0$ (where each entry is the column in the matrix).

Note that $\gamma'(s)$ tells us the “direction” of the section at $\gamma(s)$ (the tangent vector at $\gamma(s)$ if you want) and $f(\gamma(s))$ is the vector field at $\gamma(s)$ so if the determinant is non-zero, we know the two vectors are linearly independent, giving the transversal intersection of γ with the vector field. In fact by reversing the parameterisation of γ if necessary, you can always assume the determinant to be positive.

Theorem 16.7 (Poincaré-Bendixson) *Let $f : \mathbb{R}^2 \supset D \rightarrow \mathbb{R}^2$ be C^1 and $x(t)$ a solution to $x' = f(x)$ that is contained entirely in D . Suppose that $\omega(x)$ contains no equilibrium points and is contained in a compact set. Then $\omega(x)$ is a periodic solution to the ODE. In other words, there is some $\tau > 0$, such that $y(t) = y(t + \tau)$ for all t . Moreover, we have that $\omega(x) = \{y(t) : 0 \leq t \leq \tau\}$.*

Proof. Let $y \in \omega(x)$. Since $\omega(x)$ does not contain any equilibria we know that $f(y) \neq 0$. Then by continuity of f , we can find a small section through y on which f is non-zero (in fact if we take γ to be small enough then the vector field is going to be roughly parallel on the section). The first part of this proof is roughly about understanding how solutions of the ODE behave, where the section γ will be a very useful tool for doing so. To be a bit more precise, we will see that if a solution intersects a section multiple times, then it must do so in a very particular way. We will then find a section which y will intersect repeatedly but then argue that this is always through the same point (curiously, it's only this last point that will require the fact that $y \in \omega(x)$). This will allow us to conclude that the orbit of y is periodic and contained in $\omega(x)$. Finally we will show that this inclusion is in fact an equality.

Now suppose $\tilde{x}(t)$ is any solution to the ODE that intersects the section γ at times t_1, t_2, t_3, \dots with $t_i < t_{i+1}$. That is there are some s_1, s_2, s_3, \dots such that $\tilde{x}(t_1) = \gamma(s_1), \tilde{x}(t_2) = \gamma(s_2), \dots$. Then we claim that $s_i < s_{i+1}$ for all i or $s_i > s_{i+1}$ for all i (in other words either $s_1 < s_2 < s_3 < \dots$ or $s_1 > s_2 > s_3 > \dots$). In order to verify this consider [Figure 24](#).




Figure 24: Solutions are trapped inside C after entering it

Let C be the region bounded by $\{\tilde{x}(t) : t \in [t_1, t_2]\} \cup S$ where S is the portion of the section that is between $\tilde{x}(t_1)$ and $\tilde{x}(t_2)$ (in other words $S = \{\gamma(s) : s \in [s_2, s_1]\}$ assuming $s_2 < s_1$ as in the figure). It is clear that the region C is positively invariant: solutions cannot exit by crossing the solution curve x since this would violate uniqueness and they cannot exit via S since the vector field is pointing into C (this also highlights why it's important that we're in the plane, clearly the argument above breaks down for higher dimensions). Thus if $\tilde{x}(t)$ intersects γ again at $\gamma(s_3)$, it must do so inside C implying that $s_3 < s_2$. There remain 3 other cases to consider (1 other if we ignore mirror images) and the arguments for all other cases are near identical.

We now have some understanding of how solutions behave if they intersect a section multiple times. We are interested in the solution that passes through y . So in order to study it further, let us try and find a section that this solution (the one through y) will intersect repeatedly.




Figure 25: Other cases for showing monotonicity. Note in the second two cases it is the ‘complement’ of C that is positively invariant.

Let $z \in \omega(y)$ and let $\tilde{\gamma}$ be a section through z (note how we are using sections and limit sets to understand the behaviour of the solution through y . Hopefully this serves as motivation for why these are useful tools). We define a map $h : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ by $h(s, t) = \Phi_t(\tilde{\gamma}(s))$. In other words, we use the first coordinate to say where on the section we are and the second coordinate to say how far forward (or backward) we evolve in time. Note that

$$\frac{\partial h}{\partial s} = \Phi'_t(\tilde{\gamma}(s))\tilde{\gamma}'(s), \quad \frac{\partial h}{\partial t} = f(\Phi_t(\tilde{\gamma}(s)))$$

This means that

$$Dh(0, 0) = \begin{pmatrix} \tilde{\gamma}'(0) & f(\tilde{\gamma}(0)) \end{pmatrix}$$

which, by construction of $\tilde{\gamma}$, is invertible. Thus we can apply the Inverse Function Theorem to conclude that h is at least locally a homeomorphism (and in fact this holds for every point $(s, 0)$ since Φ_0 is always the identity).

Let V be the neighbourhood of z upon which h acts homeomorphically. Since $y(t_n) \rightarrow z$, we can easily pick t_n such that $y(t_n) \in V$. By applying the inverse of h , we find there is some (s, t) near the origin such that $h(s, t) = \Phi_t(\tilde{\gamma}(s)) = y(t_n)$. This means that $\Phi_{-t}(y(t_n)) = \Phi_{-t}(\Phi_t(\tilde{\gamma}(s))) = \tilde{\gamma}(s)$ which of course lies on the section. Since $\Phi_{-t}(y(t_n)) = y(t_n - t)$ we see that the solution through y intersects $\tilde{\gamma}$ (geometrically what we’re doing is finding a neighbourhood of the section where flow is essentially parallel. We find $y(t_n)$ is in this neighbourhood and then flow it forward or backward until it reaches the section). By repeating this for all other $y(t_n)$ which lie in V , of which there are infinitely many, we see that $y(t)$ intersects this section infinitely many times. All that remains to show is that all of these points are in fact the same.¹⁰

The claim is as follows: if a solution through y where $y \in \omega(x)$ intersects a section multiple times, then it must do so at the same point each time. Suppose not. In other words, suppose y_1 and y_2 are distinct points on the solution through y that also lie on the section $\tilde{\gamma}$. Like in the previous paragraph, we find neighbourhoods V_1, V_2 of y_1 and y_2 respectively where the homeomorphisms can be found and without loss of generality we can assume that these are disjoint (simply consider smaller neighbourhoods if not). The intersections of V_1 and V_2 with the section $\tilde{\gamma}$ are going to form sections of their own. By using the homeomorphisms as in the previous case if necessary, we can assume that $x(t_n)$ intersect $\tilde{\gamma}$ for all (sufficiently large) n . Thus x must intersect the entire section $\tilde{\gamma}$ as well as the smaller, disjoint sections monotonically in all cases. Clearly this cannot be done, see Figure 26.

Thus all the intersections pass through the same point which immediately implies (by uniqueness) that the solution through y is periodic. By Proposition 16.4, we know that the orbit of y is contained in $\omega(x)$. What remains to show is that this is in fact the entirety of $\omega(x)$. In fact there is a slightly more general

¹⁰Note that we are only trying to find a section that the solution through y intersects repeatedly. Since x intersects γ , the original section, repeatedly and $x(t_n)$ is close to y , one would hope that you can show that y must intersect γ repeatedly as well (so we wouldn’t have to look at another section to make the above argument). Unfortunately it’s not clear to me how to make this argument, but perhaps you, my dear reader, are clever enough to think of something.




Figure 26: $x(t_n)$ must eventually be in the top box but this breaks monotonicity on γ

statement that holds. If $\omega(x)$ contains a closed orbit Γ , then $\omega(x) = \Gamma$.

We will show that $d(\Phi_t(x), \Gamma) \rightarrow 0$ as $t \rightarrow \infty$. By everything we have done thus far, we know there exists an increasing sequence t_n such that $\Phi_{t_n}(x) \in \gamma$, $\Phi_{t_n}(x) \rightarrow y$ and $\Phi_t(x) \notin \gamma$ if $t_n < t < t_{n+1}$. We will use Γ to denote the orbit of y where $y \in \omega(x)$ as before.

By definition, we know that $x_n := x(t_n)$ come arbitrary close y , so we would like to say the distance between $\Phi_t(x(t_n))$ and $\Phi_t(y)$ is pretty small.¹¹ We know for any fixed t , we could prove this rigorously using continuous dependence on initial conditions (for example we know that solutions can only differ by some exponential amount as time evolves. Unfortunately, this bound is not uniform; it depends on t). This is where we use the fact that the solution through y is periodic, say with period τ . Since the x_n are pretty close to y , they should also be roughly periodic with a period of about τ . In fact this is exactly what is captured by the choice of t_n : these are the times when the solution x comes back to the section (and since the $x(t_n)$ converge we know that $x(t_{n+1})$ is close $x(t_n)$ which illustrates the ‘rough periodicity’). Since y has period τ , we should expect that t_{n+1} and t_n differ by about τ as well. But notice how we are basically done now: any t is going to lie between some t_n and t_{n+1} so we only ever need to evolve by τ (or whatever the largest difference between all the t_{n+1} and t_n is). Since we have a maximal time interval to look at, we can use continuous dependence to determine how close the initial point needs to be to y in order to ensure we remain close to y for all time in this interval. After evolving x_n by this amount (at most), we will be even closer to y , so continuous dependence will still hold! Now all that remains to do is turn the above into precise mathematics (buckle up?).

As is typical in math proofs, we will work backwards. First we find the bound for $t_{n+1} - t_n$. Let V be a neighbourhood of y such that we have the homeomorphism h from previous paragraph(s) on it. Let ϵ such that a closed rectangle of height 2ϵ centered on the x -axis and is contained in $h^{-1}(V)$ (in other words ϵ is such that $J \times [-\epsilon, \epsilon]$ is contained in $h^{-1}(V)$ where J is some closed interval itself). Let V_ϵ be the image of this rectangle under h . Since $\Phi_\tau(y) = y \in V_\epsilon$ (we take the closed interval J appropriately so y remains in V_ϵ), we know that for x_n sufficiently large we will have $\Phi_\tau(x_n) \in V_\epsilon$. Therefore $\Phi_{\tau+t}(x_n) \in \gamma$ for some $t \in [-\epsilon, \epsilon]$. Since the t_n are chosen so that successive ones intersect γ while everything between t_n and t_{n+1} doesn't, we find that $t_{n+1} - t_n \leq \tau + \epsilon$.

Now we have a bound for $t_{n+1} - t_n$. From here, as outlined above, the procedure is quite simple. Suppose $\beta > 0$ is given. By continuous dependence there exists some δ such that $|z - y| < \delta$ implies that $|\Phi_t(z) - \Phi_t(y)| < \beta$. Moreover, we can choose δ so that this holds for all t where $|t| \leq \tau + \epsilon$ (I turn especially to [Theorem 13.2](#). We are working with a C^1 function on a compact set so we know it is Lipschitz and we're working in the simpler case with $\epsilon_1 = \epsilon_2 = 0$). Since the x_n converge to y , we find some n_0 such that $|y - x_n| < \delta$ for all $n \geq n_0$. Then for $t \geq t_{n_0}$ we find n such that

$$t_n \leq t < t_{n+1}$$

¹¹Side note: $\Phi_t(x(t_n)) = \Phi_{t_n+t}(x)$. So by renaming variables, the statement above is equivalent to saying that $\Phi_{t-t_n}(x_n)$ is close to $\Phi_{t-t_n}(y)$. This is not terribly important but we will use this at the very end so hopefully it won't seem too mysterious.

With this n , we find that

$$\begin{aligned} d(\Phi_t(x), \Gamma) &\leq |\Phi_t(x) - \Phi_{t-t_n}(y)| \\ &= |\Phi_{t-t_n}(x_n) - \Phi_{t-t_n}(y)| \\ &< \beta \end{aligned}$$

Here we use the fact that $t - t_n \leq \tau + \epsilon$. Thus by taking t to be sufficiently large, we can make $d(\Phi_t(x), \Gamma)$ arbitrarily small, finally, *finally*, proving the theorem. \square

I apologise, the previous proof should probably have been split into multiple lemmas. But then again, if I was good at making decisions, this wouldn't exist so :)

17 Techniques for solving ODEs

17.1 Separation of variables

Suppose we have a differential equation of the following form

$$\frac{dx}{dt} = f(t)g(x)$$

where f, g are continuous. Suppose we are also given that $x(0) = x_0$ as our initial condition.

If $g(x_0) = 0$ then $x(t) \equiv x_0$ is a solution. So suppose that $g(x_0) \neq 0$, implying that g is non-zero in some neighbourhood around x_0 . We can thus do the following

$$\begin{aligned} \frac{dx}{dt} &= f(t)g(x) \\ \frac{x'(t)}{g(x(t))} &= f(t) \\ \int \frac{x'(t)}{g(x(t))} dt &= \int f(t) dt \\ \int \frac{1}{g(u)} &= \int f(t) dt \\ G(x) &= F(t) + c \end{aligned}$$

where in the penultimate line we substitute $u = x(t)$ and in the final line we take G to be the anti-derivative of $\frac{1}{g}$ and F to be the anti-derivative of f . Recall we are solving around $(0, x_0)$ and we know that g is non-zero in a neighbourhood of it. Thus $G'(x) = \frac{1}{g(x)}$ is also non-zero (either all positive or all non-negative) implying that G is invertible on this neighbourhood. We can thus take $x = G^{-1}(F(t) + c)$. The solution passing through (t_0, x_0) corresponds to $c = G(x_0) - F(t_0)$.

Remark 17.1. The concrete construction of the solution implies that the solution to a differential equation with separable values is unique in a small neighbourhood of (t_0, x_0) , given that $g(x_0) \neq 0$.

The above steps are often written more simply as

$$\begin{aligned} \frac{dx}{dt} &= f(t)g(x) \\ \frac{1}{g(x)} dx &= f(t) dt \\ \int \frac{1}{g(x)} dx &= \int f(t) dt \\ G(x) &= F(t) + c \end{aligned}$$

17.1.1 Example

A simple example of using separation of variables is [Equation 2.3](#) (no more guesswork required!). A more interesting example can be found in [Section 4](#).

17.2 Homogeneous Functions

We say a function F on \mathbb{R}^2 is homogeneous of degree α (with $\alpha \in \mathbb{R}$) if $F(tx, ty) = t^\alpha F(x, y)$.

Suppose we are given

$$P(x, y)dx + Q(x, y)dy = 0 \quad (17.1)$$

Where P and Q are homogeneous and of the same degree. Then substituting $y = xu$ (giving us $dy = udx + xdu$) or $x = yv$ (giving us $dx = vdy + ydv$) changes our differential equation into one where we can separate variables.

17.2.1 Example

Suppose we have the following equation

$$\underbrace{\left[xe^{\frac{y}{x}} - y \sin\left(\frac{y}{x}\right) \right]}_{P(x,y)} dx + \underbrace{x \sin\left(\frac{y}{x}\right)}_{Q(x,y)} dy = 0 \quad (17.2)$$

It is easy to verify P and Q are homogeneous functions (of degree 1). We substitute $y = xu$ (either substitution works but often one is easier than the other. In this case it seems quite apparent that one would want to replace $\frac{y}{x}$). The substitution gives us

$$\begin{aligned} [xe^u - xu \sin(u)]dx + x \sin(u)(udx + xdu) &= 0 \\ e^u dx + x \sin(u)du &= 0 \end{aligned}$$

(note we divide by x without worry since the original equation [Equation 17.2](#) doesn't allow for $x = 0$) The variables can now clearly be separated.

17.3 Exact Differential Equations

Suppose we have an ODE of the form

$$P(x, y)dx + Q(x, y)dy = 0$$

is exact if there exists a function $f(x, y)$ such that

$$\frac{\partial f}{\partial x} = P(x, y), \quad \frac{\partial f}{\partial y} = Q(x, y)$$

The general solution to such ODEs is given by the one-parameter family

$$f(x, y) = c$$

It would be nice to know when an ODE is exact. Luckily we have the following theorem which not only tells us when an equation is exact but even gives a construction for the function f .

Theorem 17.2 Suppose we have an ODE of the form

$$P(x, y)dx + Q(x, y)dy = 0$$

where P, Q are defined on a simply connected region and are C^1 . Then the equation is exact if and only if

$$\frac{\partial P}{\partial y} = \frac{\partial Q}{\partial x}$$

Proof. The forward direction follows from a theorem in analysis. Namely if a function f is smooth (or even just C^1) then

$$\frac{\partial^2 f}{\partial x \partial y} = \frac{\partial^2 f}{\partial y \partial x}$$

Thus we need show the reverse direction. Thus we assume that

$$\frac{\partial P}{\partial y} = \frac{\partial Q}{\partial x}$$

Suppose such an f were to exist. Then we know that f must satisfy

$$\frac{\partial f}{\partial x} = P(x, y) \Rightarrow f(x, y) = \int_{x_0}^x P(s, y) ds + R(y)$$

Typically we get a constant term when integrating, but in this case the constant term depends on y , hence why R becomes a function of y . By assumption the derivative of f with respect to y is Q . Thus we get that

$$\begin{aligned} Q(x, y) &= \frac{\partial f}{\partial y} = \frac{\partial}{\partial y} \left(\int_{x_0}^x P(s, y) ds + R(y) \right) \\ &= \int_{x_0}^x \frac{\partial}{\partial y} P(s, y) ds + R'(y) \\ &= \int_{x_0}^x \frac{\partial}{\partial x} Q(s, y) ds + R'(y) \\ &= Q(x, y) - Q(x_0, y) + R'(y) \end{aligned}$$

where of course x_0 is some arbitrary constant in the domain. We then conclude that $R'(y) = Q(x_0, y)$. Therefore

$$R(y) = \int_{y_0}^y Q(x_0, s) ds$$

(we ignore the integration constant for now since account for it later when giving the general solution to the ODE). Therefore

$$f(x, y) = \int_{x_0}^x P(s, y) ds + \int_{y_0}^y Q(x_0, s) ds$$

All that remains to check is that this f does indeed satisfy the conditions. Now it is clear that

$$\frac{\partial f}{\partial x} = P(x, y)$$

The other one takes slightly more work

$$\begin{aligned} \frac{\partial f}{\partial y} &= \frac{\partial}{\partial y} \int_{x_0}^x P(s, y) ds + Q(x_0, y) \\ &= \int_{x_0}^x \frac{\partial}{\partial y} P(s, y) ds + Q(x_0, y) \\ &= \int_{x_0}^x \frac{\partial}{\partial x} Q(s, y) ds + Q(x_0, y) \\ &= Q(x, y) - Q(x_0, y) + Q(x_0, y) \\ &= Q(x, y) \end{aligned}$$

□

17.3.1 Example

Suppose we have the equation

$$(2x + y \cos x)dx + (2y + \sin x - \sin y)dy = 0 \quad (17.3)$$

In this case we have $P(x, y) = 2x + y \cos x$ and $Q(x, y) = 2y + \sin x - \sin y$. Since

$$\frac{\partial P}{\partial y} = \cos x = \frac{\partial Q}{\partial x}$$

we can conclude using the previous theorem that the differential equation is exact. Thus we can construct f as given by the proof as well. We start with

$$f(x, y) = \int P(x, y)dx + R(y) = x^2 + y \sin x + R(y)$$

Then we know that

$$\begin{aligned} \frac{\partial f}{\partial y} &= Q(x, y) \\ \sin x + R'(y) &= 2y + \sin x - \sin y \end{aligned}$$

Therefore

$$R(y) = y^2 + \cos y$$

Finally we conclude that the general solution to [Equation 17.3](#) is

$$f(x, y) = x^2 + y \sin x + y^2 + \cos y = c$$

17.3.2 Integrating Factors

Sometimes an equation is not exact, but we can make it exact by multiplying it with a function. Such a function is called the integration factor. As an example consider the equation

$$(t^2 x - t)dx + xdt = 0 \quad (17.4)$$

Suppose there was some function $h(t)$ (we assume h is a function of t because we know this works. In general it's difficult to know what the integration factor should be a function of). Then we would have

$$(h(t)(t^2 x - t))dx + (xh(t))dt = 0$$

which we assume to be exact. This means that

$$\begin{aligned} \frac{\partial}{\partial t}(h(t)(t^2 x - t)) &= \frac{\partial}{\partial x}(xh(t)) \\ h'(t)(t^2 x - t) + h(t)(2tx - 1) &= h(t) \\ \frac{h'(t)}{h(t)} &= \frac{2 - 2tx}{t^2 x - t} \\ &= -\frac{2}{t} \end{aligned}$$

Since the right hand side is a function of t , we can integrate to conclude that

$$\log(h(t)) = -2 \log t$$

or in other words that

$$h(t) = \frac{1}{t^2}$$

Similar manipulations can be performed if h is function of $x, xt, \frac{x}{t}, \frac{t}{x}$, etc. Often it is not obvious what h should be a function and requires some trial and error. Hence why although we know that any ODE $x' = f(x, t)$ with $f \in C^1$ in a neighbourhood of (x_0, t_0) admits an integration factor, in principle this is a useless technique because it is incredibly hard to find the integration factor.

However there do exist some special cases where integration factors *can* be found and these serve as useful examples. Suppose we have an equation of the form

$$\frac{dy}{dx} + P(x)y = Q(x) \quad (17.5)$$

Since the derivative and y are raised to the power of 1, we call this a linear differential equation and specifically a linear differential equation of order 1. One way of solving this equation is to solve the homogeneous version (i.e. set $Q = 0$) using separation of variables and then solve the ODE using variation of constants. A second method, however, is to use integration factors.

We can first rewrite the equation to get

$$[P(x)y - Q(x)]dx + dy = 0$$

Suppose we had an integrating factor $u(x)$ (in this case we know that the integrating factor is always a function of the dependent variable, in this case x). Then we would have

$$\begin{aligned} \frac{\partial}{\partial y}[u(x)P(x)y - u(x)Q(x)] &= \frac{\partial}{\partial x}u(x) \\ u(x)P(x) &= u'(x) \\ \frac{u'(x)}{u(x)} &= P(x) \\ u(x) &= e^{\int P(x)dx} \end{aligned}$$

Note that when integrating the exponent we don't need to worry about the integration constant since we don't need the general solution. Substituting this integration factor into [Equation 17.5](#) we get

$$\begin{aligned} e^{\int P(x)dx} \frac{dy}{dx} + P(x)e^{\int P(x)dx} y &= e^{\int P(x)dx} Q(x) \\ \frac{d}{dx}(ye^{\int P(x)dx}) &= e^{\int P(x)dx} Q(x) \\ y &= e^{-\int P(x)dx} \left(\int e^{\int P(x)dx} Q(x) \right) + ce^{-\int P(x)dx} \end{aligned}$$

We apologise for the horror we have bestowed upon the reader.

17.3.3 Bernoulli Equations

Bernoulli equations are slight variations of the linear differential equations seen before which aren't themselves linear equations but can be turned into ones. Specifically, a Bernoulli equation is of the form

$$\frac{dy}{dx} + P(x)y = Q(x)y^n \quad (17.6)$$

In the case of $n = 1$ we can solve by separation of variables so let us assume that $n \neq 1$. Then we multiply both sides with $(1 - n)y^{-n}$ to get

$$(1 - n)y^{-n} \frac{dy}{dx} + (1 - n)y^{1-n}P(x) = (1 - n)Q(x)$$

Now if we define $u = y^{1-n}$ we see the above equation becomes

$$\frac{du}{dx} + (1 - n)P(x)u = (1 - n)Q(x)$$

which is a linear differential equation of order 1 we know how to solve.